

Ruggero Rossi
r.rossi@opencomplexity.com

A Gentle Principled Introduction to Deep Reinforcement Learning

version 1.1.3t
© 2022-2025 Ruggero Rossi

Preface

This book is an attempt to introduce Deep Reinforcement Learning to those who possess a solid knowledge of Deep Learning. The aim is to explain the functioning of selected Deep Reinforcement Learning algorithms that are both practically effective and didactically representative, providing the reader with a theoretical justification and a balanced level of proofs and mathematical details, without losing sight of the general picture. The order of chapters has been chosen so to present at first the algorithms that are preparatory to the others. Also, when possible, the chapter order tries to follow a hypothetical path that could lead naturally from one algorithm to another that represents an evolution of the former, or at least from an algorithm to another one that shares some common points. Anyway, since many algorithms have common grounds with more than one other algorithm, sometimes other different chapter orders could be possible, in those cases I just picked my favourite one among them. When choosing which mathematical proofs to insert, I decided to make room only for (I) the ones that are essentials to understand the nature of the algorithm, and (II) the ones that are essential for the implementation of some detail of the algorithm and that in original papers may appear somehow obscure to the novice, and hence needed a more detailed explanation because the referenced paper alone would have been too hard to the unexperienced reader. The prerequisite to comfortably read this book is the knowledge of Deep Learning, that implies knowledge of basic Calculus, Linear Algebra, Probability and Statistics.

Contents

Preface	2
1. Introduction	7
1.1 Reinforcement Learning Definition	7
1.2 Reinforcement Learning Formalization	11
1.3 Variety of Methods	15
2. Q-Learning	17
2.1 Essential Algorithm	17
2.2 Minibatch Improvement	21
2.3 Saving Transitions into a Buffer	21
2.4 Bootstrapping and Convergence	22
2.5 Increasing Gradient Stability: Target Network	23
2.6 Polyak Averaging	24
2.7 Action-Value Overestimation	26
2.8 Double Q-Learning	29
2.9 Issues with Continuous Actions	32
2.10 Sarsa	33
3. Policy Gradient	36
3.1 Policy Based Methods	36
3.2 Policy Gradient Basics	37
3.3 Rewards-To-Go	40
3.4 Bootstrapping the Value	44
3.5 Beyond Rewards-To-Go	45
3.6 Advantage Actor-Critic (A2C)	46
3.7 Benefits of Policy Gradients	49
3.8 Improvements	50
4. Off-Policy Policy Gradient	51
4.1 Motivations for Off-Policy	51
4.2 The Approximated Off-Policy Policy Gradient	52
4.3 Alternative Derivation	56
5. Generalized Advantage Estimation	61
6. Natural Policy Gradient	64
6.1 Motivation	64
6.2 Constraining the New Policy	65
6.3 Derivation of the Formula	68

7. Trust Region Policy Optimization	72
7.1 Problem formalization	72
7.2 Ideal Objective.....	74
7.3 Local Approximation of Objective	75
7.4 Objective Lower Bound.....	75
7.5 Monotonic Improvement	82
7.6 Surrogate Objective	83
7.7 Constraining the Kullback-Leibler Divergence	86
7.8 Inverting the Fisher Information Matrix via Conjugate Gradient.....	88
7.8.1 Product of Average Hessian of KL	90
7.8.2 Average Gradient of Product by KL Gradient	95
7.9 Using the Conjugate Gradient.....	96
7.10 Vine Sampling Scheme	98
8. Proximal Policy Optimization	100
8.1 Base Idea	100
8.2 Surrogate Objective Clipping	100
8.3 Multiple Steps of Gradient Ascent.....	102
8.4 Kullback-Leibler divergence check.....	102
9. Deep Deterministic Policy Gradient	105
9.1 Base Idea	105
9.2 Action-Value Gradient.....	106
9.3 On-Policy Deterministic Actor-Critic.....	110
9.4 Off-Policy Deterministic Actor-Critic.....	112
9.5 Compatible Function Approximator.....	114
9.6 Deep Deterministic Policy Gradient	119
10. TD3 (Twin Delayed DDPG)	122
10.1 Base Idea	122
10.2 Twin Networks for Q-Value estimation.....	122
10.3 Computing Target Actions	122
10.4 Delayed Policy Updates.....	123
11. Soft Actor-Critic.....	126
11.1 Base Idea	126
11.2 Entropy Regularization	127
11.3 Policy Update	129
12. Reward Shaping and Curriculum Learning.....	132
13. Imitation Learning.....	134

13.1 Behavioral Cloning	134
13.2 Dataset Aggregation (DAgger).....	135
13.3 Implementation	136
14. Reinforcement Learning from Human Feedback	140
14.1 Introduction.....	140
14.2 Method	140
14.3 Implementation	142
15. Reinforcement Learning in Large Language Models	146
15.1 Introduction.....	146
15.2 RLHF in LLMs.....	147
15.2.1 Analogy between LLM and RL	147
15.2.2 Differences from standard RLHF: Human Evaluation	148
15.2.3 Reward Model.....	151
15.2.4 Value function	151
15.2.5 Kullback-Leibler Penalty for the Reward in RL Stage	152
15.2.6 Reward-to-go in RL Stage.....	153
15.2.7 Training process.....	155
15.2.8 Pretrain Likelihood Term in RL Objective	157
15.2.9 PPO implementation details	157
15.3 Group Relative Policy Optimization.....	158
15.3.1 Advantage surrogate without Value Function: Outcome Supervision.....	158
15.3.2 Process Supervision	160
15.3.3 Kullback-Leibler Divergence for the Loss	161
15.4 Decoupled Clip and Dynamic sAmpling Policy Optimization (DAPO).....	162
15.4.1 No Kullback-Leibler Divergence penalty	162
15.4.2 Higher Clipping Threshold.....	163
15.4.3 Dynamic Sampling	163
15.4.3 Token-Level Policy Gradient Loss.....	163
15.4.4 Overlong Reward Shaping	164
15.5 RLAIF (Constitutional AI)	165
15.5.1 Initial Setup	165
15.5.2 Supervised Phase	166
15.5.3 Reinforcement Learning Phase	167
15.6 Other RL without Human Feedback.....	169
16. Afterword	170
Appendix A: Information Theory Refresh.....	171

A.1 Kullback-Leibler Divergence.....	171
A.2 Total Variation Distance	173
A.3 Score (of the Log-Likelihood)	175
A.4 Fisher Information Matrix.....	177
A.5 Fisher Information Matrix equivalence to negative expectation of Hessian Matrix of log-probability	178
A.6 Kullback-Leibler divergence approximation by second order Taylor expansion using Fisher Information Matrix.....	180
A.7 Relationship between the Hessian of Kullback-Leibler divergence and Fisher Information Matrix.....	182
A.8 Information Entropy	184
A.9 Differential Entropy of the Gaussian	189
A.10 Mutual Information	192
Appendix B: Conjugate Gradient Algorithm	193
B.1 Aim.....	193
B.2 The Equivalent Problem	194
B.3 The Steepest Descent.....	195
B.4 The Conjugate Gradient	197
Appendix C: Importance Sampling	200
References	202

1. Introduction

1.1 Reinforcement Learning Definition

Reinforcement Learning is a method for an artificial agent to do automated learning using the outcomes caused by its actions. A definition from [Sutton & Barto 2018] reports: “Reinforcement learning is learning what to do - how to map situations to actions - so as to maximize a numerical reward signal. The learner is not told which actions to take, but instead must discover which actions yield the most reward by trying them”.

So, paraphrasing Sutton and Barto, Reinforcement Learning is: learning which action to take, depending on the current situation, with the purpose of maximizing the total reward.

In that definition we have some of the base components of a Reinforcement Learning system:

- an *agent* able to decide which action to take
- a *state* of the world – or better an *observation* about the state of the world
- a *reward*

The *agent* is obviously the artificial intelligence that is trying to learn how to act, it may be just a program inside a simulation as well as a physical robot in a real environment.

The *observation* about the state of the world is the data that the agent possesses about the state of world, and that may be potentially incomplete, noisy, or delayed. The observation may be given to the agent by the simulator software, in case of simulated environment, or may be obtained by the agent through sensors in case of physical robot.

The *reward* is a number that represents how well or how badly the agent is behaving (usually it is positive for good rewards and negative for bad rewards). It is something similar to the concept of utility. In case of a simulated environment the reward is given to the agent by the simulator itself, together with the timed information updates such as the observation of the state of the world. In case of physical robot in a real environment instead, a software layer must be programmed in the robot in order to understand what is happening in the world and calculate a reward that is related to how much good or bad the situation is.

In both cases the reward is a number whose value is communicated to the agent by a human-made program: either in the simulation or in the robot software. There is not such a thing in the environment as a “reward number”, the reward is just a fictitious number that permits to do numerical optimization (humans and animals feel pain and pleasure and that may be seen as a reward system that has been evolved, but for our artificial agents it has to be built by us). Anyway in some cases that number may be somehow directly available from the real world or from the final purpose objective: think about an artificial intelligence trying to maximize the gains in the stock market, the reward may be directly the market returns, or think about of a robot trying to pick up trash from the ground, a reward proportional to garbage weight may be given every time that garbage is actually collected.

If the task to be accomplished by the agent is never ending it is said to be “infinite-horizon”, otherwise if a begin and an end of the sequence of actions can be identified, it is said “finite-horizon” and we may call “episode” what happens within the begin and the end. For instance, if a robot is trying to shoot a basketball inside a hoop, an episode starts when the robot tries to grab the ball and ends when the ball, after being shot, lands, either passing through the hoop or outside it. An episode contains a “*trajectory*” made of a sequence of states and actions. The sum of the rewards obtained in an “*episode*”, possibly discounted by a time distance value, is called “*return*” (the time discount represents the preference of receiving rewards immediately rather than in the future).

Another useful concept when talking about reinforcement learning is the concept of “*Policy*”: with that term we mean a strategy that associates an action to a world state, for each possible state. That means that in each situation an agent is, the Policy will tell him what action to take. A good Policy is one that makes the agent collect a good amount of rewards, an optimal policy is one that makes the agent collect the maximum amount of rewards (there may be more than one optimal policy). A policy may be stochastic and define, for each state, a probabilistic mixture of actions to take, such as “take action A with probability 80%, take action B with probability 20%”.

So, what finally an agent wants to learn is an optimal policy (or at least one good enough ! If finding the optimal policy takes too much time or resources, it may be better to find just a decent policy and use that, which is the topic “exploration vs exploitation” discussed later).

The agent may or may not have knowledge of how the world works and what happens if he takes a certain action, in which state he will end up and what reward will be obtained because of that action. When the agent has that kind of knowledge it is possible to apply the so-called *model-based* reinforcement learning algorithms. The model of the world is generally considered stochastic, with deterministic models being just particular cases. The model may

be a fixed, prior knowledge of the agent. If instead the agent has not a prior model of the world mechanics, it is possible either to learn one in the process, or to use *model-free* algorithms that don't require a model of the world. In some case the model of the world is available (at least partially), in some other cases the learning algorithm may be trying to learn a model of the world through experience and then use that model to apply a model-based algorithm. When the model of the world is complete, and when there are a finite number of states it is theoretically possible to calculate the best action for each state just with numerical optimization (e.g. with dynamic programming), using the “Bellman Equation” [Bellman 1952] [Sutton and Barto 2018] without the need to make the agent taking real actions. Usually that is not the case because often the model of the world is not known (and sometimes difficult to be learned) and there is an infinite number of states, hence agents must advance through trial and errors to discover what is the best action in which situation.

Trying actions in the environment with the purpose of learning exposes the agents to the risk of obtaining very bad rewards: in a real environment that would mean damaging seriously the robot, or even worse creating risky situations outside the experimental environment. So, there is a trade-off between exploration (trying new actions to figure out if they bring better rewards) and exploitation (doing only the actions that so far showed to be sufficiently rewarding). Exploration exposes to risks but permits to optimize the policy, exploitation permits to gain the fruits of past exploration but avoids further learning. Agents usually are given a greater degree of exploration at the beginning, and the exploration degree is decreased as the learning progresses, favouring exploitation.

Some Reinforcement Learning algorithms (but not all) use the concept of “Value Function”: a function that takes the current state (or current observation) as input and evaluates the hypothetical goodness of it.

The evaluation score returned by the Value Function is related to how good it is expected to be the future situation in the long run, not just in the immediate next moments. If a certain state is expected to make the agent obtaining a reward, the value function will take in account the reward expected in that situation and all the rewards expected from what hypothetically may be the future situations if the agent acts as its “policy” suggests, possibly discounting the future rewards by a time-dependent distance value (rewards further in time may weight less than immediate rewards).

So, as I will describe later formally with the Bellman Equation, the Value function is a recursive concept: saying that the value at some state depends not only on the reward obtained in that state but also on all the rewards that may be obtained afterwards, is the same of saying that

the value at a certain state depends not only on the goodness of that state “per se” but also on the values of the states that are reachable from there. This also implies that the value function depends on the policy: two different policies may reach different states from the same starting point and hence they may have two different value functions. So, in general each value function is tied to a certain policy: changing the policy changes the value function.

There are two different types of value functions: proper “value functions” and “action-value functions”.

The proper “value function”, returns the value of a state, considering that the agent will follow the policy from that state, so its only input is the state (and implicitly the policy). It is sometimes called also “state-value” function.

The “action-value” function takes as input the state and also an action, and returns a value considering that the agent will take that particular action in that state (even if it’s a different action with respect to what the policy would suggest), and after that it will follow his policy.

It has to be noted that when in literature it is used the term “Value Function” it may be referred to two very different things: the first is the value function as described above, the second is the estimate of the value function that an agent is approximating, and that may be far from the real value (and it should be better referred to with the terms “value function estimate”). When there is a finite number of states the estimate of the value function is often saved in a table with an estimate for each state (or, in case of “action-value”, an estimate for each state-action pair), when instead there is a great number of states that cannot be done and a function approximator is used, such as a neural network. That may be the case when the state/observation is described by continuous variables which if discretized in a table would need too much memory, or which would lose relevant details.

The term “Deep Reinforcement Learning” is used to describe Reinforcement Learning methods that use neural networks, especially neural networks with many hidden layers (“deep”). For instance, neural networks may be used to approximate value function estimates, or, as we will see later, for policy functions. The usage of deep neural networks allows to learn complex policies that are not possible with linear methods as function approximators. Moreover, they allow to learn end-to-end, that means having a neural network which, without the hand-made engineering of different other software modules, receives the sensory input and internally learns all the necessary functionalities to process it and take the output action, as opposed to having a separate module that analyses the input, a next module that plans what to do, a further module that learns how to actuate the robot engines, etc.




Figure 1.1

A hypothetical Policy neural network with 16 input units, three hidden layers of respectively 12, 10 and 5 units, and 2 output units. The input state may consist of the values of 16 sensors that are fed into the 16 input units, and after the computation the values of the two output units will represent the values of 2 different action dimensions. For instance, imagine a simple robotic arm, the first output unit may indicate how much to rotate clockwise/counterclockwise, the second output unit may indicate how much to extend. Figure created with the online tool at <https://alexlenail.me/NN-SVG/>.

1.2 Reinforcement Learning Formalization

A formalization of the Reinforcement Learning problem may be done expressing it as a Markov Decision Process (MDP). A MDP is a way to describe how an agent passes from one state to another and possibly obtains a reward as a consequence of its actions. One requirement is that states must have the “Markov Property”: the probabilities to move from one state to another and to obtain rewards depend only on the knowledge of the current state and the

chosen action, and not on any previously accessed state. In other words, a state description has the Markov property if it contains all the necessary information from the past and from the present to determine the transition probabilities to other states and the rewards.

This means that theoretically, non-Markov states can be turned in Markov states if all the history of past states and interactions are added to them (but this way to describe states as long sequences of the past are not simple to be managed by algorithms).

In a MDP, the mechanism that makes an agents passing from a state to another as a consequence of its actions is described by a “*transition function*” that is generally probabilistic and it is not necessarily known by the agent (the agent may just experience the change in state after an action). The mechanism that assigns a reward depending on a certain state, an action executed in that state, and a consequent state, is described by a “*reward function*”, that is not necessarily known by the agent (the agent may just notice the obtained reward after every action).

The theoretical exposition of RL that follows, when a different source is not referenced, is taken from [Sutton & Barto 2018] and [Achiam & OpenAi 2020A], with some change of variables names to have a consistent description.

Formally, following OpenAI definition, a MDP is a 5-tuple $\langle S, A, R, P, \rho_0 \rangle$:

- S is the set of all states
- A is the set of actions
- $R : S \times A \times S \rightarrow \mathbb{R}$ is the reward function, with $r_t = R(s_t, a_t, s_{t+1})$
- $P : S \times A \rightarrow \mathcal{P}(S)$ is the transition function, with $P(s'|s, a)$ being the probability of transitioning from state s to state s' after taking action a , and $\mathcal{P}(S)$ is the set of probability measures on S
- ρ_0 is the distribution of initial state

The policy that decides which action to take depending on the state, is a stochastic function π , such that if at time t the state is s_t , it will return a probability distribution over the actions, from which it may be sampled the action to take a_t :

$$a_t \sim \pi(\cdot | s_t) \tag{1.1}$$

A sequence of states and actions is called “trajectory”, identified by the symbol τ

$$\tau = (s_0, a_0, s_1, a_1, \dots) \quad (1.2)$$

The return $G(\tau)$ of a trajectory is the (potentially infinite) sum of all rewards of the trajectory, with $\gamma \in [0,1]$, and it is time-discounted if $\gamma < 1$:

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t \quad (1.3)$$

Given a policy π , the probability of following any trajectory τ made of T steps is:

$$P(\tau|\pi) = \rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi(a_t|s_t) \quad (1.4)$$

And the expected return $J(\pi)$ is :

$$J(\pi) = \int_{\tau} P(\tau|\pi) G(\tau) = E_{\tau \sim \pi}[G(\tau)] \quad (1.5)$$

Where a notation like $E_{x \sim q}[f(x)]$ means: Expected Value of $f(x)$ where x follows the probability density function (or probability mass function) $q(x)$.

Later we will see also a notation like $E_{x \sim q}[f(x)|y = m, u = n]$ that means: Expected Value of $f(x)$ where x follows the pdf/pmf $q(x)$, given that y is equal to m and u is equal to n .

The problem of Reinforcement Learning is to find the policy that maximizes that expectation in eq. 1.5. Hence:

$$\pi^* = \arg \max_{\pi} J(\pi) \quad (1.6)$$

Where π^* is the optimal policy.

The Value Function for a given policy π is:

$$V^\pi(s) = E_{\tau \sim \pi}[G(\tau)|s_0 = s] \quad (1.7)$$

The Action-Value function (also named “Q-Value”) is:

$$Q^\pi(s, a) = E_{\tau \sim \pi}[G(\tau)|s_0 = s, a_0 = a] \quad (1.8)$$

It must be noted that:

$$V^\pi(s) = E_{a \sim \pi}[Q^\pi(s, a)] \quad (1.9)$$

As previously anticipated, the value function is a recursive concept because the value of a state is dependent on the value of the next state, recursively. The Bellman Equations express this relationship:

$$V^\pi(s) = E_{\substack{a \sim \pi \\ s' \sim P}}[R(s, a, s') + \gamma V^\pi(s')] \quad (1.10)$$

$$Q^\pi(s, a) = E_{s' \sim P}[R(s, a, s') + \gamma E_{a' \sim \pi}[Q^\pi(s', a')]] \quad (1.11)$$

Using eq. 1.9 in 1.11 we obtain:

$$Q^\pi(s, a) = E_{s' \sim P}[R(s, a, s') + \gamma V^\pi(s')] \quad (1.12)$$

Some algorithms use the “Advantage Function” [Baird 1993], that is a function indicating how much it is better or worse to take a certain action (and after that follow the policy) instead of taking the action suggested by the policy, i.e. it indicates the relative advantage of taking a certain action with respect to the policy. The advantage function is the following:

$$A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s) \quad (1.13)$$

It must be noted that an alternative way of denoting $J(\pi)$ is:

$$J(\pi) = E_{\tau \sim \pi, s_0 \sim \rho_0} [V^\pi(s_0)] \quad (1.14)$$

1.3 Variety of Methods

Among RL methods, a first distinction can be made between **model based** and **model free** methods. Model based methods use a model of the world, consisting generally in the knowledge of the reward and transition functions, and the initial state distribution: R, P, ρ_0 .

This model is either already existing from a priori knowledge or learnt by the agent.

Model free methods instead do not use a model (generally in real environments a model is not available), and are agnostic of initial state distribution, transition function, and reward function (they are aware only of the rewards that the agent receive when it receives any, without knowing the underlying function or rule).

A second distinction is between on-policy and off-policy methods: **on-policy** methods evaluate and improve the same policy that is followed by the agent during training, while in **off-policy** methods the training is done following a policy that is different from the one that is evaluated or improved. In off-policy methods the policy that is being learned is named “*target policy*” and the policy that is used to generate samples is named “*behavior policy*”. A particular version of off-policy reinforcement learning is the **offline** reinforcement learning, where the data needed to feed the algorithm has been previously collected and no new data is produced [Levine et al. 2020].

A third distinction may be made between value based methods and policy based methods. In **value based** methods the algorithm aims at estimating a value function or an action-value function and consequently modify the policy to prefer actions that lead to better values. The “*Policy Improvement Theorem*” [Sutton & Barto 2018, Ch.4] guarantees that modifying the policy for a state in a way that obtains a value improvement in that state, strictly improves the expected return $J(\pi)$ and hence improves the overall policy.

If a world model is not available, “proper” value functions (i.e. state-only values) do not contain enough information to improve the policy, hence the only model-free value based methods are

those that use action-values: so they can choose the action that is associated with the best outcome for each state.

In **policy based** methods instead, the action is directly chosen by a policy function (that outputs a probability distribution over actions) and value or action-value functions may not be present, or may be used only to improve learning.

A fourth distinction of RL methods can be done between “tabular” and “approximate” methods. **Tabular** methods are those whose state space is composed by a finite set of discrete states, that are hence representable in memory by tables, for instance the Value Function estimation may be simply a table that associates a state to a value. **Approximate** methods are those whose state space is continuous or too big, and so its representation in the algorithm is made by some function approximation. For instance, the Value Function could be approximated (or estimated) by a neural network that takes the information about the state/observation (e.g. sensor data, webcam frame image, etc.) as input, and computes an approximate value as output. It is common to use convolutional neural networks as approximators when the state/observation input is an image. While tabular and approximate methods share much theory and algorithms, they also have important differences, the most notable of which is the fact that an approximate value based algorithm, if it is also using bootstrapping (i.e. using the approximate value function estimation as a replacement for the true value of value function in the Bellman equation, or equivalently using the approximated q-value estimation instead of true q-value) and if it is off-policy, it is not guaranteed to be convergent, while a tabular method would be, see [Sutton and Barto 2018, Ch. 11].

When a deep neural network is used as a function approximator in a RL algorithm, it is said to be Deep Reinforcement Learning.

2. Q-Learning

2.1 Essential Algorithm

An example of (action-)value based method is the Q-Learning [Watkins 1989], an algorithm for discrete actions. Let us see how it works.

The basic idea of Q-Learning is to estimate the value of any action-state pair, and then the target policy would be the policy that, given a state, selects the action with greater action-state value estimate. But the important detail in all this is that during learning, Q-Learning “bootstraps”, that means that it estimates the value of a past action summing the obtained reward for the past action and the estimate of the value function of the next state. This will be explained better later.

Q-Learning is an off-policy method, that means that the behavior policy (the policy used to generate trajectories) does not have to coincide with the target policy (the policy that we want to improve until it is optimal). In fact they could be very different, but for practical reasons it is better that they are not too much: since a function approximator is used to evaluate state-action value estimates, it is better to train that approximator on state-action samples coming from a state-action distribution that is similar to the one of the target policy. This is because the approximation capacity of a function approximator is limited and it is more precise when it is run on inputs similar to the ones used for training.

For this reason in Q-Learning the policy followed by agents during training (the *behavior policy*) is often made not too different from the target policy: the behavior policy sometimes chooses the action with the greatest action-value (exploitation) and sometimes a random action (exploration). A way to do that is to have an ϵ -greedy policy: with probability ϵ a random action is taken, with probability $1 - \epsilon$ instead the action with greater action-value (Q-value) is taken. Usually, ϵ is progressively decreased during learning.

In Q-Learning the Q-value is computed with respect to a policy (the *target policy*) that would always choose the best action, and never a random one. This is expressed with the formula:

$$Q^*(s, a) = E_{s' \sim P}[R(s, a, s') + \gamma \max_{a'} Q^*(s', a')] \quad (2.1)$$

I wrote Q^* instead of Q^π to make explicit that it is not computed over a generic policy π but ideally over the optimal policy (the asterisk means “optimal policy”). You can compare the differences between equations 2.1 and 1.11, where it was:

$$Q^\pi(s, a) = E_{s' \sim P} [R(s, a, s') + \gamma E_{a' \sim \pi} [Q^\pi(s', a')]]$$

(remarkably, in Q-Learning the target policy always selects the best action, so it has $\gamma \max_{a'} Q^*(s', a')$ instead of $E_{a' \sim \pi} [Q^\pi(s', a')]$).

Hence the Q-Value is referred to a tentatively optimal policy, that is asymptotically better or equal than the behavior policy. This makes Q-Learning formally an off-policy method because it evaluates a different policy than the one used to generate trajectories.

In Q-Learning the true Q^* function is unknown, and the algorithm tries to estimate it during the course of learning. So, for the Q-value estimate instead of Q^* we will use the symbol Q_ψ that means that it is an estimate using parameters ψ .

$$Q_\psi(s, a) = E_{s' \sim P} [R(s, a, s') + \gamma \max_{a'} Q_\psi(s', a')] \quad (2.2)$$

The execution of Q-Learning is as follows: at the beginning all $Q_\psi(s, a)$ for all states and actions are set to arbitrary values (if Q_ψ is in tabular form the values can be set close to zero for faster convergence, while if Q_ψ is a function approximator such as a neural network the parameters ψ may be set to a small random noise such as in [Glorot and Bengio 2010] or [Saxe et al. 2013]). Then using the behavior policy (for instance the ϵ -greedy policy specified before) an action a is taken, new state s' is reached, and the reward $r = R(s, a, s')$ is obtained.

Now, if we want our estimate of the Q function to respect the Bellman Equation 2.2, we have to make $Q_\psi(s, a)$ move towards the right-hand side value, to be more similar to it. So, calling the “target value” as y :

$$y = r + \gamma \max_{a'} Q_\psi(s', a') \quad (2.3)$$

Obviously, if s' is a terminal state, i.e. the trajectory ends there and no action a' is taken, the target value is only equal to r , since we must take in account that no additional rewards are possible to be obtained from that transition.

If Q_ψ is in tabular format, to move it closer to the target value, it is updated in the following way:

$$Q_\psi(s, a) \leftarrow Q_\psi(s, a) + \eta (y - Q_\psi(s, a)) \quad (2.4)$$

Where η is the learning rate. This is called “*Temporal-Difference Learning*” (or “TD Learning”) because the estimate of the q-values Q_ψ is updated using the difference between the bootstrapped target value after having done the action and the q-value estimate before doing the action.

If instead Q_ψ is not in tabular format but it is estimated using a function approximator such as a neural network, a Loss is computed for the discrepancy of the approximated Q-value :

$$L = (y - Q_\psi(s, a))^2 \quad (2.5)$$

Then the parameters of Q_ψ are updated with a gradient descent step (or with batch gradient descent) using the gradient of that Loss function.

It must be noted that in equation 2.5 the value y is to be considered as a constant, not as a value depending on Q_ψ , so that the gradient of the loss with respect to ψ will not include the effect of $\max_{a'} Q_\psi(s', a')$ but only of $Q_\psi(s, a)$. For this reason, the Q-Learning update is considered “semi-gradient”, and not complete gradient descent. In other words, the update in case of function approximation would be the following:

$$\psi \leftarrow \psi + \eta (y - Q_\psi(s, a)) \nabla_\psi Q_\psi(s, a) \quad (2.6)$$

Then, the algorithm continues in the same way both in tabular and approximated version: a new action a' is taken following the policy, and Q_ψ is updated either with the tabular or the

gradient descent method using the new $y' = r' + \gamma \max_{a''} Q_{\psi}(s'', a'')$ as described above, and so on.

In the following pseudo-code algorithms, the variable *is_terminal* is a Boolean which indicates if reached state s_{i+1} is a terminal state, i.e. the trajectory from which the transition is sampled ends there.

Algorithm 2.1 Deep Q-Learning (simple version)

Require: Step size η

Require: Initialize parameters ψ of network Q_{ψ} with small random values

Do until (supposed) convergence:

using some behavior policy collect a transition $\langle s_i, a_i, s_{i+1}, r_i, is_terminal \rangle$

if not *is_terminal* then:

$$\text{padding-left: 80px; } y_i = r + \gamma \max_{a_{i+1}} Q_{\psi}(s_{i+1}, a_{i+1})$$

else:

$$\text{padding-left: 80px; } y_i = r$$

endif

$$\text{padding-left: 40px; } \psi \leftarrow \psi + \eta (y_i - Q_{\psi}(s_i, a_i)) \nabla_{\psi} Q_{\psi}(s_i, a_i)$$

end of do

When the agent is not in the learning phase, i.e. when it must use the Q-value to choose an action during its operational task, each q-value can be read directly from the q-function estimate Q_{ψ} , that is just calling $Q_{\psi}(s, a)$. The target policy in q-learning is the greedy policy, that mean that in a state s , the q-value for all actions is read, and the action with greater q-value is the chosen one:

$$a_{chosen} = \operatorname{argmax}_a Q_{\psi}(s, a) \tag{2.7}$$

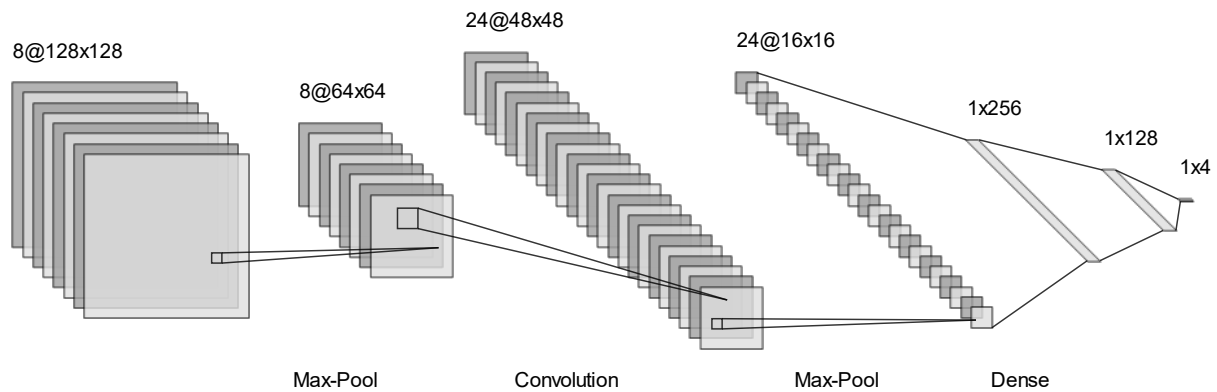


Figure 2.1

A hypothetical convolutional deep Q-Network modeled after a LeNet-like network [LeCun et al. 1998], with an additional output layer of 4 units, to be used for a problem where the agent can choose among 4 discrete actions. The input layer must be fed with a picture of the state (256x256 pixels, 3 channels RGB), then the network does its computations, and in the 4 output neurons we will have the 4 q-values for the 4 different actions in that state. Figure created with the online tool at <https://alexlenail.me/NN-SVG/>.

2.2 Minibatch Improvement

The algorithm above is a simple version and in the next part of the chapter we will see how to improve it. The first thing to notice is that it collects a transition and computes the respective change of Q-value estimate immediately. But to improve learning speed it is possible to collect a certain number of transitions, called “minibatch”, and calculate at once all the q-value changes for all the transitions in the minibatch by vectorized computations in a GPU or in a Tensor Processing Unit (note: in Machine Learning literature some authors use the word “batch” as a synonym for “minibatch”, some authors instead reserve the word “batch” to describe the whole set of samples available. Hence for clarity’s sake I used the word “minibatch”, to indicate a certain number of samples but not the whole set).

2.3 Saving Transitions into a Buffer

As wrote above, Q-Learning is an off-policy algorithm because it trains under a policy that includes exploration, but computes a Q-function estimate for a policy that only assumes exploitation. In fact, because of how the algorithm is designed, theoretically the behavior policy

could be even a policy very different from the target policy, we are not obliged to use an ϵ -greedy policy to generate samples. That means that all past trajectories may be retained and reused in future optimization iterations. Saving past trajectories in a buffer allows us to subsequently sample them randomly so to have uncorrelated transitions instead of transitions that are consecutive and hence correlated and close in state space (when transitions are correlated, learning can be unstable). Nonetheless there is a practical limit: if the q-value estimate is computed through a neural network, it has a finite approximation capacity. A neural network can approximate better when the inputs at inference time are of the same distribution of the training samples (this, very trivially, because the number of parameters is finite and the network cannot retain all the information that is presented during training, and it will approximate better for the kind of inputs that went through a greater number of optimization steps). So, it is better to train the q-network with input states that are on the same distribution that the q-network would experience if using the target policy (the tentatively optimal/greedy policy). To have similar distribution of states it is necessary to use a behavior policy that is similar to the target policy. That motivates the usage of an ϵ -greedy policy and motivates the practice of re-using the old samples (saved in a buffer) only for a certain number of steps and then discard them. A simple mechanism to achieve that is, when the buffer is full, to discard eldest samples to make room for new ones, namely a FIFO (first-in-first-out) buffer. Nonetheless, the bigger is your neural network, and hence the network approximation capacity, the longer the old samples may be retained and reused by the algorithm.

2.4 Bootstrapping and Convergence

Q-learning uses “Bootstrapping” in Q-value updates: it uses Q-value estimate of next state to update the Q-value estimate of current state, instead of using only rewards or complete returns (methods that only use complete returns are instead called “Monte Carlo methods”).

Since Q-Learning is both off-policy and bootstrapping, it is not guaranteed to converge when using function approximators such as neural networks for the Q-function [Szepesvári 2010]. In fact, the three conditions “function approximation”, “bootstrapping”, “off-policy training” are named “The Deadly Triad” when they occur together in value based methods, see [Sutton and Barto 2018, Ch.11]. Despite that, it is believed that Q-Learning may converge if the behavior policy is sufficiently close to the target policy, like in ϵ -greedy policies with small ϵ . In many experiments Q-Learning showed to work well, such as in [Mnih et al. 2013] where a variant of

Q-Learning with a convolutional neural network as a function approximator was trained to play Atari video games.

2.5 Increasing Gradient Stability: Target Network

When Q_ψ is a function approximator such as a neural network, the fact that it is used both to compute the prediction (and hence changes at each gradient descent iteration) and the target, makes the bootstrapped value y a “moving target” (literally!). This slows down learning because the gradient descent trajectory is wandering too much. To stabilize the trajectory of the gradient descent it is better to do a minibatch update with a big number of samples. To further stabilize it, it is useful to observe the fact that it is an off-policy algorithm: each tuple of $\langle s_t, a_t, s_{t+1}, r_t \rangle$ can be saved to a “replay buffer” and used in future (as we wrote in section 2.3): sampling randomly from the replay buffer instead of using the recently experienced trajectory will avoid to having correlated samples that would bias the gradient. In addition to this, to decrease even more the “moving target” effect it would be better to run some cycles of minibatch updates using a “target network” which is a copy of the Q-network that is not updated at every gradient descent step, only after a certain number of steps, and hence it is more stable.

Algorithm 2.2 Deep Q-Learning with Replay Buffer and Target Network

Require: an empty replay buffer B

Require: Step size η

Require: Initialize parameters ψ of network Q_ψ with small random values

do:

 copy target network parameters $\psi' \leftarrow \psi$

 do N times:

 using some behavior policy collect transitions $\{(s_i, a_i, s_{i+1}, r_i, is_terminal)\}$,

 add transition to buffer B

 do K times:

 sample a minibatch of M transitions $\langle s_i, a_i, s_{i+1}, r_i \rangle$ from B

 for each i of M transitions do (all can be done at once if vectorized):

 if not $is_terminal$ then:

$$y_i = r + \gamma \max_{a_{i+1}} Q_{\psi'}(s_{i+1}, a_{i+1})$$

 else:

$$y_i = r$$

 endif

$$\psi \leftarrow \psi + \eta (y_i - Q_\psi(s_i, a_i)) \nabla_\psi Q_\psi(s_i, a_i)$$

 end of for each

 end of do

 end of do

end of do

2.6 Polyak Averaging

The stability of the algorithm with target network can be improved if, instead of copying completely the parameters from original q-value network to target network after some round of updating, the target network gets a small update towards the original q-network every time that the original q-network is updated. The target network updates are such to make the target network a little more similar to the original network, but not too fast. For instance an update of the target network may be:

$$\psi' \leftarrow \varsigma \psi' + (1 - \varsigma) \psi$$

with $0 \leq \varsigma \leq 1$ (a good value may be $\varsigma = 0.99$)

(2.8)

This is named Polyak Averaging.

Algorithm 2.3 Deep Q-Learning with Replay Buffer, Target Network, Polyak Averaging

Require: an empty replay buffer B

Require: Step size η

Require: Polyak Step size ς

Require: Initialize parameters ψ of network Q_ψ with small random values

copy target network parameters $\psi' \leftarrow \psi$

do N times:

using some behavior policy collect transitions $\{ \langle s_i, a_i, s_{i+1}, r_i, is_terminal \rangle \}$,
add transition to buffer B

do K times:

sample a minibatch of M transitions $\langle s_i, a_i, s_{i+1}, r_i \rangle$ from B

for each i of M transitions do (all can be done at once if vectorized):

if not $is_terminal$ then:

$$y_i = r + \gamma \max_{a_{i+1}} Q_{\psi'}(s_{i+1}, a_{i+1})$$

else:

$$y_i = r$$

endif

$$\psi \leftarrow \psi + \eta (y_i - Q_\psi(s_i, a_i)) \nabla_\psi Q_\psi(s_i, a_i)$$

$$\psi' \leftarrow \varsigma \psi' + (1 - \varsigma) \psi$$

end of for each

end of do

end of do

2.7 Action-Value Overestimation

Another problem of Q-Learning is that it learns a Q-value that instead of being the maximum of the expectation of the action-value is biased towards the estimate of the maximum of the action-value and hence overestimates the action-value (the maximum of an expectation is different from the expectation of the maximum) [Thrun and Schwartz 1993].

To see how that happens, let us create a toy scenario. Imagine we are in a certain state s_0 in which you have five actions available: a_1, a_2, a_3, a_4, a_5 . Taking a_1 will give a very low reward and lead to a state s_1 with $\max_{a'} Q(s_1, a')$ equal to some constant K . Taking a_2 will give sometimes a very high reward (50% of the times) and sometimes a medium reward (50% of the times) and lead to a state s_2 with a $\max_{a'} Q(s_2, a')$ equal to the same K as s_1 . Actions a_3, a_4 and a_5 will behave in the same way as a_2 , except they will lead respectively to states s_3, s_4 and s_5 , which have the same $\max_{a'} Q(s', a')$ as s_1 and s_2 (that is K). Let us put that in numbers:

Action in s_0	Reward	Leads to State	$\max_{a'} Q(s', a')$
a_1	0 (Probability 1.0)	s_1	$\max_{a'} Q(s_1, a') = K$
a_2	30 (Probability 0.5)	s_2	$\max_{a'} Q(s_2, a') = K$
	10 (Probability 0.5)		
a_3	30 (Probability 0.5)	s_3	$\max_{a'} Q(s_3, a') = K$
	10 (Probability 0.5)		
a_4	30 (Probability 0.5)	s_4	$\max_{a'} Q(s_4, a') = K$
	10 (Probability 0.5)		
a_5	30 (Probability 0.5)	s_5	$\max_{a'} Q(s_5, a') = K$
	10 (Probability 0.5)		

Table 2.1

So we can compute the expected values for each action:

$$\begin{aligned}
 EQ(s_0, a_1) &= 0 * 1.0 + \gamma \max_{a'} Q(s_1, a') = 0 + \gamma K = \gamma K \\
 EQ(s_0, a_2) &= 30 * 0.5 + 10 * 0.5 + \gamma \max_{a'} Q(s_2, a') = 20 + \gamma K \\
 EQ(s_0, a_3) &= 30 * 0.5 + 10 * 0.5 + \gamma \max_{a'} Q(s_3, a') = 20 + \gamma K \\
 EQ(s_0, a_4) &= 30 * 0.5 + 10 * 0.5 + \gamma \max_{a'} Q(s_4, a') = 20 + \gamma K \\
 EQ(s_0, a_5) &= 30 * 0.5 + 10 * 0.5 + \gamma \max_{a'} Q(s_5, a') = 20 + \gamma K
 \end{aligned}$$

Hence, the *maximum of expectation of Q with respect to the actions* for the state s_0 is:

$$\max_{a'} EQ(s_0, a') = 20 + \gamma K$$

But if instead we run the Q-Learning algorithm we will find a different value. For instance imagine we do 15 runs starting from state s_0 , we choose each action three times, and because of random chances we get the following rewards:

Run	Chosen action	Reward r
1	a_1	0
2	a_1	0
3	a_1	0
4	a_2	10
5	a_2	30
6	a_2	30
7	a_3	10
8	a_3	10
9	a_3	10
10	a_4	30
11	a_4	30
12	a_4	30
13	a_5	10
14	a_5	30
15	a_5	10

Table 2.2

Given these outcomes, we can see that randomly, due to the probabilistic nature of rewards, in certain cases the target value $y = r + \gamma \max_{a'} Q_\psi(s', a')$ (eq. 2.3) is greater than the expected value $EQ(s_0, a')$ and sometimes it is smaller. For instance for action a_5 the target value is twice $y = 10 + \gamma K$ and only once $y = 30 + \gamma K$. The Q-learning algorithm will then tend to compute a value for $Q_\psi(s_0, a_5)$ that is smaller than $EQ(s_0, a_5)$. For action a_4 instead,

all three runs obtained the maximum reward, hence the target value is $y = 30 + \gamma K$ all the three times, that is greater than the expected value of $20 + \gamma K$, so the Q-learning algorithm will tend to compute a value for $Q_\psi(s_0, a_4)$ that is bigger than $EQ(s_0, a_4)$. At first glance this may seem ok, because we expect that due to random nature of runs, some values will have an estimate greater than their true expectation, and some a lesser one. But the problem is that the Q-Value is the maximum among all the estimates, so if the Q-value estimate for the optimal action (the one that has actually the greater $EQ(s_t, a')$) is greater than its expected value, it will drive the value of s_t up. But if instead the Q-value estimate for the optimal action is lesser than its expected value, it is not certain that it will drive the value of s_t down, because another action (with expected value lower or equal), due to randomness, may have had an estimate greater than the estimate of the optimal action, and that will still drive the value of s_t up. The value of a state (computed as the maximum among its Q-values) is used to bootstrap the Q-values of the states that reach it, propagating the value overestimation to other states. In our example we know that:

$$\max_a EQ(s_0, a) = EQ(s_0, a_2) = EQ(s_0, a_3) = EQ(s_0, a_4) = EQ(s_0, a_5)$$

We also know that, from the samples in our example:

$$\begin{aligned} Q_\psi(s_0, a_2) &> EQ(s_0, a_2) \\ Q_\psi(s_0, a_3) &< EQ(s_0, a_3) \\ Q_\psi(s_0, a_4) &> EQ(s_0, a_4) \\ Q_\psi(s_0, a_5) &< EQ(s_0, a_5) \\ Q_\psi(s_0, a_4) &> Q_\psi(s_0, a_2) \end{aligned}$$

And we know that $V(s_0)$ is estimated by:

$$\max_a Q_\psi(s_0, a) = Q_\psi(s_0, a_4) > \max_a EQ(s_0, a)$$

Hence even if for optimal actions we had an equal number of runs with target greater and with target lesser than expected, the estimate of s_0 value $\max_a Q_\psi(s_0, a)$ is greater than the real value, and through bootstrapping will propagate the overestimation to any other Q-values of states that reach s_0 .

So in this toy example we see how random fluctuations may give returns that are either lesser or greater than true average returns but while all times in which the return of optimal action is greater than its expectation it drives the $\max_a Q_\psi(s, a)$ up, not all times in which the return of

optimal action is lesser than its expectation it drives the $\max_a Q_\psi(s, a)$ down. This creates an overestimation bias that propagates to other Q-values through bootstrapping.

2.8 Double Q-Learning

A formal mathematical proof of the overestimation, as well as the proposal of an algorithm named “*Double Q-Learning*” that mitigates the problem can be found in [van Hasselt 2010], the application of that algorithm to deep neural networks can be found in [van Hasselt et al. 2016].

In Double Q-Learning the algorithm uses two different action-value function estimators, e.g. two neural networks that we may call A and B. When computing the target value, one neural network (let us say A) is used to check which action has the maximum value, while the other network (let us say B) is used to actually obtain the Q-value of that action, to compute the target value. Then the target value is used to update the Q-value of the former network (A, in this case).

At each iteration the estimators may swap roles: the one that previously was used to evaluate which action has the maximum value may now used to obtain the Q-value of the optimal action, while the estimator who was previously used to get the Q-value of the optimal action may now used to find which action has the maximum value, and be subject to Q-value update. The roles of the two networks A and B may be decided randomly or in a fixed way after a certain number of iterations. For clarity's sake in the pseudocode algorithm we use the names “network Q_{ψ_1} ” and “network Q_{ψ_2} ” instead of A and B.

Algorithm 2.5 Deep Double Q-Learning with Replay Buffer

Require: an empty replay buffer B

Require: Step size η

Require: Initialize parameters ψ_1 of network Q_{ψ_1} with small random values

Require: Initialize parameters ψ_2 of network Q_{ψ_2} with small random values

do:

 do N times:

 using some behavior policy collect transitions $\{ \langle s_i, a_i, s_{i+1}, r_i, is_terminal \rangle \}$,
 add transition to buffer to B

 do K times:

 sample a minibatch of M transitions $\langle s_i, a_i, s_{i+1}, r_i \rangle$ from B

 choose (e.g. randomly) either UPDATE(1) or UPDATE(2)

 for each i of M transitions do (all can be done at once if vectorized):

 if UPDATE(1) then:

 if not $is_terminal$ then:

$$a_{i+1}^* = \arg \max_{a_{i+1}} Q_{\psi_1}(s_{i+1}, a_{i+1})$$

$$y_i = r + \gamma Q_{\psi_2}(s_{i+1}, a_{i+1}^*)$$

 else :

$$y_i = r$$

 endif

$$\psi_1 \leftarrow \psi_1 + \eta (y_i - Q_{\psi_1}(s_i, a_i)) \nabla_{\psi_1} Q_{\psi_1}(s_i, a_i)$$

 else if UPDATE(2) then:

 if not $is_terminal$ then:

$$a_{i+1}^* = \arg \max_{a_{i+1}} Q_{\psi_2}(s_{i+1}, a_{i+1})$$

$$y_i = r + \gamma Q_{\psi_1}(s_{i+1}, a_{i+1}^*)$$

 else :

$$y_i = r$$

 endif

$$\psi_2 \leftarrow \psi_2 + \eta (y_i - Q_{\psi_2}(s_i, a_i)) \nabla_{\psi_2} Q_{\psi_2}(s_i, a_i)$$

 endif

 end of for each

end of do

end of do

end of do

The Double Q-Learning may be used together with the Target Network mechanism (the one where the network used to compute the q-value for the target value is a copy of the actual network, but it is updated only after the whole learning round, see section 2.5).

Algorithm 2.6 Deep Double Q-Learning with Replay Buffer & Target Network

Require: an empty replay buffer B

Require: Step size η

Require: Initialize parameters ψ_1 of network Q_{ψ_1} with small random values

Require: Initialize parameters ψ_2 of network Q_{ψ_2} with small random values

do:

 copy target network parameters $\psi_1' \leftarrow \psi_1$

 copy target network parameters $\psi_2' \leftarrow \psi_2$

 do N times:

 using some behavior policy collect transitions $\{ \langle s_i, a_i, s_{i+1}, r_i, is_terminal \rangle \}$,

 add transition to buffer to B

 do K times:

 sample a minibatch of M transitions $\langle s_i, a_i, s_{i+1}, r_i \rangle$ from B

 choose (e.g. randomly) either UPDATE(1) or UPDATE(2)

 for each i of M transitions do (all can be done at once if vectorized):

 if UPDATE(1) then:

 if not $is_terminal$ then:

$$a_{i+1}^* = \arg \max_{a_{i+1}} Q_{\psi_1'}(s_{i+1}, a_{i+1})$$

$$y_i = r + \gamma Q_{\psi_2'}(s_{i+1}, a_{i+1}^*)$$

 else :

$$y_i = r$$

 endif

$$\psi_1 \leftarrow \psi_1 + \eta (y_i - Q_{\psi_1}(s_i, a_i)) \nabla_{\psi_1} Q_{\psi_1}(s_i, a_i)$$

 else if UPDATE(2) then:

 if not $is_terminal$ then:

$$a_{i+1}^* = \arg \max_{a_{i+1}} Q_{\psi_2'}(s_{i+1}, a_{i+1})$$

$$y_i = r + \gamma Q_{\psi_1'}(s_{i+1}, a_{i+1}^*)$$

```

else :
     $y_i = r$ 
endif
 $\psi_2 \leftarrow \psi_2 + \eta (y_i - Q_{\psi_2}(s_i, a_i)) \nabla_{\psi_2} Q_{\psi_2}(s_i, a_i)$ 
endif
end of for each
end of do
end of do
end of do

```

2.9 Issues with Continuous Actions

It has to be noted that the Q-Learning algorithms presented above can be used only with a set of discrete actions. The problem with continuous actions is the computation of the maximum q-value among all possible actions: when we have a set of discrete actions it is possible to check the q-value of all actions and pick the maximum, but if the actions are continuous we would need an analytic way to compute the maximum of the q-function, and that is not possible if the q-function estimate is represented by a neural network. There have been many attempts to use workarounds to extend Q-Learning to continuous actions, but for a matter of space we will not describe them here: some of them are viable in certain scenarios but not effective in others, such as discretizing the action space which suffers of the curse of dimensionality. A better way to do Reinforcement Learning with continuous actions is to use Policy Based methods, that are detailed in the following chapter. In particular, one of those methods named Deep Deterministic Policy Gradient (Ch. 9), is considered to be applying to continuous actions the same basic strategy that Q-Learning applies to discrete actions: the ideas of having a deterministic target policy and of improving it greedily.

2.10 Sarsa

If you do not need to learn the optimal policy but you only want to evaluate the q-values for a fixed policy π , you can use Sarsa: an algorithm that is similar to Q-Learning where instead of bootstrapping with the maximum among the $Q_\psi(s_{i+1}, a_{i+1})$ for all possible $a_{i+1} \in A$, you just sample a_{i+1} following the policy $\pi(\cdot | s_{t+1})$. The trajectory of actual steps followed by the agent can be decided by the fixed policy π (on-policy) but can also follow a different behavior policy (off-policy), but a behavior policy too different from target policy may increase divergence risk. Given a sample transition $\langle s, a, s', r \rangle$ obtained by any policy (same of π , or different from π), the update rule is:

$$\begin{aligned} a' &\sim \pi(\cdot | s') \\ y &= r + \gamma Q_\psi(s', a') \end{aligned} \tag{2.9}$$

The update rule for the tabular version (TD Learning) is:

$$Q_\psi(s, a) \leftarrow Q_\psi(s, a) + \eta (y - Q_\psi(s, a)) \tag{2.10}$$

The update rule for the function approximator version is:

$$\begin{aligned} L &= (y - Q_\psi(s, a))^2 \\ \psi &\leftarrow \psi + \eta (y - Q_\psi(s, a)) \nabla_\psi Q_\psi(s, a) \end{aligned} \tag{2.11}$$

Since to estimate q-values we use a tuple $\langle s, a, r, s', a' \rangle$, we understand where the name of the algorithm comes from.

Algorithm 2.5 Deep Sarsa – Q-Value Estimation Only (simple version)

Require: Step size η

Require: Initialize parameters ψ of network Q_ψ with small random values

Do until (supposed) convergence:

 using some behavior policy collect a transition $\langle s_i, a_i, s_{i+1}, r_i, is_terminal \rangle$

 sample $a_{i+1} \sim \pi(\cdot | s_{i+1})$

 if not $is_terminal$ then:

$$y_i = r_i + \gamma Q_\psi(s_{i+1}, a_{i+1})$$

 else :

$$y_i = r_i$$

 endif

$$\psi \leftarrow \psi + \eta (y_i - Q_\psi(s_i, a_i)) \nabla_\psi Q_\psi(s_i, a_i)$$

end of do

If the actions are discrete, it is also possible to improve the policy π : for each state the improved policy can select the action that gives the greater q-value, as we were doing with q-value in eq. 2.7 : $a_{chosen} = \operatorname{argmax}_a Q_\psi(s, a)$

This way of improving the policy anyway is much slower than q-learning because it bootstraps the q-value using next-state q-value of the action chosen by the fixed policy and not the maximum q-value, so the information of maximum q-value is slower to propagate.

It is not possible (or better, it is not easy nor recommended) to improve policies in this way when actions are continuous, as we already discussed about q-learning that would be possible only with workarounds and there are better alternatives.

If we have discrete actions, an improved version of Sarsa is possible. “Expected Sarsa” bootstraps with the expectation of $Q_\psi(s_{i+1}, a_{i+1})$: instead of just sampling a_{i+1} from the target policy and compute $Q_\psi(s_{i+1}, a_{i+1})$.

$$y = r + \gamma \sum_a \pi(a|s') Q_\psi(s', a') \quad (2.12)$$

This implies a lower variance and faster convergence.

When the target policy is the greedy policy and the exploration policy is e-greedy, Expected Sarsa (with policy improvement) coincides with Q-Learning.

Algorithm 2.6 Deep Expected Sarsa – Q-Value Estimation Only (simple version)

Require: Step size η

Require: Initialize parameters ψ of network Q_ψ with small random values

Do until (supposed) convergence:

 using some behavior policy collect a transition $\langle s_i, a_i, s_{i+1}, r_i, is_terminal \rangle$

 if not $is_terminal$ then:

$$y_i = r_i + \gamma \sum_a \pi(a|s_{i+1}) Q_\psi(s_{i+1}, a_{i+1})$$

 else :

$$y_i = r_i$$

 endif

$$\psi \leftarrow \psi + \eta (y_i - Q_\psi(s_i, a_i)) \nabla_\psi Q_\psi(s_i, a_i)$$

end of do

If we do not need to improve/learn a policy, but we only need to estimate its q-values, we can use Sarsa estimation also with continuous actions (but we must use a function approximator such as a neural network, since we cannot use tabular methods with continuous actions). This can be useful in other algorithms, such as in Deep Deterministic Policy Gradient (Ch. 9).

3. Policy Gradient

3.1 Policy Based Methods

Policy based methods, also called “Policy Gradient methods” and “Policy Optimization methods” (those two expressions have diverse meaning but often are used interchangeably), are characterized by having a parametrized policy function that, given the state/observation as input, returns a probability for each action, without explicitly assigning expected return values. Differently from the value based methods, policy based methods may or may not use value or action-value functions.

In detail, the policy function $\pi(\cdot | s_t)$ is a function that is explicitly parametrized by θ (for instance θ may be the weights of a neural network), such that if at time t the state is s_t , it will return the probability distribution over which action to take a_t :

$$a_t \sim \pi_{\theta}(\cdot | s_t) \tag{3.1}$$

That is like equation 1.1 with the addition of parameters θ .

One way in which a neural network can output a probability distribution over a set of discrete actions is to have a last layer with as many neurons as the number of different actions, and on top of that operate a Softmax. If instead the actions are continuous then the last layer of the neural network will have as many neurons as continuous parameters in the action, and each neuron will be considered the mean of the distribution of that action parameter (the distribution type may be arbitrary, for instance a normal usually works well). When actions are continuous the covariance matrix of the distribution is usually considered diagonal (so to sample each action parameter independently from the others), and it may be of a fixed value established a priori, or it may be a value outputted by other neurons of the network. If the variance of each continuous action is outputted by neurons, it is better to output it in form of log-standard deviation, so there is no need to do any workaround to ensure that it is non-negative. When deciding which action to take, the actual action to be taken will be sampled from the distribution outputted by the neural network as we just described.

The problem of Reinforcement Learning is still the same, we want to maximize the expected return $J(\pi)$, (detailed in equation 1.5), but this time we would like to modify directly the policy parameters θ . One way of doing that could be to do gradient ascent using the gradient of the expected return $J(\pi)$ with respect to the parameters θ :

$$\theta_{k+1} \leftarrow \theta_k + \eta \nabla_{\theta} J(\pi_{\theta_k}) \quad (3.2)$$

With η learning rate.

So it is necessary to compute the gradient of the expected return $J(\pi_{\theta_k})$ with respect to θ .

Apparently that could seem problematic because the expected return $J(\pi_{\theta_k})$ is an expectation with respect to the distribution of the trajectories, that depends both on the policy and on the transition function, but we assume to not know the transition function so we do not know the distribution of trajectories (or the distribution of states). Hence, we do not even know how a change of the policy may change the distribution of trajectories. But we need to estimate how the expected return changes as a consequence of changes of policy (due to changes of θ), and this seems to be problematic because, as just said, the effect of the policy change on the trajectories distribution is unknown. Fortunately, as we will see in next chapter, there is a way to derive the gradient of the expected return $J(\pi_{\theta_k})$ that does not involve the effect of the change of policy on states/trajectories distribution, so this is not really a problem.

3.2 Policy Gradient Basics

A basic algorithm that improves the policy through gradient ascent on the expected return $J(\pi_{\theta})$ with respect to the policy parameters θ is the one called “**Vanilla Policy Gradient**” [Achiam & OpenAi 2020A] (or just “Policy Gradient”), that I describe below. It is very similar to the earlier “REINFORCE” [Williams 1992], with the difference that REINFORCE updates the policy parameters at each step while Vanilla Policy Gradient updates parameters using minibatches of steps. It is an on-policy algorithm. Some use the term “Vanilla Policy Gradient” as a family of algorithms, with REINFORCE being one of them (see [Peters and Schaal 2008]), where the policy gradient is not manipulated (as it happens instead in Natural Policy Gradient or in Proximal Policy Optimization, which I describe in later chapters).

Following the method in [Achiam & OpenAi 2020A] it is useful to recall from calculus that the derivative of $\log(x)$ with respect to x is $1/x$. Then, because of chain rule, the derivative of $\log(f(x))$ with respect to x is the derivative of $f(x)$ divided by $f(x)$, that is $\log(f(x))' = f'(x)/f(x)$. Applying that using the trajectory distribution $P(\tau|\theta)$ instead of $f(x)$, and rearranging, we obtain the so-called “log derivative trick”:

$$\nabla_{\theta} P(\tau|\theta) = P(\tau|\theta) \nabla_{\theta} \log P(\tau|\theta) \quad (3.3)$$

Now, taking the definition of the probability of a trajectory τ made of T steps from equation 1.4, we compute the log of it:

$$\log P(\tau|\theta) = \log \rho_0(s_0) + \sum_{t=0}^{T-1} \log P(s_{t+1}|s_t, a_t) + \log \pi_{\theta}(a_t|s_t) \quad (3.4)$$

At this point we want to compute $\nabla_{\theta} \log P(\tau|\pi)$. We note that the terms $P(s_{t+1}|s_t, a_t)$ and $\rho_0(s_0)$ are environmental and that do not depend on θ , so their gradient is zero. Hence:

$$\nabla_{\theta} \log P(\tau|\theta) = \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \quad (3.5)$$

Now, computing the gradient of equation 1.5 we obtain:

$$\begin{aligned} \nabla_{\theta} J(\pi_{\theta}) &= \nabla_{\theta} E_{\tau \sim \pi_{\theta}} [G(\tau)] \\ &= \nabla_{\theta} \int_{\tau} P(\tau|\theta) G(\tau) \\ &= \int_{\tau} \nabla_{\theta} P(\tau|\theta) G(\tau) \end{aligned} \quad (3.6)$$

And applying equation 3.3 (log derivative trick):

$$\begin{aligned}
&= \int_{\tau} P(\tau|\theta) \nabla_{\theta} \log P(\tau|\theta) G(\tau) \\
&= E_{\tau \sim \pi_{\theta}} [\nabla_{\theta} \log P(\tau|\theta) G(\tau)]
\end{aligned}
\tag{3.7}$$

Applying equation 3.5 :

$$\nabla_{\theta} J(\pi_{\theta}) = E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau) \right]$$

■

(3.8)

The same formula can be obtained also in a different way through the *Policy Gradient Theorem* [Sutton et al. 1999], [Sutton and Barto 2018, Ch.13] and then *REINFORCE* [Williams 1992] which I will not detail here.

What is remarkable here is that even though the expected return depends on the state distribution (which in turn depends on the policy and hence on the parameters θ), the gradient of the expected return does not depend on the gradient of the state distribution (which we do not know).

We obtained the gradient of the expected return in form of an expectation. This means that if we can sample what is contained inside the expectation, we can use it to estimate the gradient. It turns out that we can do that: the expectation is with respect to the trajectory τ , whose probability distribution is determined by the environment and by our policy π_{θ} , so by making the agent follow the policy we can collect the returns $G(\tau)$ and compute the estimate of the expectation. Since the agent must follow the policy π_{θ} the algorithm is on-policy.

If we run N different episodes or trajectories $\{\tau_0, \tau_1, \dots, \tau_{N-1}\}$, with respective trajectory lengths $\{T_0, T_1, \dots, T_{N-1}\}$ we can do minibatch gradient ascent. If we call the actions and states of trajectory i at time t respectively with $a_{i,t}$ and $s_{i,t}$ the following is mean gradient:

$$\hat{g} = \frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) G(\tau_i)$$

(3.9)

Then we can update the weights of the policy neural network:

$$\theta' \leftarrow \theta + \eta \hat{g} \quad (3.10)$$

Algorithm 3.1 **Policy Gradient**

Require: Policy network step size η

Require: Initialize parameters θ of network π_θ with small random values

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using policy $\pi_k(\theta_k)$

 compute total rewards $G(\tau_i)$ for each τ_i

 estimate policy gradient as:

$$\hat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau)$$

 update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \hat{g}_k$$

end of for

3.3 Rewards-To-Go

A notable fact is that the gradient of the log probability of each action a_t is multiplied by the complete return $G(\tau)$. The complete return of the trajectory includes the rewards that have been received before that action, and for which the action has not any responsibility. From a logical standpoint we would expect that the gradient of the log probability of a certain action is multiplied only by the sum of the rewards obtained after that action, instead also the rewards obtained before that action are used. As we will see, the part of gradient of the expected return that involves the multiplication for the rewards obtained before the action is equal to zero in expectation, so we can use only the rewards obtained after the action. To prove it, I must digress a little and introduce the Expected Grad-Log-Prob (EGLP) lemma, from [Achiam & OpenAI 2020A].

EGLP Lemma: suppose that f_{θ} is a parametrized distribution over a random variable x , then:

$$E_{x \sim f_\theta} [\nabla_\theta \log f_\theta(x)] = 0 \quad (3.11)$$

Proof (noting that the integral of a probability distribution is always = 1 by definition):

$$\nabla_\theta \int_x f_\theta(x) = \nabla_\theta 1 = 0 \quad (3.12)$$

Now, applying the “log derivative trick” equation 3.3 :

$$\begin{aligned} 0 &= \nabla_\theta \int_x f_\theta(x) \\ &= \int_x \nabla_\theta f_\theta(x) \\ &= \int_x f_\theta(x) \nabla_\theta \log f_\theta(x) \\ \therefore 0 &= E_{x \sim f_\theta} [\nabla_\theta \log f_\theta(x)] \end{aligned} \quad \blacksquare \quad (3.13)$$

(Optional: this proof is analogous to the proof that the expected value of statistical score of a distribution conditioned on the true value of its parameters is equal to zero. In fact since the statistical score is the gradient of the log-likelihood, when θ are the true values of the parameters of the distribution we could substitute the likelihood function $\mathcal{L}(\theta; X)$ in place of $f_\theta(x)$ in equation 3.11 , that means that the probability of X would be $\mathcal{L}(\theta; X)$, and we will obtain the same result, that in this case would mean that expected value of statistical score conditioned on true θ is zero).

We can decompose $G(\tau)$ in 2 sums, one of the rewards before the action and one after the action happened at time t (this proof is elaborated from the one by [Soemers 2019], there exists also another proof by [Achiam & OpenAi 2020B]).

We use $G(\tau_{0:t-1})$ for the past rewards:

$$G(\tau_{0:t-1}) = \sum_{k=0}^{t-1} \gamma^k r_k \quad (3.14)$$

And we call $G(\tau_t)$ “reward-to-go”.

$$G(\tau_t) = \sum_{k=t}^{T-1} \gamma^k r_k \quad (3.15)$$

$$G(\tau) = G(\tau_{0:t-1}) + G(\tau_t) \quad (3.16)$$

We can then decompose equation 3.8 in the same way:

$$\begin{aligned} \nabla_{\theta} J(\pi_{\theta}) &= E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau_{0:t-1}) + \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau_t) \right] \\ &= E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau_{0:t-1}) \right] + E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau_t) \right] \end{aligned} \quad (3.17)$$

Now, this can be rewritten as:

$$part1 = E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau_{0:t-1}) \right] \quad (3.18)$$

$$part2 = E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau_t) \right] \quad (3.19)$$

$$\nabla_{\theta} J(\pi_{\theta}) = part1 + part2$$

(3.20)

Now, because of EGLP lemma:

$$E_{a_t \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | s_t)] = 0 \quad (3.21)$$

Now, from 3.18 we have:

$$\begin{aligned} part1 &= E_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) G(\tau_{0:t-1}) \right] \\ &= \sum_{t=0}^{T-1} E_{\tau \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | s_t) G(\tau_{0:t-1})] \end{aligned} \quad (3.22)$$

We know that $\nabla_\theta \log \pi_\theta(a_t | s_t)$ and $G(\tau_{0:t-1})$ are independent because the action depends only on the state at time t by definition, and $G(\tau_{0:t-1})$ depends only on states and actions before time t . So, we can separate the expectation of the multiplication into a multiplication of expectations:

$$= \sum_{t=0}^{T-1} E_{\tau \sim \pi_\theta} [G(\tau_{0:t-1})] E_{\tau \sim \pi_\theta} [\nabla_\theta \log \pi_\theta(a_t | s_t)] \quad (3.23)$$

Now we plug equation 3.21 in 3.23:

$$\begin{aligned} &= \sum_{t=0}^{T-1} E_{\tau \sim \pi_\theta} [G(\tau_{0:t-1})] \cdot 0 \\ &\therefore part1 = 0 \end{aligned} \quad (3.24)$$

So we obtain:

$$\nabla_\theta J(\pi_\theta) = E_{\tau \sim \pi_\theta} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) G(\tau_t) \right]$$

(3.25)

$$\nabla_{\theta} J(\pi_{\theta}) = E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \sum_{l=t}^{T-1} \gamma^l r_l \right] \quad (3.26)$$

This permits us to use, for each action, only the rewards obtained after it (the rewards-to-go) to calculate the gradient of the expected return. Using the complete returns would not be wrong, because it is equal in expectation, but it would introduce much more variance in the gradient and that would slow down learning. Using $\gamma = 1$ would assure that each action has the same importance, it can be used in finite-horizon (episodic learning).

Algorithm 3.2 Policy Gradient with Rewards-To-Go

Require: Policy network step size η

Require: Initialize parameters θ of network π_{θ} with small random values

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using policy $\pi_k(\theta_k)$

 compute rewards-to-go $G(\tau_{i,t})$, for each t of each τ_i

 estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) G(\tau_t)$$

 update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g}_k$$

end of for

3.4 Bootstrapping the Value

In infinite time horizon it is not possible to collect the whole return (the trajectory never ends), so in that case it is necessary to bootstrap $G(\tau_t)$ using only the reward for the first step (or for some steps) and then summing a discounted estimate of the state value.

For instance, in infinite-horizon, a k -step bootstrap, using a function approximator $V_\psi(s)$, parametrized by ψ , to estimate the value function, would be:

$$\widehat{G}_k(\tau_t) = \sum_{l=t}^{t+k-1} \gamma^l r_l + \gamma^{t+k} V_\psi(s_{t+k}) \quad (3.27)$$

$\widehat{G}_k(\tau_t)$ is equal in expectation to $G(\tau_t)$, so it is theoretically correct to use it, but since in any algorithm we are not using the real state value but an approximation, it will introduce bias (but it will allow to work with infinite-horizon cases).

In infinite horizon case moreover, since we don't use the whole infinite trajectory at once but only n -step of it a time, we have to consider the infinite trajectory as equivalent to an infinite set of n -step episodes. So, every n -step we will reset t to 0, that is also necessary to avoid that the discount makes progressively less relevant the subsequent rewards. Or, if we don't want to reset t to 0, just change the equation a little, subtracting t to the exponent:

$$\widehat{\widehat{G}}_k(\tau_t) = \sum_{l=t}^{t+k-1} \gamma^{l-t} r_l + \gamma^k V_\psi(s_{t+k}) \quad (3.28)$$

In this way the finite and infinite horizon are included in the same framework.

3.5 Beyond Rewards-To-Go

Also $Q^{\pi_\theta}(s_t, a_t)$ is theoretically correct to be used instead of $G(\tau_t)$, because since inside the expectation the actions are distributed following the policy π_θ , it implies that $Q^{\pi_\theta}(s_t, a_t)$ is equal in expectation to $G(\tau_t)$ (another formal proof by OpenAI may be found in [Achiam & OpenAI 2020C]).

In fact, it is theoretically correct also to use, instead of $G(\tau_t)$, any other expression that is equal in expectation. Every expression that starts with $G(\tau_t)$ or $\widehat{G}_k(\tau_t)$ and then adds any other expression that does not depend on the current action (but may depend on current state, since the current action and any function depending only on the current state are conditionally

independent given the current state), is equal in expectation. That is because of EGLP lemma, and we already used it to show that $G(\tau_{0:t-1})$ may be or may be not added to $G(\tau_t)$ without changing the expectation (see how we obtained equation 3.25).

So, more generally:

$$\nabla_{\theta} J(\pi_{\theta}) = E_{\tau \sim \pi_{\theta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \Phi_t \right] \quad (3.29)$$

With Φ_t that may be one of:

$$\begin{aligned} & G(\tau_t) \\ & \widehat{G}_k(\tau_t) \\ & \widehat{\widehat{G}}_k(\tau_t) \\ & Q^{\pi_{\theta}}(s_t, a_t) \\ & G(\tau_t) + b(s_t) \\ & \widehat{G}_k(\tau_t) + b(s_t) \\ & \widehat{\widehat{G}}_k(\tau_t) + b(s_t) \\ & Q^{\pi_{\theta}}(s_t, a_t) + b(s_t) \end{aligned} \quad (3.30)$$

Where $b(s_t)$, usually called “baseline”, is an additional expression that may depend on state and may be negative. For example, it may be $V^{\pi_{\theta}}(s_t)$ (in practice we would use its estimate $V_{\psi}(s_t)$).

So, many different Φ_t are possible. For instance if we start from $Q^{\pi_{\theta}}(s_t, a_t)$ and we add $b(s_t) = -V^{\pi_{\theta}}(s_t)$, we end up with the Advantage function $\Phi_t = A^{\pi_{\theta}}(s_t, a_t) = Q^{\pi_{\theta}}(s_t, a_t) - V^{\pi_{\theta}}(s_t)$, see equation 1.13.

3.6 Advantage Actor-Critic (A2C)

Using the advantage function as Φ_t makes sense because the policy gradient update is intended to increase the probability of actions that perform better than the policy and to

decrease the probability of actions that perform worse than the policy, if we use the Advantage function as Φ_t we are sure to have positive updates when the action is better than current policy choice, and negative updates when it is worse. It will have the effect of reducing variance in the gradient updates and hence speed up learning.

To implement an advantage function it is necessary to have a neural network $V_\psi(s)$, parametrized by ψ , that estimates the value function. That neural network will be updated with a minibatch gradient descent, similar to the Q-network gradient descent of equations 2.3 and 2.4, where the target value y may be computed by $G(\tau_t)$, the Monte Carlo rewards-to-go of state t (this would have a great variance) or may be computed bootstrapping the value using $V_\psi(s')$ (in that case it will not be full gradient descent but semi-gradient, since the value of the next state will be considered as a constant and not as a function of ψ).

The value network update will be the following:

$$\begin{aligned}
y &= r + \gamma V_\psi(s') \text{ if bootstrapping} \\
&\text{or} \\
y &= G(\tau_t) = \sum_{k=t}^{T-1} \gamma^k r_k \text{ if using Monte Carlo} \\
L &= (y - V_\psi(s))^2 \\
\psi &\leftarrow \psi + \eta (y - V_\psi(s)) \nabla_\psi V_\psi(s)
\end{aligned} \tag{3.31}$$

Policy gradient methods that use an estimate of the value function to bootstrap the estimate of the total return such as in equation 3.27 and 3.28 are called “**actor-critic**”, where actor is referred to the policy function and critic to the value function [Sutton and Barto 2018, Ch.13]. Some authors call “actor-critic” every policy method that uses an estimate of the value function as component of Φ_t , but Sutton and Barto specify that the term should be exclusive for those that use value function to bootstrap the estimate of the total return. That means when generally, for a fixed k and a certain baseline $b(s_t)$:

$$\Phi_t = \sum_{l=t}^{t+k-1} \gamma^{l-t} r_l + \gamma^k V_\psi(s_{t+k}) - b(s_t) \tag{3.32}$$

It is common to do that with $k = 1$ and with the estimate of the value function used as baseline, so that:

$$\Phi_t = r_t + \gamma V_\psi(s_{t+1}) - V_\psi(s_t) , \quad (3.33)$$

Which is equivalent to a bootstrapped advantage function.

Because of the usage of the value function to compute the advantage function, this algorithm is usually known as “**Advantage Actor-Critic**”, (A2C).

Algorithm 3.3 Policy Gradient with rewards-to-go and advantage function (A2C)

Require: Policy network step size η

Require: Value network step size ω

Require: Initialize parameters θ of network π_θ with small random values

Require: Initialize parameters ψ of network V_ψ with small random values

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using policy $\pi_k(\theta_k)$

 compute rewards-to-go $G(\tau_{i,t})$, for each t of each τ_i

 compute advantage estimates \hat{A}_t using current estimate of value function V_{ψ_k} :

$$\hat{A}_t = G(\tau_t) - V_{\psi_k}(s_t)$$

 estimate policy gradient as:

$$\hat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \hat{A}_t$$

 update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \hat{g}_k$$

For $n = 0, 1, 2, \dots, N - 1$ do a value function gradient descent iteration:

 estimate the value function gradients as:

$$\hat{h}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \frac{1}{T} \sum_{t=0}^{T-1} \nabla_\psi (V_\psi(s_t) - G(\tau_t))^2$$

 update the value function with a gradient descent step (or other method):

$$\psi_{k+1} \leftarrow \psi_k + \omega \hat{h}_k$$

 end of for

end of for

In [Mnih et al. 2016] the advantage function is used in a policy gradient actor-critic algorithm in a parallel learning setting: many agents, sharing the same policy and advantage function estimators, run simultaneously in separated environments (in separate cores or CPUs) and then asynchronously after a certain number of epochs share their accumulated gradients with a centralized parameters server, and retrieve updated models. This algorithm is named “A3C”, that stands for “Asynchronous Advantage Actor-Critic policy gradient”. Being asynchronous implies that sometimes the gradients sent to the central parameters server may be outdated because computed with a policy that is not updated with respect to the one on the central server, but in practice this algorithm seems to have good performance.

3.7 Benefits of Policy Gradients

A remarkable aspect of policy gradient methods is that compared to value based methods they have a more "natural" distribution over actions: the policy function outputs "by design" a distribution on actions, and during training the distribution smoothly changes towards one that improves the returns. In value based methods instead the estimate of value function or q-function only gives expected (action-)state returns, so we have to choose an artificial way to create a distribution on actions, such as with ϵ -greedy. Hence in value base methods when some action that previously seemed not optimal starts to seem better than all the other actions (that is when its q-value estimate becomes the greater among all actions), we have an abrupt change in the policy: the policy stops suggesting the previously-considered-best action (except for a small probability depending on ϵ) and starts suggesting almost greedily the new best action. In policy gradient methods instead the change in distribution is smooth because it is due to gradual changes through the policy gradient ascent, and this also implies an automatic gradual passage from exploration to exploitation.

This “natural” distribution over actions of policy methods, is also great in adversarial applications where an antagonist may understand that the agent is always deterministically taking a certain action which is deemed the best, and prepare a counter-action. If the agent has a value based policy, it will have a deterministic or a quasi-deterministic policy, so it may be predictable, while if the agent instead has a probability distribution over actions (a stochastic policy) the antagonist cannot be sure about which action is going to be taken by the agent (in Game Theory distributions over actions at state s are named “*mixed actions*”).

3.8 Improvements

It must be noted that in Policy Gradient algorithm the policy ascent direction is often not the most efficient because each parameter of θ may have a different magnitude of influence over the final policy distribution, but we are updating all parameters by the same step size. We will describe the problem and present an algorithm that computes a better direction in Ch. 6 with Natural Policy Gradient.

One thing to be careful of in Policy Gradient is the fact that when the policy step size η is too little, the learning will be slow. So, it should be big enough to make the learning progress quick, but not so big to make the new policy worse than the old. In fact the new policy may end up being quite close in parameters space to the old policy, but even small differences in parameters may result in big differences in policies. Hence a small step in parameter space may correspond to a big step in policy space, in a way that worsens the policy instead of improving it. An algorithm that aims at computing the right step size (together with a better direction) is Trust Region Policy Optimization, in Ch. 7 .

A further improvement for all Policy Gradient methods is the Generalized Advantage Estimation, that is a way to compute a better Advantage Function which can be tuned in bias and variance. It is not a policy gradient algorithm in its own, it is just an improved Advantage Function that can be used in any policy gradient algorithm. It is presented in Chapter 5.

4. Off-Policy Policy Gradient

4.1 Motivations for Off-Policy

As it is evident watching the Policy Gradient algorithm, each trajectory is used only for one step of optimization in the policy network gradient ascent. This is because it is an on-policy algorithm: since the gradient of the expected return (equation 3.29) is an expectation with respect to the distribution of samples generated by the policy π_θ , the generated samples can be used only once, to improve π_θ . They cannot be used again to improve the next policy $\pi_{\theta'}$ (the new policy resulted from the optimization step) because they come from a different distribution, the one generated by the old policy. It may be tempting to use each trajectory more than once, but as [Schulman et al. 2017] report: "*doing so is not well-justified, and empirically it often leads to destructively large policy updates*". A method that instead allows to use each trajectory more than once is Proximal Policy Optimization, described in Ch. 8, which permits us to use the samples generated by the previous policy distribution if the new optimized policy distribution is not too different. It is still considered an on-policy algorithm, but you will see that it is also *slightly off-policy*, because it reuses recent samples, if the policy is not too different. Proximal Policy Optimization is based on the same theoretical framework of Trust Region Policy optimization, so before reading Ch. 8 it is necessary to read Ch. 7 at least until eq. 7.31

It would be even better if we could use not only the trajectories generated by recent policies (which do not differ much from the current policy), but also use trajectories generated by any policy, without caring of the difference with our target policy. In other words, it would be great to do off-policy learning. That could allow us to reuse all past trajectories, or to choose behavior policies that differ from the target policy in order to explore some particular state space, or even to use trajectories obtained by just observing other agents (Offline Reinforcement Learning). This could greatly speed up learning.

Intuitively, to do policy gradient using samples from a different policy, it is necessary to do *importance sampling* (explained in Appendix C), as described by [Degris et al. 2012], [Levine and Koltun 2013] and [Levine 2021b]: since the trajectory has been generated by the behavior policy but has a different probability of occurring if we used the target policy, we must do importance sampling using the ratio of the probability of the target policy π_θ to the behavior policy π_β (where β are the parameters of the neural network for the behavior policy).

4.2 The Approximated Off-Policy Policy Gradient

Let us derive the off-policy gradient formally following [Levine 2021c]:

Recalling from equation 1.3 the total return of a trajectory τ :

$$G(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t$$

And the probability of a trajectory τ from eq. 1.4:

$$P(\tau|\pi) = \rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi(a_t|s_t)$$

The expected return under the behavior policy π_β is :

$$J(\pi_\beta) = E_{\tau \sim \pi_\beta} [G(\tau)] \quad (4.1)$$

Using importance sampling under the behavior policy π_β we can express the expected return for the target policy π_θ as:

$$J(\pi_\theta) = E_{\tau \sim \pi_\beta} \left[\frac{P(\tau|\pi_\theta)}{P(\tau|\pi_\beta)} G(\tau) \right] \quad (4.2)$$

From that we can compute the gradient:

$$\begin{aligned} \nabla_\theta J(\pi_\theta) &= \nabla_\theta E_{\tau \sim \pi_\beta} \left[\frac{P(\tau|\pi_\theta)}{P(\tau|\pi_\beta)} G(\tau) \right] \\ &= E_{\tau \sim \pi_\beta} \left[\nabla_\theta \left(\frac{P(\tau|\pi_\theta)}{P(\tau|\pi_\beta)} G(\tau) \right) \right] \end{aligned}$$

and since only $P(\tau|\pi_\theta)$ depends on θ :

$$= E_{\tau \sim \pi_\beta} \left[\frac{\nabla_\theta P(\tau|\pi_\theta)}{P(\tau|\pi_\beta)} G(\tau) \right]$$

applying the log – derivative trick of eq. 3.1:

$$\begin{aligned}
&= E_{\tau \sim \pi_\beta} \left[\frac{P(\tau|\pi_\theta)}{P(\tau|\pi_\beta)} \nabla_\theta \log P(\tau|\pi_\theta) G(\tau) \right] \\
&= E_{\tau \sim \pi_\beta} \left[\frac{\rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi_\theta(a_t|s_t)}{\rho_0(s_0) \prod_{t=0}^{T-1} P(s_{t+1}|s_t, a_t) \pi_\beta(a_t|s_t)} \nabla_\theta \log P(\tau|\pi_\theta) G(\tau) \right]
\end{aligned}$$

initial state and transition probabilities are the same for both policies and get cancelled:

$$\begin{aligned}
&= E_{\tau \sim \pi_\beta} \left[\left(\prod_{t=0}^{T-1} \frac{\pi_\theta(a_t|s_t)}{\pi_\beta(a_t|s_t)} \right) \nabla_\theta \log P(\tau|\pi_\theta) G(\tau) \right] \\
&= E_{\tau \sim \pi_\beta} \left[\left(\prod_{t=0}^{T-1} \frac{\pi_\theta(a_t|s_t)}{\pi_\beta(a_t|s_t)} \right) \left(\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t|s_t) \right) G(\tau) \right]
\end{aligned} \tag{4.3}$$

That is basically the common policy gradient of eq. 3.8 with importance sampling applied using the trajectory probabilities, where only the policy probabilities for actions are used in the ratio. We can improve that formula using the rewards-to-go (as in section 3.2) to decrease variance: instead of multiplying the weighted gradient of log probability by the total return $G(\tau)$, we multiply it only by the sum of the rewards in which the current action has a causality relationship, that is from the current time t to the end of the trajectory. This means that also the weighing must be changed accordingly, multiplying each reward by the multiplication of only the policy ratios from the beginning of the trajectory to the time of the reward:

$$\nabla_\theta J(\pi_\theta) = E_{\tau \sim \pi_\beta} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t|s_t) \left(\sum_{l=t}^{T-1} \gamma^l r_l \left(\prod_{t'=0}^l \frac{\pi_\theta(a_{t'}|s_{t'})}{\pi_\beta(a_{t'}|s_{t'})} \right) \right) \right]$$

common multiplication terms before t may be grouped together out of summation

$$= E_{\tau \sim \pi_\beta} \left[\sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t|s_t) \left(\prod_{t'=0}^t \frac{\pi_\theta(a_{t'}|s_{t'})}{\pi_\beta(a_{t'}|s_{t'})} \right) \left(\sum_{l=t}^{T-1} \gamma^l r_l \left(\prod_{t''=t+1}^l \frac{\pi_\theta(a_{t''}|s_{t''})}{\pi_\beta(a_{t''}|s_{t''})} \right) \right) \right] \tag{4.4}$$

Let us name:

$$\begin{aligned}\Pi_{current} &= \prod_{t'=0}^t \frac{\pi_{\theta}(a_{t'}|s_{t'})}{\pi_{\beta}(a_{t'}|s_{t'})} \\ \Pi_{after} &= \prod_{t''=t+1}^l \frac{\pi_{\theta}(a_{t''}|s_{t''})}{\pi_{\beta}(a_{t''}|s_{t''})}\end{aligned}\tag{4.5}$$

The first multiplication grouping $\Pi_{current}$ is the importance weight for the actions until current time t , without any future action weighing, while the second multiplication grouping Π_{after} is inside the inner of the summation and multiplies each reward r_l by the importance weight of actions after the current time t until the time of the reward l (included).

This way of computing the gradient has still some issue: the multiplication of the policy probability ratios $\Pi_{current}$ and Π_{after} can grow exponentially large.

To mitigate that it is possible to ignore Π_{after} and drop that multiplication: the resulting formula will not be a proper gradient anymore but it may still improve the policy because it will be equivalent to an update of policy iteration [Levine 2021c], even if with some exception that we will analyze later.

$$\nabla_{\theta}^{\approx} J(\pi_{\theta}) = E_{\tau \sim \pi_{\beta}} \left[\sum_{t=0}^{T-1} \nabla_{\theta} \log \pi_{\theta}(a_t|s_t) \left(\prod_{t'=0}^t \frac{\pi_{\theta}(a_{t'}|s_{t'})}{\pi_{\beta}(a_{t'}|s_{t'})} \right) \left(\sum_{l=t}^{T-1} \gamma^l r_l \right) \right]\tag{4.6}$$

I use the symbol $\nabla_{\theta}^{\approx}$ to make it clear that it is an approximation. That is equivalent to the gradient of expected return following the target policy π_{θ} until (and including) the “current” timestamp t and then supposing that the agent will follow the behavior policy π_{β} (because importance sampling is not applied to those steps).

To see if this still improves the target policy, even if after t there is not importance sampling and it is like the expectation follows the behavior policy β , let us examine three cases:

- A) $V^{\theta}(s_{t+1}) = V^{\beta}(s_{t+1})$
- B) $V^{\theta}(s_{t+1}) < V^{\beta}(s_{t+1})$
- C) $V^{\theta}(s_{t+1}) > V^{\beta}(s_{t+1})$

In case (A) the value of the next state will be the same, so the approximate gradient will be equal to the true gradient.

In case (B) the approximated gradient is a half-way between the gradient of the target policy and the gradient of a better policy ! So, in some sense it will update the current policy towards choosing the action in s_t that would be the best if the actions for states after t were chosen by the better policy. So basically, in case (B) the update will make the target policy behave more like a better policy, as if actions for next states were already improved, “believing” that in future also the actions taken in next states will converge towards the ones of the better policy.

In case (C) the approximate gradient is smaller than the true gradient. We can identify three subcases. Subcase (C1): if the approximate gradient has the same sign of the true gradient, but a smaller magnitude, the target policy will still improve, even if slower. Subcase (C2): if the approximate gradient has the same sign but a bigger magnitude, that may happen only if the true gradient is negative, the update may be too big, even if in the right direction, and that may slow down learning (maybe even distort the policy). Subcase (C3): if the approximate gradient has the opposite sign of the true gradient, that may happen when $V^\beta(s_{t+1})$ is negative, and so big in magnitude to cancel the positivity of the part of the approximate gradient computed following the target policy. In this case the update would go in the opposite direction of the correct update, and this will jeopardize learning.

So, we see that in the two subcases (C2) and (C3) we may incur in some issues.

We may wonder if, adding $V^\beta(s_{t+1})$ (that is the expectation of $\sum_{l=t+1}^{T-1} \gamma^l r_l$) instead of $V^\theta(s_{t+1})$ (that is the expectation of $\sum_{l=t+1}^{T-1} \gamma^l r_l \left(\prod_{t''=t+1}^l \frac{\pi_\theta(a_{t''}|s_{t''})}{\pi_\beta(a_{t''}|s_{t''})} \right)$) to $\gamma^t r_t$, it is like adding $V^\theta(s_{t+1})$ plus a baseline = $(V^\beta(s_{t+1}) - V^\theta(s_{t+1}))$ that would mean that we are just using a baseline and so what we add would be zero in expectation. But this is not the case, because what we are adding depends on a_t , not only on s_t , (because we are adding the expected value of s_{t+1} , not of s_t) so we are not adding a proper baseline and this means that we are actually distorting the gradient (as in subcase C2, or in subcase C3 with the possibility of an upgrade in the wrong direction !).

Now, a further approximation would be to use as importance sampling ratio only the ratio of the probability of choosing the action a_t under the target policy to the behavior policy, dropping the multiplication of the ratios for the previous time. That is:

$$\nabla_{\theta}^{\approx} J(\pi_{\theta}) = E_{\tau \sim \pi_{\beta}} \left[\sum_{t=0}^{T-1} \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\beta}(a_t | s_t)} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) \left(\sum_{l=t}^{T-1} \gamma^l r_l \right) \right] \quad (4.7)$$

Since $E_{\tau \sim \pi_{\beta}} [\sum_{l=t}^{T-1} \gamma^l r_l]$ is equivalent to $E_{\tau \sim \pi_{\beta}} [Q^{\beta}(a_t | s_t)]$ we can write it also as:

$$\nabla_{\theta}^{\approx} J(\pi_{\theta}) = E_{\tau \sim \pi_{\beta}} \left[\sum_{t=0}^{T-1} \frac{\pi_{\theta}(a_t | s_t)}{\pi_{\beta}(a_t | s_t)} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) Q^{\beta}(a_t | s_t) \right] \quad (4.8)$$

Also this approximation has not a theoretical justification.

4.3 Alternative Derivation

There is another way of computing the same approximation of the off-policy policy gradient, presented by [Degrís et al. 2012] (where it is named “Off-PAC”, short for “Off-Policy Actor-Critic”).

If we name $d_{\beta}(s)$ the hypothetical stationary distribution of behavior policy π_{β} :

$$d_{\beta}(s) = \lim_{t \rightarrow \infty} \text{Prob}(s_t = s | s_0; \beta) \quad (4.9)$$

Let us define a performance indicator where the states are occurring following the behavior policy π_{β} but the action-value is computed as if actions followed the target policy π_{θ} :

$$\begin{aligned} J_{\beta}(\pi_{\theta}) &= \sum_s d_{\beta}(s) \sum_a Q^{\theta}(a, s) \pi_{\theta}(a | s) \\ &= E_{s \sim d_{\beta}} \left[\sum_a Q^{\theta}(a, s) \pi_{\theta}(a | s) \right] \end{aligned} \quad (4.10)$$

We can rewrite it to make more evident that in $Q^{\theta}(a, s)$ the action a is distributed with policy π_{θ} as well as any subsequent action in the computation of Q^{θ} :

$$J_{\beta}(\pi_{\theta}) = E_{s \sim d_{\beta}} \left[\sum_a^A E_{s' \sim P} \left[R(s, a, s') + \gamma E_{a' \sim \theta} [Q^{\theta}(s', a')] \right] \pi_{\theta}(a|s) \right] \quad (4.11)$$

Its gradient with respect to θ is:

$$\nabla_{\theta} J_{\beta}(\pi_{\theta}) = \nabla_{\theta} E_{s \sim d_{\beta}} \left[\sum_a^A Q^{\theta}(a, s) \pi_{\theta}(a|s) \right]$$

applying the chain rule of calculus:

$$= E_{s \sim d_{\beta}} \left[\sum_a^A \left(Q^{\theta}(a, s) \nabla_{\theta} \pi_{\theta}(a|s) + \pi_{\theta}(a|s) \nabla_{\theta} Q^{\theta}(a, s) \right) \right] \quad (4.12)$$

The equation 4.12 above is the basis on which [Degris et al. 2012] derive their version of off-policy policy gradient. Unfortunately, it is not theoretically justified to compute the gradient in that way if the policy function is approximated (as with a neural network).

Let us give an intuitive explanation about that.

Affirming that doing gradient ascent on $\nabla_{\theta} J_{\beta}(\pi_{\theta})$ optimizes policy π_{θ} is like asserting that:

$$J_{\beta}(\pi_{\theta'}) \geq J_{\beta}(\pi_{\theta}) \text{ implies } J(\pi_{\theta'}) \geq J(\pi_{\theta}) \quad (4.13)$$

Where $J(\pi_{\theta})$ is, as defined in equation 1.5, the expected return with states obtained following policy π_{θ} , while $J(\pi_{\theta'})$ is the expected return with states obtained following policy $\pi_{\theta'}$.

We know that when the policy function is approximated, as when using a neural network to express the policy function, a change in parameters may not only change the policy function for the states in the trajectory used by the algorithm, but also for other states (this is called “aliasing”). Maybe that parameters change will improve the returns from a state that is frequently visited under that policy, but at the same time it will worsen the returns from a state that is rarely visited under the same policy, and as a whole the policy will improve because the rarely visited state has little impact to the whole expected return.

Imagine that to increase $J_\beta(\pi_\theta)$ the changes we make in the policy π_θ are such that they increase the value $V^\theta(s_a)$ of some state s_a that occurs very frequently on π_β and very rarely on π_θ but at the same time those policy changes decrease the value $V^\theta(s_b)$ of some state s_b that occurs very rarely on π_β and very frequently on π_θ , that means we would have:

$$J_\beta(\pi_{\theta'}) > J_\beta(\pi_\theta) \text{ and } J(\pi_{\theta'}) < J(\pi_\theta) \quad (4.14)$$

Which would contradict eq. 4.13.

Is that possible to happen with a gradient ascent if we have tabular policy ? No, because in tabular policies each state policy is independent from every other states' policies: changing the policy of a state will not change the policy of other states (there is not "aliasing"). That implies that increasing the value of a state s_a through a change in the policy on that state s_a will not decrease the value of another state s_b : either it will let the value of s_b the same (if from the state s_b it is not possible to reach the state s_a) or it will increase the value of s_b (if from the state s_b it is possible to reach the state s_a following the policy π_θ , that means that from s_b there is a probability greater than zero to reach a state that has increased its values, so also the total $V^\theta(s_b)$ has increased).

On the opposite, when using a function approximator like a neural network as a policy function, the change of the parameters θ resulting from a gradient ascent step that increases the value of some state s_a may result in modifying also the way the function approximator computes the policy for another state s_b (as already wrote, a phenomenon named by some authors "*aliasing*" or "*state-aliasing*"), possibly also decreasing the value $V^\theta(s_b)$.

So, the case of eq. 4.14 may happen and contradict eq. 4.13, that means that doing gradient ascent on $\nabla_\theta J_\beta(\pi_\theta)$ is not theoretically justified to improve policy π_θ when using neural networks as policy function approximators.

We do not need to do a dedicated analysis of the cases in which, as result of maximizing $J_\beta(\pi_\theta)$ (eq. 4.10), some states of policy π_θ decrease their frequency while other increase it. The reason is that if a state s_a decreases its frequency, another state s_b must increase its own, and if $V^\theta(s_a) > V^\theta(s_b)$ it will lead to decreasing the values of some states from which s_a was reachable, because some of them will reach more frequently than before the lower value state s_b and less frequently s_a . On the opposite, if $V^\theta(s_a) < V^\theta(s_b)$, some states from which s_a was reachable will increase their value. Anyway all those occurrences are just cases of either state value increase or decrease, and so they are already included in the analysis we just made when discussing if $J_\beta(\pi_{\theta'}) \geq J_\beta(\pi_\theta)$ implies $J(\pi_{\theta'}) \geq J(\pi_\theta)$.

Another discussion of the mismatch of using a behavior policy to estimate the target policy, together with a proposed solution, can be found in [Liu et al. 2019].

Let us continue now with the derivation of off-policy policy gradient by [Degris et al. 2012] from eq 4.12 . Now, it is quite difficult to compute $\nabla_{\theta} Q^{\theta}(a, s)$, so an easier approximation could be to drop the second term in the summation: $\pi_{\theta}(a|s) \nabla_{\theta} Q^{\theta}(a, s)$. This will give us the following approximation:

$$\begin{aligned}
\nabla_{\theta}^{\approx} J_{\beta}(\pi_{\theta}) &= E_{s \sim d_{\beta}} \left[\sum_a^A \left(Q^{\theta}(a, s) \nabla_{\theta} \pi_{\theta}(a|s) \right) \right] \\
&= E_{s \sim d_{\beta}} \left[\sum_a^A \left(\frac{\pi_{\beta}(a, s) \pi_{\theta}(a, s)}{\pi_{\beta}(a, s) \pi_{\theta}(a, s)} Q^{\theta}(a, s) \nabla_{\theta} \pi_{\theta}(a|s) \right) \right] \\
&= E_{s \sim d_{\beta}} \left[\sum_a^A \left(\pi_{\beta}(a, s) \frac{\pi_{\theta}(a, s)}{\pi_{\beta}(a, s)} Q^{\theta}(a, s) \frac{\nabla_{\theta} \pi_{\theta}(a|s)}{\pi_{\theta}(a, s)} \right) \right] \\
&= E_{s \sim d_{\beta}, a \sim \pi_{\beta}} \left[\frac{\pi_{\theta}}{\pi_{\beta}} Q^{\theta}(a, s) \nabla_{\theta} \log \pi_{\theta}(a|s) \right] \\
&= E_{\beta} \left[\frac{\pi_{\theta}(a, s)}{\pi_{\beta}(a, s)} Q^{\theta}(a, s) \nabla_{\theta} \log \pi_{\theta}(a|s) \right]
\end{aligned} \tag{4.15}$$

At this point [Degris et al. 2012] introduce a further approximation replacing $Q^{\theta}(a, s)$ with the off-policy return, that in expectation is equivalent to $Q^{\beta}(a, s)$. So we obtain:

$$\nabla_{\theta}^{\approx} J_{\beta}(\pi_{\theta}) = E_{\beta} \left[\frac{\pi_{\theta}(a, s)}{\pi_{\beta}(a, s)} Q^{\beta}(a, s) \nabla_{\theta} \log \pi_{\theta}(a|s) \right] \tag{4.16}$$

It is easy to see that sampling that is equivalent to sampling $\nabla_{\theta}^{\approx} J(\pi_{\theta})$ in eq. 4.8.

[Degris et al. 2012] give a partial justification for this approximation that is correct only for tabular representations of the policy and not for versions that use function approximations such as neural networks, that interests us when doing Deep Reinforcement Learning.

In Ch. 9 we will see the Deterministic Policy Gradient algorithm [Silver et al. 2014], which is an off-policy policy gradient method that does not use importance sampling.

Algorithm 4.1 Off-Policy Policy Gradient with Rewards-To-Go

Require: Policy network step size η

Require: Initialize parameters θ of network π_θ with small random values

Require: A behaviour policy π_β

For $k = 0, 1, 2, \dots$ do:

collect trajectories $D_k = \{ \tau_i \}$ using policy π_β

compute rewards-to-go $G(\tau_{i,t})$, for each t of each τ_i

estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \frac{\pi_\theta(a_t|s_t)}{\pi_\beta(a_t|s_t)} \nabla_\theta \log \pi_\theta(a_t|s_t) G(\tau_t)$$

update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g}_k$$

end of for

5. Generalized Advantage Estimation

[Schulman et al. 2016] introduced a particular version of the Advantage Function, the “Generalized Advantage Estimation”, with proven variance reduction property. It is intended to be used with the policy gradient family of algorithms, using undiscounted rewards. So there is not a time-dependent parameter γ to discount the rewards, but there is another parameter named γ with a different meaning but same mathematical behaviour: it is meant to downweigh rewards corresponding to delayed effects. Thus, the meaning is different from the “discount”, it does not mean that delayed effects are less “useful” at current time (that would be the meaning of the classical time-discount), but it still discounts the delayed effects, and doing so it introduces bias, because in this way the total return will be smaller in absolute value. On the other hand, it will decrease variance, for the same reason: smaller absolute value means smaller variance among different trajectories. So, the meaning of this parameter $\gamma \in [0,1]$ here is to be a switch for smoothly parametrize a trade-off between bias and variance, with bias increasing and variance decreasing when γ decreases, and bias decreasing and variance increasing when γ tends to 1.

Then this Generalized Advantage Estimation has another characteristic that decreases variance but introduces bias: it computes the estimate of the returns as an exponentially weighted average of k different n -step bootstrap estimates, where n goes from 1 to k . To be clearer, if we define:

$$\delta_t^V = r_t + \gamma V(s_{t+1}) - V(s_t) \quad (5.1)$$

Imagine having a trajectory of length at least $k > 1$, then we could build a 1-step bootstrap estimate of the advantage function at time t :

$$\hat{A}_t^{(1)} = r_t + \gamma V(s_{t+1}) - V(s_t) = \delta_t^V \quad (5.2)$$

But, if $k > 2$ we could also build a 2-step bootstrap estimate, or if $k > 3$ a 3-step estimate, or a k -step estimate etc. :

$$\hat{A}_t^{(2)} = r_t + \gamma r_{t+1} + \gamma^2 V(s_{t+2}) - V(s_t) = \delta_t^V + \gamma \delta_{t+1}^V$$

$$\hat{A}_t^{(3)} = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \gamma^3 V(s_{t+3}) - V(s_t) = \delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V$$

...

$$\hat{A}_t^{(k)} = -V(s_t) + \sum_{l=t}^{t+k-1} \gamma^{l-t} r_l + \gamma^k V(s_{t+k}) = \sum_{l=t}^{t+k-1} \gamma^{l-t} \delta_t^V \quad (5.3)$$

If we recall eq. 3.28, we may also notice that:

$$\hat{A}_t^{(k)} = \widehat{\widehat{G}}_k(\tau_t) - V(s_t) \quad (5.4)$$

In infinite horizon, k may tend to ∞ .

$$\hat{A}_t^{(\infty)} = -V(s_t) + \sum_{l=t}^{\infty} \gamma^{l-t} r_l = \sum_{l=t}^{\infty} \gamma^{l-t} \delta_l^V \quad (5.5)$$

Now, a 1-step estimate has low variance but high bias, while a 5-step estimate has bigger variance and lower bias, the more steps we include in the estimate the bigger the variance and the lower the bias. It is possible to compute an exponentially weighted average of all n -steps estimators, parametrized by λ such that when $\lambda = 0$ the weighted average coincides with $A_t^{(1)}$ and when $\lambda = 1$ it coincides with $A_t^{(\infty)}$, so that λ is another parameter that may be used to tune the bias/variance trade-off.

To obtain that, let us note that the series $\sum_{k=0}^{\infty} \lambda^k$, with $\lambda \in [0,1]$, is equal to $1/(1-\lambda)$. Hence $(1-\lambda) \sum_{k=0}^{\infty} \lambda^k = 1$. So, a summation of infinite terms, each weighted by $(1-\lambda)\lambda^k$, has the total sum of weights = 1.

The “Generalized Advantage Function” hence is:

$$\hat{A}_t^{GAE(\gamma, \lambda)} = (1-\lambda) \sum_{k=1}^{\infty} \lambda^{k-1} \hat{A}_t^{(k)} \quad (5.6)$$

And a more practical formulation may be obtained:

$$\begin{aligned}
\hat{A}_t^{GAE(\gamma, \lambda)} &= (1 - \lambda)(\hat{A}_t^{(1)} + \lambda \hat{A}_t^{(2)} + \lambda^2 \hat{A}_t^{(3)} + \dots) \\
&= (1 - \lambda)(\delta_t^V + \lambda(\delta_t^V + \gamma \delta_{t+1}^V) + \lambda^2(\delta_t^V + \gamma \delta_{t+1}^V + \gamma^2 \delta_{t+2}^V) + \dots) \\
&= (1 - \lambda)(\delta_t^V(1 + \lambda + \lambda^2 + \lambda^3 + \dots) + \gamma \delta_{t+1}^V(\lambda + \lambda^2 + \lambda^3 + \lambda^4 + \dots) \\
&\quad + \gamma^2 \delta_{t+2}^V(\lambda^2 + \lambda^3 + \lambda^4 + \lambda^5 + \dots) + \dots) \\
&= (1 - \lambda)(\delta_t^V \left(\frac{1}{1 - \lambda} \right) + \gamma \delta_{t+1}^V \left(\frac{\lambda}{1 - \lambda} \right) + \gamma^2 \delta_{t+2}^V \left(\frac{\lambda^2}{1 - \lambda} \right) + \dots) \\
&= \sum_{l=t}^{\infty} (\lambda \gamma)^{l-t} \delta_t^V
\end{aligned} \tag{5.7}$$

This GAE hence can be used instead of the Advantage Function and put in place of Φ_t in the policy gradient (see equation 3.29). This permits to tune variance and bias with two parameters: γ controls how far in time to consider the rewards relevant (the closer to 1, the farther in time), and λ that decides how much the return component of the advantage has to be similar to a 1-step bootstrap (if $\lambda=0$) or progressively to a Monte Carlo return (if $\lambda=1$).

6. Natural Policy Gradient

6.1 Motivation

To improve the speed of training in Policy Gradient one right thing to do can be to scale the gradient in a way that gradient ascent works better. It is possible that the gradient is not pointing to the direction of steepest ascent, and it may be convenient to modify or scale the gradient differently for each parameter in order to make it point to the direction of greater policy improvement, through the so-called “Natural Policy Gradient” [Kakade 2002], [Peters et al. 2003], [Bagnell and Schneider 2003] (also known as “Covariant Policy Gradient”), similarly to what has been studied in supervised learning [Amari 1998].

In fact, some parameters of the neural network affect the policy more than others: when we are doing gradient ascent with a certain learning rate, we are putting a limit on how much parameters can be changed, but this is not equivalent to put a limit on how much the policy can be changed. A faster learning would happen if we were able to fix the rate at which the policy is changed, and have larger rates for the parameters with little influence on the policy and smaller rates for the parameters with greater influence on the policy.




Figure 6.1

An hypothetical example of how gradients are computed in standard/vanilla Policy Gradient (on the left), and in Natural Policy Gradient (on the right). We have only two parameters here: θ_1 , represented on the horizontal axis, and θ_2 , represented on the vertical axis. Blue lines are zones of equivalent

expected return (return landscape). Black arrows are normalized gradients. The black half-circle in the middle of the bottom is representative of the optimum expected return. In standard Policy Gradient we see how the gradient used for the update is not directed at the optimum. The original example can be found in [Peters et al. 2003].




Figure 6.2

The analogous problem as it is usually depicted in Supervised Learning using gradient descent, where the dashed lines are zones of equivalent error function, the central dot is the optimum, the red line would be the steepest descent, while the taken path usually follows a more wandering trajectory, represented by the small circles connected by segments. The two axes represent the two parameters θ_1 and θ_2 of the machine learning system.

6.2 Constraining the New Policy

One way to compute a better gradient consists of putting a constraint to each gradient ascent step, such that the new policy and the old policy are not too different as distributions. An

appropriate computation for that could be the Kullback-Leibler divergence (see Appendix A.1), enforcing the constraint $D_{KL}(\pi_{\theta'}, || \pi_{\theta}) < \epsilon$.

The Kullback-Leibler divergence is defined as $D_{KL}(\pi_{\theta'}, || \pi_{\theta}) = E_{\theta'}[\log \pi_{\theta'} - \log \pi_{\theta}]$.

Second order Taylor expansion of KL divergence is approximated by (see Appendix A.6, eq. a.25):

$$D_{KL}(\pi_{\theta'}, || \pi_{\theta}) \approx \frac{1}{2}(\theta' - \theta)^T \mathbf{F}_{\pi_{\theta}}(\theta' - \theta) \quad (6.1)$$

Where $\mathbf{F}_{\pi_{\theta}}$ is the Fisher-information matrix of the policy π_{θ} .

The Fisher-information matrix of a distribution is the expected value of the covariance matrix of the score, given parameters θ . The score is the gradient of log-likelihood with respect to the parameters (see Appendix A.3 and A.4). Hence:

$$\mathbf{F}_{\pi_{\theta}} = E_{\pi_{\theta}}[\nabla_{\theta} \log \pi_{\theta}(a|s) \nabla_{\theta} \log \pi_{\theta}(a|s)^T] \quad (6.2)$$

That can be approximated with samples from the trajectory:

$$\widehat{\mathbf{F}}_{\pi_{\theta}} = \frac{1}{N} \sum_{i=1}^N \nabla_{\theta} \log \pi_{\theta}(a_i|s_i) \nabla_{\theta} \log \pi_{\theta}(a_i|s_i)^T \quad (6.3)$$

The common gradient ascent step for policy gradient is:

$$\theta' \leftarrow \theta + \eta \nabla_{\theta} J(\pi_{\theta}) \quad (6.4)$$

That can be seen as optimizing a constrained first order Taylor expansion of $J(\pi_{\theta})$:

$$\begin{aligned} \theta' &\leftarrow \arg \max_{\theta'} \nabla_{\theta} J(\pi_{\theta})^T (\theta' - \theta) \\ &\text{such that } \|\theta' - \theta\|^2 \leq \epsilon \end{aligned} \quad (6.5)$$

Explanation: maximizing a first order Taylor expansion of $J(\pi_\theta)$ over θ' means maximizing $J(\pi_\theta) + \nabla_\theta J(\pi_\theta)^T(\theta' - \theta)$, but $J(\pi_\theta)$ is fixed for any θ' , so the solution is equivalent to only maximize $\nabla_\theta J(\pi_\theta)^T(\theta' - \theta)$. Now, that is the maximization of a dot product between two vectors: $\nabla_\theta J(\pi_\theta)$ and $(\theta' - \theta)$. The first vector is fixed, we can choose θ' so to have a second vector that maximizes the dot product. The constraint says that the second vector must have maximum Euclidean norm of $\sqrt{\epsilon}$. The maximum dot product from a fixed first vector and a second vector of fixed length happens when the second vector is a scaled version of the first one (precisely, scaled to the point of having length = $\sqrt{\epsilon}$).

That would imply:

$$\theta' \leftarrow \theta + \sqrt{\frac{\epsilon}{\|\nabla_\theta J(\pi_\theta)\|^2}} \nabla_\theta J(\pi_\theta) \quad (6.6)$$

But as I wrote before, instead of imposing a constraint over $\|\theta' - \theta\|^2$ it would be better to impose a constraint over the distributions, such as $D_{KL}(\pi_{\theta'}, \pi_\theta) < \epsilon$.

Given the relation between $D_{KL}(\pi_{\theta'}, \pi_\theta)$ and Fisher-information matrix from equation 6.1, and given that we can compute an approximation of \mathbf{F}_{π_θ} from samples (as in eq. 6.3), such a constraint in the KL divergence can be imposed in a gradient descent step with the following update [Peters and Schaal 2008], [Levine 2021a]:

$$\theta' \leftarrow \theta + \eta \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)$$

$$\text{with } \eta = \sqrt{\frac{2\epsilon}{\nabla_\theta J(\pi_\theta)^T \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)}} \quad (6.7)$$

This is called “Natural gradient”, and its usage permits a faster learning with a smoother gradient ascent.

The computation of η as in eq. 6.7 is necessary only if you want to adhere to a strict a-priori fixed ϵ , otherwise you can use an arbitrarily small value for η : doing so will not change the direction of the change.

6.3 Derivation of the Formula

The derivation of η as in eq. 6.7 can be found in Appendix A.1 of [Peters 2007] and we propose it here in an equivalent form in the following lines, showing how to obtain eq. 6.7.

Let us define $\delta = (\theta' - \theta)$. So we can redefine the optimization problem, starting from eq. 6.5, using the Kullback-Leibler constraint, and the eq. 6.1 approximation, as:

$$\begin{aligned} \delta &\leftarrow \arg \max_{\delta} \nabla_{\theta} J(\pi_{\theta})^T \delta \\ \text{such that } &\frac{1}{2} (\theta' - \theta)^T \mathbf{F}_{\pi_{\theta}} (\theta' - \theta) < \epsilon \end{aligned} \tag{6.8}$$

This is a constrained optimization problem, solvable by Lagrangian Multiplier method. The resulting Lagrangian is the following:

$$\Lambda(\delta, \lambda) = \nabla_{\theta} J(\pi_{\theta})^T \delta + \lambda \left(\epsilon - \frac{1}{2} \delta^T \mathbf{F}_{\pi_{\theta}} \delta \right) \tag{6.9}$$

Now, all partial derivatives of $\Lambda(\delta, \lambda)$ should be zero.

$$\nabla_{\delta} \Lambda(\delta, \lambda) = \nabla_{\theta} J(\pi_{\theta}) - \lambda \frac{1}{2} 2 \mathbf{F}_{\pi_{\theta}} \delta = 0$$

$$\nabla_{\theta} J(\pi_{\theta}) = \lambda \mathbf{F}_{\pi_{\theta}} \delta$$

$$\lambda^{-1} \mathbf{F}_{\pi_{\theta}}^{-1} \nabla_{\theta} J(\pi_{\theta}) = \lambda^{-1} \mathbf{F}_{\pi_{\theta}}^{-1} \lambda \mathbf{F}_{\pi_{\theta}} \delta$$

$$\delta = \lambda^{-1} \mathbf{F}_{\pi_{\theta}}^{-1} \nabla_{\theta} J(\pi_{\theta}) \tag{6.10}$$

Now we plug the δ found in eq. 6.10 into the constraint $\frac{1}{2} \delta^T \mathbf{F}_{\pi_{\theta}} \delta = \epsilon$ and we have the dual function:

$$\frac{1}{2} \lambda^{-1} (\mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta))^T \mathbf{F}_{\pi_\theta} \lambda^{-1} \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta) = \epsilon$$

$$\frac{1}{2\epsilon} (\mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta))^T \nabla_\theta J(\pi_\theta) = \lambda^2$$

because of matrix and vector property $(Mv)^T = v^T M^T$

$$\lambda^2 = \frac{1}{2\epsilon} \nabla_\theta J(\pi_\theta)^T (\mathbf{F}_{\pi_\theta}^{-1})^T \nabla_\theta J(\pi_\theta)$$

since \mathbf{F}_{π_θ} is symmetric, also $\mathbf{F}_{\pi_\theta}^{-1}$ is symmetric, and for symmetric matrices $M = M^T$

$$\lambda = \sqrt{\frac{\nabla_\theta J(\pi_\theta)^T \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)}{2\epsilon}}$$

(6.11)

With λ the Lagrange multiplier. Now we plug it into eq. 6.10:

$$\delta = \sqrt{\frac{2\epsilon}{\nabla_\theta J(\pi_\theta)^T \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)}} \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)$$

(6.12)

Since δ is the name we used for $(\theta' - \theta)$, we know that:

$$\theta' - \theta = \sqrt{\frac{2\epsilon}{\nabla_\theta J(\pi_\theta)^T \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)}} \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)$$

$$\theta' = \theta + \sqrt{\frac{2\epsilon}{\nabla_\theta J(\pi_\theta)^T \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)}} \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)$$

now if we name $\eta = \sqrt{\frac{2\epsilon}{\nabla_\theta J(\pi_\theta)^T \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)}}$ we have:

$$\theta' = \theta + \eta \mathbf{F}_{\pi_\theta}^{-1} \nabla_\theta J(\pi_\theta)$$

The final 2 lines of eq. 6.13 are equivalent to eq. 6.7 .

Now, let us see the algorithm for Natural Policy Gradient, that as you may already understood is an on-policy algorithm because it updates the policy that has been used to generate samples and must use the latest policy to generate new trajectories.

Algorithm 6.1 Natural Policy Gradient

Require: Policy network step size η (optional)

Require: Value network step size ω

Require: Initialize parameters θ of network π_θ with small random values

Require: Initialize parameters ψ of network V_ψ with small random values

For $k = 0, 1, 2, \dots$ do:

collect trajectories $D_k = \{ \tau_i \}$ using policy $\pi_k(\theta_k)$

compute rewards-to-go $G(\tau_{i,t})$, for each t of each τ_i

compute advantage estimates \widehat{A}_t using current estimate of value function $V_{\psi k}$:

$$\widehat{A}_t = G(\tau_t) - V_{\psi k}(s_t)$$

estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \widehat{A}_t$$

compute approximate average Fisher Information Matrix:

$$\widehat{F}_{\pi_\theta} = \frac{1}{\sum_{\tau \in D_k} N_\tau} \sum_{\tau \in D_k} \sum_{t=0}^{N_\tau-1} \nabla_\theta \log \pi_\theta(a_t | s_t) \nabla_\theta \log \pi_\theta(a_t | s_t)^T$$

compute update vector for policy parameters

$$\Delta_k = \widehat{F}_{\pi_\theta}^{-1} \widehat{g}_k$$

compute the step size η either with npg $\eta = \sqrt{\frac{2\epsilon}{\widehat{g}_k^T \widehat{F}_{\pi_\theta}^{-1} \widehat{g}_k}}$ or with other method

update the policy network:

$$\theta_{k+1} \leftarrow \theta_k + \eta \Delta_k$$

For $z = 0, 1, 2, \dots, Z-1$ do a value function gradient descent iteration:

estimate the value function gradients as:

$$\widehat{h}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \frac{1}{T} \sum_{t=0}^{T-1} \nabla_{\psi} (V_{\psi}(s_t) - G(\tau_t))^2$$

update the value function with a gradient descent step (or other method):

$$\psi_{k+1} \leftarrow \psi_k + \omega \widehat{h}_k$$

end of for

end of for

7. Trust Region Policy Optimization

7.1 Problem formalization

The Trust Region Policy Optimization is an on-policy algorithm in which an alternative optimization problem is posed, such that allows to maximize the length of the step of the policy gradient ascent.

The Reinforcement Learning problem here is framed in a different formulation.

In the classical Policy Gradient formulation, we aim at maximizing $J(\pi_\theta)$, that for ease of reading we may now write $J(\theta)$. But equivalently, we may maximize the performance of a policy (the one to be maximized) with respect to a fixed policy (the one used until now to generate trajectories). If we call $\pi_{\theta'}$ the optimized policy (parametrized by the new parameters θ') and π_θ the fixed policy, we may want to maximize $J(\theta') - J(\theta)$.

Following [Schulman et al. 2015], it is useful to note that:

$$J(\theta') - J(\theta) = E_{\tau \sim \pi_{\theta'}} \left[\sum_t \gamma^t A^{\pi_\theta}(s_t, a_t) \right] \quad (7.1)$$

Proof:

Recall equation 1.12 and 1.13:

$$Q^{\pi_\theta}(s_t, a) = E_{s_{t+1} \sim P} [R(s_t, a, s_{t+1}) + \gamma V^{\pi_\theta}(s_{t+1})]$$

$$A^{\pi_\theta}(s_t, a) = Q^{\pi_\theta}(s_t, a) - V^{\pi_\theta}(s_t)$$

Then:

$$\begin{aligned} & E_{\tau \sim \pi_{\theta'}} \left[\sum_t \gamma^t A^{\pi_\theta}(s_t, a_t) \right] \\ &= E_{\tau \sim \pi_{\theta'}} \left[\sum_t \gamma^t (R(s_t, a, s_{t+1}) + \gamma V^{\pi_\theta}(s_{t+1}) - V^{\pi_\theta}(s_t)) \right] \end{aligned}$$

$$\begin{aligned}
&= E_{\tau \sim \pi_{\theta'}} \left[-V^{\pi_{\theta}}(s_0) + \sum_t \gamma^t R(s_t, a, s_{t+1}) \right] \\
&= E_{\tau \sim \pi_{\theta'}} [-V^{\pi_{\theta}}(s_0)] + E_{\tau \sim \pi_{\theta'}} \left[\sum_t \gamma^t R(s_t, a, s_{t+1}) \right]
\end{aligned}$$

distribution of s_0 does not depend on policy but only on initial state distribution $\rho_0 \Rightarrow$

$$\begin{aligned}
&= E_{s_0 \sim \rho_0} [-V^{\pi_{\theta}}(s_0)] + E_{\tau \sim \pi_{\theta'}} \left[\sum_t \gamma^t R(s_t, a, s_{t+1}) \right] \\
&= E_{\tau \sim \pi_{\theta}} [-V^{\pi_{\theta}}(s_0)] + E_{\tau \sim \pi_{\theta'}} \left[\sum_t \gamma^t R(s_t, a, s_{t+1}) \right] \\
&= -J(\theta) + J(\theta')
\end{aligned}$$

■

(7.2)

Now, equation 7.1 means that the difference between the expected return of the policy parametrized by θ' and the expected return of the policy parametrized by θ is equal to the expectation of the advantage function $A^{\pi_{\theta}}(s_t, a_t)$ over trajectories distributed by the policy parametrized by θ' . (Since the expectation of the trajectories follows the distribution by policy θ' , it implies that the actions a_t are distributed by $\pi_{\theta'}$, so $A^{\pi_{\theta}}(s_t, a_t)$ in eq. 7.1 it is the advantage function of the actions selected by the new policy θ' with respect to the fixed policy θ).

We can improve the policy if we can maximize the right-hand side of eq. 7.1 (because the greater is the difference between the expected return of new policy and the expected return of the old policy, the greater is the right-hand side). To do that, one strategy is to have an equivalent equation in a form that we can sample, so that we can run gradient ascent with respect to the policy parameters.

7.2 Ideal Objective

We can define the probability of being in state s at time t , depending on policy parametrized by θ as $Prob(s_t = s | \theta)$.

Then we can define the frequency $\xi_\theta(s)$ of a state s as the number of times that s is expected to be visited, computed as unnormalized (it is a frequency, not a probability) and time-discounted, under the policy π_θ parametrized by θ :

$$\xi_\theta(s) = Prob(s_0 = s | \theta) + \gamma Prob(s_1 = s | \theta) + \gamma^2 Prob(s_2 = s | \theta) + \dots \quad (7.3)$$

We can rewrite eq. 7.1 in a way to sum over time, then over states, and then over actions. Then in a second passage we can write it in a way to sum over states first, and in the last passage we plug eq. 7.3, referred to policy θ' , into it:

$$\begin{aligned} J(\theta') - J(\theta) &= \sum_t \sum_s Prob(s_t = s | \theta') \sum_a \pi_{\theta'}(a_t | s_t) \gamma^t A^{\pi_\theta}(s_t, a_t) \\ J(\theta') - J(\theta) &= \sum_s \sum_t \gamma^t Prob(s_t = s | \theta') \sum_a \pi_{\theta'}(a_t | s_t) A^{\pi_\theta}(s_t, a_t) \\ J(\theta') - J(\theta) &= \sum_s \xi_{\theta'}(s) \sum_a \pi_{\theta'}(a_t | s_t) A^{\pi_\theta}(s_t, a_t) \\ J(\theta') &= J(\theta) + \sum_s \xi_{\theta'}(s) \sum_a \pi_{\theta'}(a_t | s_t) A^{\pi_\theta}(s_t, a_t) \end{aligned} \quad (7.4)$$

This equation entails that any new policy parametrized by θ' which in every state has a better or equal expected advantage function with respect to the previous policy parametrized by θ , i.e. $\sum_a \pi_{\theta'}(a_t | s_t) A^{\pi_\theta}(s_t, a_t) \geq 0$, improves the total expected return (or leaves the total expected return unchanged if the expected advantage function is zero at each state).

7.3 Local Approximation of Objective

Now, the right-hand side seems in a more manageable form to be maximized than eq. 7.1, but if we look carefully, we see that to sample it we would need to follow the frequency of states $\xi_{\theta'}(s)$, referred to the new policy θ' , but at that point we only know the old policy θ and we can sample only with that. Hence, we use a local approximation that uses the old policy for the frequencies (please note the usage of $\xi_{\theta}(s)$ instead of $\xi_{\theta'}(s)$):

$$L_{\xi_{\theta}}(\theta') = J(\theta) + \sum_s \xi_{\theta}(s) \sum_a \pi_{\theta'}(a_t|s_t) A^{\pi_{\theta}}(s_t, a_t) \quad (7.5)$$

It has to be noted that (from eq. 7.1, 7.4 and eq. 7.5):

$$J(\theta') = J(\theta) + E_{\tau \sim \pi_{\theta'}} \left[\sum_t \gamma^t A^{\pi_{\theta}}(s_t, a_t) \right] = J(\theta) + E_{s \sim P(\tau|\pi_{\theta'})} \left[\sum_t \gamma^t E_{a \sim \pi_{\theta'}(a|s)} [A^{\pi_{\theta}}(s_t, a)] \right] \quad (7.6)$$

$$L_{\xi_{\theta}}(\theta') = J(\theta) + E_{s \sim P(\tau|\pi_{\theta})} \left[\sum_t \gamma^t E_{a \sim \pi_{\theta'}(a|s)} [A^{\pi_{\theta}}(s_t, a)] \right] \quad (7.7)$$

This approximation $L_{\xi_{\theta}}(\theta')$ can be used in place of $J(\theta')$ only when the new distribution $\pi_{\theta'}(a_t|s_t)$ is not too different from the old distribution $\pi_{\theta}(a_t|s_t)$, it must stay inside a “trust region” (hence the name of the algorithm).

7.4 Objective Lower Bound

Now we need to use the concept of Kullback-Leibler divergence so the reader who does not know about it may read Appendix A.1 before proceeding.

Considering the status s where the Kullback-Leibler divergence of $\pi_{\theta}(\cdot|s)$ from $\pi_{\theta'}(\cdot|s)$ is maximal, we define the value of such divergence as D_{KL}^{max} :

$$D_{KL}^{max}(\pi_{\theta}, \pi_{\theta'}) = \max_s [D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\theta'}(\cdot | s))] \quad (7.8)$$

Then, we define the maximum absolute value that the Advantage function can possibly assume, in any possible pairs of state and action, as v :

$$v = \max_{s,a} |A^{\pi_{\theta}}(s, a)| \quad (7.9)$$

Then it can be proved, that

$$J(\theta') \geq L_{\xi\theta}(\theta') - C D_{KL}^{max}(\pi_{\theta}, \pi_{\theta'})$$

$$\text{where } C = \frac{4 v \gamma}{(1 - \gamma)^2} \quad (7.10)$$

We may consider C as a penalty coefficient: the bigger the difference between π_{θ} and $\pi_{\theta'}$, the bigger the penalty.

Now we define:

$$\bar{A}(s) = E_{a' \sim \pi_{\theta'}(\cdot | s)} [A^{\pi_{\theta}}(s, a')] \quad (7.11)$$

So we can rewrite eq. 7.6 and 7.7 as:

$$J(\theta') = J(\theta) + E_{s \sim P(\tau | \pi_{\theta'})} \left[\sum_t \gamma^t \bar{A}(s_t) \right] \quad (7.12)$$

$$L_{\xi\theta}(\theta') = J(\theta) + E_{s \sim P(\tau | \pi_{\theta})} \left[\sum_t \gamma^t \bar{A}(s_t) \right] \quad (7.13)$$

To prove eq. 7.10, as in [Schulman et al. 2017], we must bound the difference between $J(\theta')$ and $L_{\xi\theta}(\theta')$.

To do that we begin measuring how much the two policies agree: we create a “coupled policy pair” that means that we create a joint distribution over a pair of policy actions (a, a') , where $a \sim \pi_\theta(\cdot | s)$ and $a' \sim \pi_{\theta'}(\cdot | s)$.

We define $(\pi_\theta, \pi_{\theta'})$ as an “ α – coupled policy pair” if for all states s the joint distribution $(a, a')|s$ is such that $Prob(a \neq a' | s) \leq \alpha$.

(Please note the graphical difference between a , the first letter of latin alphabet used to denote an action, and α , that is “alpha” from greek alphabet, used to denote the value of a probability boundary).

Now:

$$\bar{A}(s) = E_{a' \sim \pi_{\theta'}(\cdot | s)}[A^{\pi_\theta}(s, a')] = E_{(a, a') \sim (\pi_\theta, \pi_{\theta'})}[A^{\pi_\theta}(s, a')]$$

since $E_{a \sim \pi_\theta(\cdot | s)}[A^{\pi_\theta}(s, a)] = 0$ we can insert it into the expectation

$$\bar{A}(s) = E_{(a, a') \sim (\pi_\theta, \pi_{\theta'})}[A^{\pi_\theta}(s, a') - A^{\pi_\theta}(s, a)]$$

when $a = a'$, we have $A^{\pi_\theta}(s, a') - A^{\pi_\theta}(s, a) = 0$, so we can consider only when $a \neq a'$

$$\bar{A}(s) = Prob(a \neq a' | s) E_{(a, a') \sim (\pi_\theta, \pi_{\theta'}) | a \neq a'}[A^{\pi_\theta}(s, a') - A^{\pi_\theta}(s, a)]$$

$$\bar{A}(s) \leq \alpha E_{(a, a') \sim (\pi_\theta, \pi_{\theta'}) | a \neq a'}[A^{\pi_\theta}(s, a') - A^{\pi_\theta}(s, a)]$$

$$|\bar{A}(s)| \leq \alpha \cdot 2 \max_{s, a} (|A^{\pi_\theta}(s, a)|)$$

(7.14)

Now, if $(\pi_\theta, \pi_{\theta'})$ is an α – coupled policy pair also the trajectories τ and τ' (respectively generated by π_θ and $\pi_{\theta'}$, using the same random seed) can be coupled. We can compute the advantage of $\pi_{\theta'}$ over π_θ at each timestamp t and decompose this expectation depending on whether $\pi_{\theta'}$ agrees with π_θ at all timestamps $i < t$. We use n_t to count the occurrences in which $\pi_{\theta'}$ disagrees with π_θ (i.e. $a \neq a'$) before t . To simplify the notation, we use $s_t \sim \pi_\theta$ instead of the previous $s \sim P(\tau | \pi_\theta)$.

$$E_{s_t \sim \pi_{\theta'}}[\bar{A}(s_t)] = Prob(n_t = 0) E_{s_t \sim \pi_{\theta'} | n_t = 0}[\bar{A}(s_t)] + Prob(n_t > 0) E_{s_t \sim \pi_{\theta'} | n_t > 0}[\bar{A}(s_t)]$$

(7.15)

We can write the same for the trajectories generated by π_θ :

$$E_{s_t \sim \pi_\theta}[\bar{A}(s_t)] = \text{Prob}(n_t = 0)E_{s_t \sim \pi_\theta | n_t=0}[\bar{A}(s_t)] + \text{Prob}(n_t > 0)E_{s_t \sim \pi_\theta | n_t>0}[\bar{A}(s_t)] \quad (7.16)$$

Note that the terms with $n_t = 0$ are equals, because in there the two policies agree at all timestamps:

$$\begin{aligned} E_{s_t \sim \pi_{\theta'} | n_t=0}[\bar{A}(s_t)] &= E_{s_t \sim \pi_\theta | n_t=0}[\bar{A}(s_t)] \\ \Rightarrow \\ E_{s_t \sim \pi_{\theta'}}[\bar{A}(s_t)] - E_{s_t \sim \pi_\theta}[\bar{A}(s_t)] &= \\ = \text{Prob}(n_t > 0)(E_{s_t \sim \pi_{\theta'} | n_t>0}[\bar{A}(s_t)] - E_{s_t \sim \pi_\theta | n_t>0}[\bar{A}(s_t)]) \end{aligned} \quad (7.17)$$

Now, by definition, $\pi_{\theta'}$ agrees with π_θ at each timestamp with probability $\geq 1 - \alpha$, hence we have that:

$$\text{Prob}(n_t = 0) \geq (1 - \alpha)^t$$

$$\text{Prob}(n_t > 0) \leq 1 - (1 - \alpha)^t \quad (7.18)$$

Consider that:

$$|E_{s_t \sim \pi_{\theta'} | n_t>0}[\bar{A}(s_t)] - E_{s_t \sim \pi_\theta | n_t>0}[\bar{A}(s_t)]| \leq |E_{s_t \sim \pi_{\theta'} | n_t>0}[\bar{A}(s_t)]| + |E_{s_t \sim \pi_\theta | n_t>0}[\bar{A}(s_t)]| \quad (7.19)$$

And, by eq 7.14:

$$|E_{s_t \sim \pi_{\theta'} | n_t>0}[\bar{A}(s_t)]| + |E_{s_t \sim \pi_\theta | n_t>0}[\bar{A}(s_t)]| \leq \alpha \cdot 4 \max_{s,a}(|A^{\pi_\theta}(s, a)|) \quad (7.20)$$

Now if we compute the absolute value of eq. 7.17 and plug eq. 7.18 and 7.20 into it:

$$|E_{s_t \sim \pi_{\theta'} | n_t=0}[\bar{A}(s_t)] - E_{s_t \sim \pi_\theta}[\bar{A}(s_t)]| = \text{Prob}(n_t > 0)|E_{s_t \sim \pi_{\theta'} | n_t>0}[\bar{A}(s_t)] - E_{s_t \sim \pi_\theta | n_t>0}[\bar{A}(s_t)]|$$

$$|E_{s_t \sim \pi_{\theta'}}[\bar{A}(s_t)] - E_{s_t \sim \pi_{\theta}}[\bar{A}(s_t)]| \leq (1 - (1 - \alpha)^t) \cdot \alpha \cdot 4 \max_{s,a} (|A^{\pi_{\theta}}(s, a)|) \quad (7.21)$$

This bounds the difference in expected advantage at each timestep t . If we sum over time, we bound the difference between $J(\theta')$ and $L_{\xi\theta}(\theta')$. Subtracting eq. 7.13 from eq. 7.12:

$$|J(\theta') - L_{\xi\theta}(\theta')| = \sum_t \gamma^t |E_{s_t \sim \pi_{\theta'}}[\bar{A}(s_t)] - E_{s_t \sim \pi_{\theta}}[\bar{A}(s_t)]| \quad (7.22)$$

Now, using $v = \max_{s,a} |A^{\pi_{\theta}}(s, a)|$ as we already defined in eq. 7.9 with eq. 7.21 and eq. 7.22:

$$\begin{aligned} |J(\theta') - L_{\xi\theta}(\theta')| &\leq \sum_t \gamma^t (1 - (1 - \alpha)^t) 4 \alpha v \\ \text{any series } \sum_{k=0}^{\infty} \gamma^k \text{ with } \gamma \in [0,1] \text{ converges to } \frac{1}{1-\gamma} &\Rightarrow \\ \sum_t \gamma^t (1 - (1 - \alpha)^t) 4 \alpha v &= 4 \alpha v \left(\sum_t \gamma^t - \sum_t (\gamma(1 - \alpha))^t \right) = 4 \alpha v \left(\frac{1}{1-\gamma} - \frac{1}{1-\gamma(1-\alpha)} \right) \Rightarrow \\ |J(\theta') - L_{\xi\theta}(\theta')| &\leq 4 \alpha v \left(\frac{1}{1-\gamma} - \frac{1}{1-\gamma(1-\alpha)} \right) \\ |J(\theta') - L_{\xi\theta}(\theta')| &\leq \frac{4 \alpha^2 \gamma v}{(1-\gamma)(1-\gamma(1-\alpha))} \\ |J(\theta') - L_{\xi\theta}(\theta')| &\leq \frac{4 \alpha^2 \gamma v}{(1-\gamma)^2} \end{aligned} \quad (7.23)$$

Now, the right-hand side of eq. 7.23 is ≥ 0 , that implies that:

$$\text{if } J(\theta') \geq L_{\xi\theta}(\theta') \Rightarrow J(\theta') - L_{\xi\theta}(\theta') \leq \frac{4 \alpha^2 \gamma v}{(1-\gamma)^2} \Rightarrow J(\theta') \leq L_{\xi\theta}(\theta') + \frac{4 \alpha^2 \gamma v}{(1-\gamma)^2} \quad (7.24)$$

$$\text{if } J(\theta') \leq L_{\xi\theta}(\theta') \Rightarrow L_{\xi\theta}(\theta') - J(\theta') \leq \frac{4 \alpha^2 \gamma v}{(1-\gamma)^2} \Rightarrow J(\theta') \geq L_{\xi\theta}(\theta') - \frac{4 \alpha^2 \gamma v}{(1-\gamma)^2} \quad (7.25)$$

If the premise of eq. 7.24 is true (*if* $J(\theta') \geq L_{\xi\theta}(\theta')$), then of course it is true also $J(\theta') \geq (L_{\xi\theta}(\theta') - \text{anything positive})$, that means that it is true also $J(\theta') \geq L_{\xi\theta}(\theta') - \frac{4 \alpha^2 \gamma v}{(1-\gamma)^2}$ (the result of 7.25), because $\frac{4 \alpha^2 \gamma v}{(1-\gamma)^2}$ is always positive. So, in any case it holds:

$$J(\theta') \geq L_{\xi\theta}(\theta') - \frac{4 \alpha^2 \gamma v}{(1-\gamma)^2} \quad (7.26)$$

Now, following [Levin et al. 2009], proposition 2.7, we know that if we have two distributions $P(x)$ and $Q(y)$ whose Total Variation Distance is α (see Appendix A.2), formally $D_{TV}(P||Q) = \alpha$, then there exists a joint distribution (X, Y) whose marginals are $P(x)$ and $Q(y)$, for which the probability of $X = Y$ is equal to $1 - \alpha$.

Hence if we have two policies $\pi_{\theta'}$ and π_{θ} such that $\max_s [D_{TV}(\pi_{\theta}(\cdot | s) || \pi_{\theta'}(\cdot | s))] \leq \alpha$ then it is possible to define an α - *coupled policy pair* $(\pi_{\theta'}, \pi_{\theta})$.

So if we name:

$$D_{TV}^{max} = \max_s [D_{TV}(\pi_{\theta}(\cdot | s) || \pi_{\theta'}(\cdot | s))] \quad (7.27)$$

And if we choose $\alpha = D_{TV}^{max}$, we can rewrite eq. 7.26:

$$J(\theta') \geq L_{\xi\theta}(\theta') - \frac{4 \gamma v}{(1-\gamma)^2} (D_{TV}^{max})^2 \quad (7.28)$$

Now, it is known (see [Pollard 2000] Ch.3) that the Total Variation distance has the following relationship with Kullback-Leibler divergence: $D_{TV}(P||Q)^2 \leq D_{KL}(P||Q)$. So in eq. 7.28 we can plug in $D_{KL}^{max}(\pi_\theta, \pi_{\theta'})$:

$$J(\theta') \geq L_{\xi\theta}(\theta') - \frac{4 \gamma v}{(1 - \gamma)^2} D_{KL}^{max}(\pi_\theta, \pi_{\theta'}) \quad (7.29)$$

And we see that eq. 7.29 is equivalent to eq. 7.10, so we just proved eq. 7.10. ■

So this proves that $L_{\xi\theta}(\theta') - C D_{KL}^{max}(\pi_\theta, \pi_{\theta'})$ is a lower bound of $J(\theta')$.




Figure 7.1

Graphic rendition of $L_{\xi\theta}(\theta') - C D_{KL}^{max}(\pi_\theta, \pi_{\theta'})$ as a lower bound for $J(\theta')$. When $\theta' = \theta$ the three curves coincide.

7.5 Monotonic Improvement

Now, looking back at eq. 7.10, if we call θ_i the policy parameters at time i , and we call ω the policy parameters for any other different policy (for instance a policy to be optimized), we define:

$$M_i(\omega) = L_{\xi\theta_i}(\omega) - C D_{KL}^{max}(\pi_{\theta_i}, \pi_{\omega}) \quad (7.30)$$

Then, if we start from eq. 7.10, substituting θ' with ω and θ with θ_i , we have:

$$\begin{aligned} J(\omega) &\geq L_{\xi\theta_i}(\omega) - C D_{KL}^{max}(\pi_{\theta_i}, \pi_{\omega}) \\ \Rightarrow J(\omega) &\geq M_i(\omega) \end{aligned} \quad (7.31)$$

When $\omega = \theta_i$ it is easy to see that since the Kullback-Leibler divergence of two identical distribution is zero, and the advantage function of two identical policies is zero (that makes $L_{\xi\theta_i}(\theta_i) = J(\theta_i)$):

$$J(\theta_i) = M_i(\theta_i) \quad (7.32)$$

Now, let us say that at time $i + 1$ we have a policy parametrized by θ_{i+1} . This can be the case such as in a loop in which at each iteration we change the policy parameters. By eq. 7.31 we have:

$$J(\theta_{i+1}) \geq M_i(\theta_{i+1}) \quad (7.33)$$

We can combine 7.32 with 7.33:

$$J(\theta_{i+1}) - J(\theta_i) \geq M_i(\theta_{i+1}) - M_i(\theta_i) \quad (7.34)$$

This means that improving M_i implies monotonically improving the expected return ! This is also called “*Monotonic improvement theorem*”.

7.6 Surrogate Objective

So, we know what we want to optimize: $L_{\xi\theta}(\theta') - C D_{KL}^{max}(\pi_{\theta}, \pi_{\theta'})$ with respect to θ' .

As [Schulman et al. 2017] notice, using that C penalty coefficient is theoretically justified but leads to very small steps.

Alternatively, we could be able to take larger steps if we just maximize $L_{\xi\theta}(\theta')$ and put as a condition that the difference between the new policy distribution and the old policy distribution stays under a certain threshold δ (the “trust region”), measuring that difference with a Kullback-Leibler divergence. Hence, considering $L_{\xi\theta}(\theta')$ as defined in eq. 7.7 :

$$\begin{aligned} \underset{\theta'}{\text{maximize}} \quad L_{\xi\theta}(\theta') &= J(\theta) + E_{s \sim P(\tau|\pi_{\theta})} \left[\sum_t \gamma^t E_{a \sim \pi_{\theta'}(a|s)} [A^{\pi_{\theta}}(s_t, a)] \right] \\ \text{s. t.} \quad D_{KL}^{max}(\pi_{\theta}, \pi_{\theta'}) &\leq \beta \end{aligned} \tag{7.35}$$

But we need also to express $L_{\xi\theta}(\theta')$ as an expectation with respect to a known policy, to be able to sample it: our agent must follow a policy with respect to which the expectation is expressed. If we look at eq. 7.35 we see that the first term of $L_{\xi\theta}(\theta')$ is $J(\theta)$, that does not depend on θ' so it is not getting optimized (we could actually get rid of it from the equation). The second term $E_{s \sim P(\tau|\pi_{\theta})} [E_{a \sim \pi_{\theta'}(a|s)} [\gamma^t A^{\pi_{\theta}}(s, a)]]$ has the outer part that is an expectation with respect to policy π_{θ} (the old policy, that we know), and the inner part that is an expectation with respect to policy $\pi_{\theta'}$, that is an unknown. To be able to sample actions, we want to have all expectations in the formula being with respect to the old policy π_{θ} . We can rewrite the equation with also the inner expectation with respect to old policy π_{θ} . In doing so we must use Importance Sampling accordingly (Importance Sampling is described in Appendix C).

$$L_{\xi\theta}(\theta') = J(\theta) + E_{s \sim P(\tau|\pi_{\theta})} \left[\sum_t \gamma^t E_{a \sim \pi_{\theta}(a|s)} \left[\frac{\pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)} A^{\pi_{\theta}}(s_t, a) \right] \right] \tag{7.36}$$

Any series $\sum_{k=0}^{\infty} \gamma^k$ with $\gamma \in [0,1]$ converges to $1/(1-\gamma)$.

Hence, the sum $\sum_t \gamma^t$ can be replaced by $\frac{1}{1-\gamma}$.

$$L_{\xi\theta}(\theta') = J(\theta) + \frac{1}{1-\gamma} E_{s \sim P(\tau|\pi_\theta)} \left[E_{a \sim \pi_\theta(a|s)} \left[\frac{\pi_{\theta'}(a|s)}{\pi_\theta(a|s)} A^{\pi_\theta}(s, a) \right] \right] \quad (7.37)$$

Since the formula will be used for optimization, we could even get rid of the constant $\frac{1}{1-\gamma}$ when optimizing with respect to θ' . So, we can write a slightly different surrogate objective:

$$\begin{aligned} L_\theta(\theta') &= J(\theta) + E_{s \sim P(\tau|\pi_\theta)} \left[E_{a \sim \pi_\theta(a|s)} \left[\frac{\pi_{\theta'}(a|s)}{\pi_\theta(a|s)} A^{\pi_\theta}(s, a) \right] \right] \\ &= J(\theta) + E_{\tau \sim \pi_\theta} \left[\frac{\pi_{\theta'}(a|s)}{\pi_\theta(a|s)} A^{\pi_\theta}(s, a) \right] \end{aligned} \quad (7.38)$$

Using that as objective is further justified by the fact that the gradient of equation 7.38 matches the gradient of $J(\theta')$ if evaluated locally at the point $\theta' = \theta$:

$$\begin{aligned} \nabla_{\theta'} L_\theta(\theta')|_{\theta' = \theta} &= E_{\tau \sim \pi_\theta} \left[\frac{\nabla_{\theta'} \pi_{\theta'}(a|s)}{\pi_\theta(a|s)} A^{\pi_\theta}(s, a) \right] \Big|_{\theta' = \theta} \\ &= E_{\tau \sim \pi_\theta} [\log \nabla_{\theta'} \pi_{\theta'}(a|s) A^{\pi_\theta}(s, a)] \Big|_{\theta' = \theta} \\ &= E_{\tau \sim \pi_{\theta'}} [\log \nabla_{\theta'} \pi_{\theta'}(a|s) A^{\pi_\theta}(s, a)] \Big|_{\theta' = \theta} \\ &= \nabla_{\theta'} J(\theta') \Big|_{\theta' = \theta} \end{aligned} \quad (7.39)$$

This means that a small change in the policy parameters from θ to θ' such that $L_\theta(\theta')$ increases, makes also the new policy's expected return $J(\theta')$ increase with respect to the old policy expected return $J(\theta)$. This “small change” in policy parameters is expressed as the constraint on the Kullback-Leibler divergence, because the two policy distributions must be close.

When maximizing eq. 7.38 with respect to θ' it is easy to see that fixed term $J(\theta)$ is not dependent on θ' so we can get rid of it when optimizing.

What we want to optimize now is:

$$\begin{aligned} \underset{\theta'}{\text{maximize}} \quad & L_\theta(\theta') = E_{\tau \sim \pi_\theta} \left[\frac{\pi_{\theta'}(a|s)}{\pi_\theta(a|s)} A^{\pi_\theta}(s, a) \right] \\ \text{s. t.} \quad & D_{KL}^{max}(\pi_\theta, \pi_{\theta'}) \leq \beta \end{aligned} \tag{7.40}$$

The constraint on the KL divergence applies to all points in space and it is not practical to be concretely applied, so a heuristic approximation may be used instead, computing the average Kullback-Leibler divergence (which at computation time will be the average of the samples).

$$\begin{aligned} \underset{\theta'}{\text{maximize}} \quad & L_\theta(\theta') = E_{\tau \sim \pi_\theta} \left[\frac{\pi_{\theta'}(a|s)}{\pi_\theta(a|s)} A^{\pi_\theta}(s, a) \right] \\ \text{s. t.} \quad & E_{\tau \sim \pi_\theta} [D_{KL}(\pi_\theta(\cdot | s) || \pi_{\theta'}(\cdot | s))] \leq \beta \end{aligned} \tag{7.41}$$

At this point the optimization problem is completely defined. We need only an effective way to solve it. The original authors [Schulman et al. 2015] used $Q^{\pi_\theta}(s, a)$ instead of $A^{\pi_\theta}(s, a)$, which leads to an equivalent optimization problem because it changes the objective only by a constant, but saves from the burden of computing $V^{\pi_\theta}(s)$ to compute $A^{\pi_\theta}(s, a)$.

$$\begin{aligned} \underset{\theta'}{\text{maximize}} \quad & E_{\tau \sim \pi_\theta} \left[\frac{\pi_{\theta'}(a|s)}{\pi_\theta(a|s)} Q^{\pi_\theta}(s, a) \right] \\ \text{s. t.} \quad & E_{\tau \sim \pi_\theta} [D_{KL}(\pi_\theta(\cdot | s) || \pi_{\theta'}(\cdot | s))] \leq \beta \end{aligned} \tag{7.42}$$

$Q^{\pi_\theta}(s, a)$ is computed by an empirical estimate, such as the sum of discounted rewards of the sampled trajectory that starts with (s, a) , that is actually $G(\tau)$ of eq. 1.3 . In our pseudocode algorithm 7.1 at the end of the chapter we anyway use the advantage function instead of the Q function, because it converges faster, so we use eq. 7.41 instead of eq. 7.42 (they both are theoretically right and should converge to the same policy).

The parameters θ are then updated with the θ' found solving the constrained optimization problem, using the conjugate gradient algorithm followed by line search (in the original paper described at point 3 of paragraph 6, and Appendix C of [Schulman et al. 2015]): we are going to explain that in detail in the following part.

7.7 Constraining the Kullback-Leibler Divergence

Now, in order to simplify the procedure to analytically find the optimization update rule satisfying the Kullback-Leibler constraint, we approximate the objective $L_\theta(\theta')$ by its first order Taylor expansion $\widehat{L}_\theta(\theta')$. Note that we do this only to find a formula that considers the Kullback-Leibler constraint, but then at computation time we will use the real $L_\theta(\theta')$ gradient, and not the one simplified using its first order Taylor expansion $\widehat{L}_\theta(\theta')$. Since in $L_\theta(\theta')$ the parameters θ are fixed and the variable parameters are θ' , and the Taylor expansion is built around the point $\theta' = \theta$, we have that:

$$\widehat{L}_\theta(\theta') = L_\theta(\theta) + \nabla_{\theta'} L_\theta(\theta')^T \Big|_{\theta' = \theta} \cdot (\theta' - \theta)$$

$$\text{since } L_\theta(\theta) = 0 \Rightarrow$$

$$\widehat{L}_\theta(\theta') = \nabla_{\theta'} L_\theta(\theta')^T \Big|_{\theta' = \theta} \cdot (\theta' - \theta)$$

we plug eq. 7.39 in

$$\widehat{L}_\theta(\theta') = \nabla_{\theta'} J(\theta')^T \Big|_{\theta' = \theta} \cdot (\theta' - \theta)$$

(7.43)

The Kullback-Leibler divergence is formulated by the approximation described in Appendix A.6 (eq. a.25), that uses an approximated Fisher Information Matrix $\widehat{\mathbf{F}}_{\pi\theta}$.

$$\widehat{D}_{KL}(\pi(\cdot | s_t; \theta) || \pi(\cdot | s_t; \theta + \delta)) = \frac{1}{2} \delta^T \widehat{\mathbf{F}}_{\pi\theta} \delta \quad (7.44)$$

What we are doing until now, that is maximizing eq. 7.43 while keeping the KL divergence constrained, seems very similar to the Natural Policy Gradient (Ch. 6). In fact at this point we are maximizing:

$$\begin{aligned} \nabla_{\theta'} J(\theta')^T |_{\theta'} &= \theta \cdot (\theta' - \theta) \\ \text{s.t. } \frac{1}{2} \delta^T \widehat{\mathbf{F}}_{\pi\theta} \delta &\leq \beta \end{aligned} \quad (7.45)$$

And it is easy to see that eq. 7.45 is equivalent to eq. 6.8 of Natural Policy Gradient ! Solving 7.45 with the method of Lagrangian multipliers would be the same as we did in Ch. 6 and will give the same solution as eq. 6.7 and eq. 6.13, with the following parameters update rule:

$$\theta' \leftarrow \theta + \sqrt{\frac{2\beta}{\nabla_{\theta} J(\pi_{\theta})^T \mathbf{F}_{\pi_{\theta}}^{-1} \nabla_{\theta} J(\pi_{\theta})}} \mathbf{F}_{\pi_{\theta}}^{-1} \nabla_{\theta} J(\pi_{\theta}) \quad (7.46)$$

But TRPO differs in the way the computation is executed. To begin with, at computation time in TRPO instead of $\nabla_{\theta} J(\pi_{\theta})$ (which was obtained from the first order Taylor approximation) we use $\nabla_{\theta'} \left(\frac{\pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)} A^{\pi_{\theta}}(s, a) \right)$ that is the gradient of the surrogate objective (see eq. 7.41). Secondly, the Fisher Information Matrix is calculated in a different way compared to Natural Policy Gradient, but we will detail it in next paragraph. In addition, there are two problematic aspects: the former is that big matrices like $\mathbf{F}_{\pi_{\theta}}$ are costly to invert, and the latter is that the approximations used may cause a destructive update of the policy. Let us see how TRPO tackles those problems.

7.8 Inverting the Fisher Information Matrix via Conjugate Gradient

To compute the update rule of eq. 7.46 we need to find the inverse of the Fisher information Matrix F_{π_θ} . Computing the inverse of a big matrix can be very expensive in terms of computation and memory, but if the matrix is symmetrical and positive definite it is possible to do it in a way that is faster and less memory demanding than with traditional linear algebra methods (especially when the matrix is sparse): with the Conjugate Gradient algorithm [Hestenes and Stiefel 1952]. It turns out that the Fisher Information Matrix is symmetrical, and it is positive semidefinite: even when a matrix is not guaranteed to be definite but only semidefinite the Conjugate Gradient algorithm is able to converge [Hayami 2018]. The Conjugate Gradients can solve a linear system:

$$\mathbf{A}x = b \quad (7.47)$$

(with \mathbf{A} and b given) exploiting the fact that, if \mathbf{A} is symmetrical and positive definite, the solution of such a system (that is $x = \mathbf{A}^{-1}b$) can be found solving the minimization problem of the equation:

$$f(x) = \frac{1}{2}x^T \mathbf{A} x - x^T b \quad (7.48)$$

The Conjugate Gradient allows us to do it fast because it computes a particular version of steepest descent in which the direction of descent is chosen efficiently. A sketch of the Gradient Descent algorithm is reported in Appendix B.

Now, if in place of matrix \mathbf{A} we use the approximation of Fisher Information Matrix $\widehat{F_{\pi_\theta}}$, and if in place of b we use $\nabla_\theta J(\pi_\theta)$, we obtain that the solution found by the Conjugate Gradient algorithm is (substituting in $x = \mathbf{A}^{-1}b$):

$$x = \widehat{F_{\pi_\theta}}^{-1} \nabla_\theta J(\pi_\theta) \quad (7.49)$$

If we watch eq. 7.46 we notice that we do not need to find also $\widehat{\mathbf{F}}_{\pi\theta}^{-1}$ alone: we can just use our x , that is $\widehat{\mathbf{F}}_{\pi\theta}^{-1}\nabla_{\theta}J(\pi_{\theta})$. In fact:

$$\sqrt{\frac{2\beta}{\nabla_{\theta}J(\pi_{\theta})^T \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta})}} \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta}) = \sqrt{\frac{2\beta}{x^T \widehat{\mathbf{F}}_{\pi\theta} x}} x \quad (7.50)$$

Proof:

$$\begin{aligned} & \sqrt{\frac{2\beta}{x^T \widehat{\mathbf{F}}_{\pi\theta} x}} x = \\ & \sqrt{\frac{2\beta}{\left(\widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta})\right)^T \widehat{\mathbf{F}}_{\pi\theta} \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta})}} \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta}) = \\ & \sqrt{\frac{2\beta}{\nabla_{\theta}J(\pi_{\theta})^T \left(\widehat{\mathbf{F}}_{\pi\theta}^{-1}\right)^T \widehat{\mathbf{F}}_{\pi\theta} \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta})}} \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta}) = \end{aligned}$$

since the Fisher Information Matrix is symmetrical, its inverse is symmetrical \Rightarrow

$$= \sqrt{\frac{2\beta}{\nabla_{\theta}J(\pi_{\theta})^T \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta})}} \widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta}) \quad (7.51)$$

This means that we do not need to compute $\widehat{\mathbf{F}}_{\pi\theta}^{-1}$, we can be happy enough if we find $\widehat{\mathbf{F}}_{\pi\theta}^{-1} \nabla_{\theta}J(\pi_{\theta})$, and this is something that the Conjugate Gradient algorithm can do for us.

To use the Conjugate Gradient algorithm all we need is a function that takes vector v as input, and returns the multiplication $\mathbf{A}v$, where in our case \mathbf{A} would be our approximated Fisher Information Matrix $\widehat{\mathbf{F}}_{\pi\theta}$.

Now, this function that returns the multiplication $\widehat{\mathbf{F}}_{\pi\theta} \cdot v$ can be realized in at least two ways: one is described by original authors [Schulman et al. 2015] and consists in finding a manageable version of an approximation of the Fisher Information Matrix. We detail it in the next section. The other method is suggested by [Achiam & OpenAi 2020D] and does not

require to compute an approximation of the Fisher Information Matrix. It is a faster method, so it is recommended, and we detail it in section 7.8.2

7.8.1 Product of Average Hessian of KL

In this section we will detail the method of computing the approximated Fisher Information Matrix as described in the original paper [appendix C1 of Schulman et al. 2015]. This section is optional and can be skipped, since there is a faster way to compute the product between Fisher Information Matrix and a vector v without using directly the Fisher Information Matrix, and it is described in next section (section 7.8.2 “Average Gradient of Product by KL Gradient”).

As explained in our Appendix A.7, given two distributions of the same family $\pi_\theta(x)$ and $\pi_{\theta'}(x)$, the Hessian of Kullback-Leibler divergence of $\pi_\theta(x)$ from $\pi_{\theta'}(x)$, with respect to θ' , evaluated at $\theta' = \theta$ is equal to the Fisher Information Matrix of $\pi_\theta(x)$. Hence this can be used to approximate the Fisher Information Matrix: using the Hessian of Kullback-Leibler divergence instead of commonly using $E_{x \sim P_\theta(x)} [\nabla_\theta \log P_\theta(x) \nabla_\theta \log P_\theta(x)^T]$ (as in Appendix A.4).

$$\widehat{F}_{P_\theta}[i, j] = \frac{1}{N} \sum_t^N \frac{\partial^2}{\partial \theta'_i \partial \theta'_j} D_{KL}(\pi_\theta(\cdot | s_t) || \pi_{\theta'}(\cdot | s_t)) \quad (7.52)$$

To compute the Hessian of Kullback-Leibler divergence of $\pi_\theta(a, s)$ from $\pi_{\theta'}(x)$, for each sampled state s we need to take in consideration all possible actions a , not only the one that has been actually taken by the sampled trajectory. An analytic estimator must integrate over a for each sampled state s .

Let us detail an analytical way to compute the Fisher Information Matrix using the Hessian of Kullback-Leibler divergence. The KL divergence of old policy $\pi_\theta(\cdot | s)$ from new policy $\pi_{\theta'}(\cdot | s)$ is $D_{KL}(\pi_\theta(\cdot | s) || \pi_{\theta'}(\cdot | s))$. The new policy is unknown, it is what we aim to find with the TRPO algorithm. Each policy function outputs a vector of real numbers, that may be the mean of a distribution (e.g. a normal) for each action parameter when the actions are continuous, while when the actions are discrete the policy function outputs a vector of probability for each action (for each action dimension). With continuous actions if you do not have fixed variances the policy function (the neural network) will also output the variance values.

So let us call the outputs of the old and new policy functions respectively $\mu_{OLD}(s)$ and $\mu_{\theta'}(s)$. You may notice that for the old policy we do not put the parameters in the notation, we just put “old”, to explicit the fact that the old policy is unaffected by θ' parameters optimization because it depends on fixed parameters (that were indicated with θ in eq. 7.41 and 7.42).

We denote the Hessian of KL divergence $\pi_{\theta}(\cdot | s)$ from $\pi_{\theta'}(\cdot | s)$ in a simplified way, using the outputs of the policy functions, so to make evident what are the variables of the KL divergence function:

$$D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\theta'}(\cdot | s)) = kl(\mu_{OLD}(s), \mu_{\theta'}(s)) \quad (7.53)$$

Now, to compute the Hessian of $kl(\mu_{OLD}(s), \mu_{\theta'}(s))$ we first compute the gradient with respect to θ' . Since $kl(\cdot)$ is a function, whose parameters are functions themselves, we apply the chain rule of calculus:

$$\begin{aligned} \frac{\partial(kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial \theta'} &= \\ \frac{\partial(kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial(\mu_{OLD}(s))} \frac{\partial(\mu_{OLD}(s))}{\partial \theta'} + \frac{\partial(kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial(\mu_{\theta'}(s))} \frac{\partial(\mu_{\theta'}(s))}{\partial \theta'} &= \\ \text{since } \frac{\partial(\mu_{OLD}(s))}{\partial \theta'} \text{ equals to 0 because } \mu_{OLD}(s) \text{ does not depend on } \theta', \text{ the first term disappears} & \\ = \frac{\partial(kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial(\mu_{\theta'}(s))} \frac{\partial(\mu_{\theta'}(s))}{\partial \theta'} & \end{aligned} \quad (7.54)$$

This may be rewritten in a clearer way making evident that $\mu_{\theta'}(s)$ may be a vector function, and θ' is (usually) a vector, using subscripts to enumerate the element of the vector, with $\mu_{\theta'}(s)$ having U elements, and θ' having V elements. Since we are applying the chain rule of calculus, the partial derivative with respect to an element of θ' must sum the derivatives of every function parameter with respect to it, that is why on every row we have a summation. To ease the reading, we write μ_{OLD} instead of $\mu_{OLD}(s)$.

$$\frac{\partial(kl(\mu_{OLD}, \mu_{\theta'}(s)))}{\partial \theta'} =$$

$$\begin{aligned}
& \left[\frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_1)} \frac{\partial(\mu_{\theta'}(s)_1)}{\partial \theta'_1} + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_2)} \frac{\partial(\mu_{\theta'}(s)_2)}{\partial \theta'_1} + \dots + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_U)} \frac{\partial(\mu_{\theta'}(s)_U)}{\partial \theta'_1} \right] \\
& \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_1)} \frac{\partial(\mu_{\theta'}(s)_1)}{\partial \theta'_2} + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_2)} \frac{\partial(\mu_{\theta'}(s)_2)}{\partial \theta'_2} + \dots + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_U)} \frac{\partial(\mu_{\theta'}(s)_U)}{\partial \theta'_2} \\
& \dots \\
& \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_1)} \frac{\partial(\mu_{\theta'}(s)_1)}{\partial \theta'_V} + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_2)} \frac{\partial(\mu_{\theta'}(s)_2)}{\partial \theta'_V} + \dots + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_U)} \frac{\partial(\mu_{\theta'}(s)_U)}{\partial \theta'_V}
\end{aligned}
\tag{7.55}$$

This may be compactly rewritten in a way that is vector-element wise, with row index i , as:

$$\begin{aligned}
& \left[\frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial \theta'_i} \right] = \\
& \left[\frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_a)} \frac{\partial(\mu_{\theta'}(s)_a)}{\partial \theta'_i} \right]
\end{aligned}
\tag{7.56}$$

Where i stands for the element of the parameters vector θ' , and a stands for “compute this equation over all elements of $\mu_{\theta'}(s)$ using a as index, and sum all the results together”, as it is evident in eq. 7.55.

If we acknowledge that $\mu_{\theta'}(s)$ is a vector function and θ' is a vector we could also write it as the product of the Jacobian of $\mu_{\theta'}(s)$ wrt to θ' for the gradient of $kl(\mu_{OLD}(s), \mu_{\theta'}(s))$ wrt to $\mu_{\theta'}(s)$:

$$\begin{aligned}
& = \left(\frac{\partial(\mu_{\theta'}(s))}{\partial \theta'} \right)^T \frac{\partial(kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial(\mu_{\theta'}(s))} \\
& = \mathbf{J}(\mu_{\theta'}(s))^T \frac{\partial(kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial(\mu_{\theta'}(s))}
\end{aligned}
\tag{7.57}$$

Eq. 7.57 is like eq. 7.54 but with the Jacobian of $\mu_{\theta'}(s)$ and the gradient of $kl(\mu_{OLD}(s), \mu_{\theta'}(s))$ in switched order, to make it possible to have the matrix multiplication. Also we used $\mathbf{J}()$ to denote a Jacobian.

Now if we compute the Jacobian of 7.56 (or equivalently 7.57 or 7.55 or 7.54) wrt θ' we have the Hessian of $kl(\mu_{OLD}(s), \mu_{\theta'}(s))$ wrt θ' , here presented for any element of row i and column j .

$$\begin{aligned}
& \left[H \left(kl(\mu_{OLD}, \mu_{\theta'}(s)) \right) \right]_{i,j} = \\
& \left[\frac{\partial(\partial kl(\mu_{OLD}, \mu_{\theta'}(s)))}{\partial \theta'_j \partial \theta'_i} \right] = \\
& \left[\frac{\partial}{\partial \theta'_j} \left(\frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_a)} \frac{\partial(\mu_{\theta'}(s)_a)}{\partial \theta'_i} \right) \right] = \\
& \left[\frac{\partial}{\partial \theta'_j} \left(\frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_a)} \right) \frac{\partial(\mu_{\theta'}(s)_a)}{\partial \theta'_i} + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_a)} \frac{\partial}{\partial \theta'_j} \left(\frac{\partial(\mu_{\theta'}(s)_a)}{\partial \theta'_i} \right) \right] = \\
& \left[\frac{\partial^2 kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_b) \partial(\mu_{\theta'}(s)_a)} \frac{\partial(\mu_{\theta'}(s)_b)}{\partial \theta'_j} \frac{\partial(\mu_{\theta'}(s)_a)}{\partial \theta'_i} + \frac{\partial kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_a)} \frac{\partial^2(\mu_{\theta'}(s)_a)}{\partial \theta'_j \partial \theta'_i} \right] =
\end{aligned}$$

the second term vanishes so only the first term remains

$$= \left[\frac{\partial^2 kl(\mu_{OLD}, \mu_{\theta'}(s))}{\partial(\mu_{\theta'}(s)_b) \partial(\mu_{\theta'}(s)_a)} \frac{\partial(\mu_{\theta'}(s)_b)}{\partial \theta'_j} \frac{\partial(\mu_{\theta'}(s)_a)}{\partial \theta'_i} \right] \quad (7.58)$$

In the equation above i and j are the indices for each element of the parameters vector θ' , and so identify the row and column of the element in the Hessian. Also, both a and b stand for “compute this equation over all elements of $\mu_{\theta'}(s)$ using a and b as indices and sum all the results together”, that since they both appear at the same time is to be interpreted as: “to compute an element of the Hessian at position (i, j) , compute eq. 7.57 for every combination of a and b , and sum together all the results (considering a and b as two separate indices of the vector $\mu_{\theta'}(s)$)”.

Writing it in a Matrix-wise way we have:

$$\begin{aligned}
& H \left(kl(\mu_{OLD}, \mu_{\theta'}(s)) \right)_{wrt \theta'} = \\
& \left(\frac{\partial(\mu_{\theta'}(s))}{\partial \theta'} \right)^T \frac{\partial^2 (kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial(\mu_{\theta'}(s)_j) \partial(\mu_{\theta'}(s)_i)} \frac{\partial(\mu_{\theta'}(s))}{\partial \theta'} = \\
& \mathbf{J}(\mu_{\theta'}(s))^T \frac{\partial^2 (kl(\mu_{OLD}(s), \mu_{\theta'}(s)))}{\partial(\mu_{\theta'}(s)_j) \partial(\mu_{\theta'}(s)_i)} \mathbf{J}(\mu_{\theta'}(s))
\end{aligned} \tag{7.59}$$

A shorter way to write it is:

$$H \left(kl(\mu_{OLD}, \mu_{\theta'}(s)) \right)_{wrt \theta'} = \mathbf{J}^T \mathbf{M} \mathbf{J}$$

with \mathbf{J} Jacobian of $\mu_{\theta'}(s)$ wrt θ' and \mathbf{M} hessian of $kl(\mu_{OLD}(s), \mu_{\theta'}(s))$ wrt $\mu_{\theta'}(s)$

(7.60)

This is a much more compact representation than the full computed hessian of $kl(\mu_{OLD}(s), \mu_{\theta'}(s))$ wrt θ' .

If you recall, we found this manageable approximation of Fisher Information Matrix with the intention of using it within the Conjugate Gradient algorithm: for that algorithm we need to have a function that multiplies the Hessian/Fisher Information Matrix for a vector v , and since we have built $\widehat{\mathbf{F}}_{\pi\theta}$ as the average of $\mathbf{J}^T \mathbf{M} \mathbf{J}$ over all states of the transitions (i.e. we plugged 7.60 into 7.52), we can just do the matrix-vector multiplication.

$$\widehat{\mathbf{F}}_{\pi\theta} \cdot v = \left(\frac{1}{N} \sum_t (\mathbf{J}(s_t))^T \mathbf{M}(s_t) \mathbf{J}(s_t) \right) \cdot v$$

(7.61)

Since this multiplication $\widehat{\mathbf{F}}_{\pi\theta} \cdot v$ is used in the conjugate gradient algorithm to compute the direction of a descent vector for an iteration, it can even be an approximate computation, and original authors [Schulman et al. 2015] suggest to compute the approximate $\widehat{\mathbf{F}}_{\mathbf{p}_\theta}$ using only 10% of samples, in order to have a faster execution.

Anyway, there seems to be a faster way to obtain the Fisher Information Matrix product, and we see it in next section.

7.8.2 Average Gradient of Product by KL Gradient

A faster way of computing the product between Fisher Information Matrix and a vector v , is explained in [Achiam & OpenAi 2020D], and it does not directly compute the Fisher Information Matrix. It consists in computing the average of the gradient of the Kullback-Leibler divergence for all states, transpose it, multiply it by the vector v (it is a dot product between vectors), and compute the gradient of that result wrt to θ' .

As in the previous section, let us call the outputs of the old and new policy functions respectively $\mu_{OLD}(s)$ and $\mu_{\theta'}(s)$. For the old policy we do not put the parameters θ in the notation, we just put “old”, to explicit the fact that the old policy is unaffected by θ' parameters optimization because it depends on fixed parameters.

Expressed in formulas:

$$\begin{aligned}\widehat{\mathbf{F}}_{\pi\theta} \cdot v &= \frac{1}{N} \sum_t \nabla_{\theta'} \left(\left(\frac{\partial(kl(\mu_{OLD}, \mu_{\theta'}(s_t)))}{\partial \theta'} \right)^T \cdot v \right) \\ &= \nabla_{\theta'} \left(\left(\frac{1}{N} \sum_t \left(\frac{\partial(kl(\mu_{OLD}, \mu_{\theta'}(s_t)))}{\partial \theta'} \right)^T \right) \cdot v \right)\end{aligned}\tag{7.62}$$

Proof (considering the vector v as given, not depending on θ'):

$$\begin{aligned}\nabla_{\theta'} \left(\left(\frac{\partial(kl(\mu_{OLD}, \mu_{\theta'}(s)))}{\partial \theta'} \right)^T \cdot v \right) &= \\ \nabla_{\theta'} \left(\left(\frac{\partial(kl(\mu_{OLD}, \mu_{\theta'}(s)))}{\partial \theta'} \right)^T \right) \cdot v + \left(\frac{\partial(kl(\mu_{OLD}, \mu_{\theta'}(s)))}{\partial \theta'} \right)^T \nabla_{\theta'}(v) &= \\ \text{since } \nabla_{\theta'}(v) \text{ is zero } \Rightarrow &\end{aligned}$$

$$\begin{aligned}
&= \nabla_{\theta'} \left(\left(\frac{\partial (kl(\mu_{OLD}, \mu_{\theta'}(s)))}{\partial \theta'} \right)^T \right) \cdot v \\
&= H(kl(\mu_{OLD}, \mu_{\theta'}(s))) \cdot v
\end{aligned}
\tag{7.63}$$

Where H denotes the Hessian.

To use this procedure for the matrix-vector multiplication, when writing the code with an automatic differentiation package it is necessary to have the values of v disconnected from the parameters θ' in the differentiation graph, so that the gradient of v wrt θ' is zero.

7.9 Using the Conjugate Gradient

Now that we have a function to do the product between the Fisher Information Matrix and a vector, we can run the Conjugate Gradient method and obtain the vector x (see eq. 7.49), that is a vector that gives the direction in which we should be going from the old parameters θ to reach the new parameters θ' . In other words we should multiply the vector x by a constant η and sum the result to old parameters θ to obtain new parameters θ' . According to eq. 7.46 and 7.50, η should be equal to $\sqrt{\frac{2\beta}{x^T \widehat{F_{\pi\theta}} x}}$.

At this point it seems that we have everything we need for the update rule. But because of the approximations, it is not sure that with $\eta = \sqrt{\frac{2\beta}{x^T \widehat{F_{\pi\theta}} x}}$ the constraint on the average Kullback-Leibler distance will be respected. If the new and old distribution differ too much there is the possibility of a destructive update. So, we should better check if $\eta = \sqrt{\frac{2\beta}{x^T \widehat{F_{\pi\theta}} x}}$ makes the new distribution complying with the Kullback-Leibler divergence constraint. Also, we need to check that the updates are not worsening the policy, so we compute for each step of the trajectories the value of the surrogate objective $\frac{\pi_{\theta'}(a_t|s_t)}{\pi_{\theta}(a_t|s_t)} \hat{A}_t$ (see eq. 7.41) and we check that the average of it computed for the new policy is greater or equal than the average of it computed for the old policy. If either the KL divergence constraint is not satisfied, or the average surrogate objective has worsened, we exponentially decrease η until we find a value that satisfies those

conditions. This is basically a line search. If in a reasonable number of iterations, we can't find any decreased value of η that satisfies those conditions, we reject this policy update.

Let us now see all the pieces together with the TRPO algorithm pseudocode.

Until now we used to write θ for the “old” parameters and θ' for the “new” ones, we named the old policy as π_θ and the new one $\pi_{\theta'}$, and we called the output of the policy (the likelihood for the action) as μ_{OLD} or μ_θ for the old policy output and $\mu_{\theta'}$ for the new policy output. But since the algorithm will run an iteration of policy changes indexed by k on the same neural network, for clarity's sake we will name the “old” parameters and policy respectively as θ_k and π_{θ_k} . The “new” parameters and policy will be respectively θ_{k+1} and $\pi_{\theta_{k+1}}$.

Be aware that when computing gradients it will be with respect to θ_{k+1} , while θ_k and π_{θ_k} will be considered as fixed, so that parameters θ_k will not accumulate gradient. This means that using automatic differentiation libraries (such as PyTorch) you need to disable gradient accumulation when computing $\pi_{\theta_k}(a_t|s_t)$.

Algorithm 7.1 Trust Region Policy Optimization

Require: Kullback-Leibler divergence constraint β

Require: backtracking coefficient b , with $0 < b < 1$

Require: maximum number of backtracking steps U

Require: Value network step size ω

Require: Initialize parameters θ of network π_{θ_k} with small random values

Require: Initialize parameters ψ of network V_ψ with small random values

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using policy π_{θ_k}

 compute rewards-to-go $G(\tau_{i,t})$, for each t of each τ_i

 compute advantage estimates \hat{A}_t using current estimate of value function V_{ψ_k} :

$$\hat{A}_t = G(\tau_t) - V_{\psi_k}(s_t)$$

 (alternatively we can use GAE):

$$\hat{A}_t^{GAE(\gamma, \lambda)} = \sum_{l=t}^{|\tau|} (\lambda \gamma)^{l-t} \delta_l^V \text{ with } \delta_l^V = r_l + \gamma V_{\psi_k}(s_{l+1}) - V_{\psi_k}(s_l)$$

 compute $\pi_{\theta_k}(a_t|s_t)$ without accumulating gradients

$$\theta_{k+1} \leftarrow \theta_k$$

 estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \nabla_{\theta_{k+1}} \left(\frac{\pi_{\theta_{k+1}}(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)} \hat{A}_t \right)$$

set a function that returns the multiplication of a vector v by Fisher Inf. Matrix: $\widehat{F}_{\pi_{\theta_k}} \cdot v$

$$\text{either by } \widehat{F}_{\pi_{\theta_k}} \cdot v = \left(\frac{1}{\sum_{\tau \in D_k} N_\tau} \sum_{\tau \in D_k} \sum_{t=0}^{N_\tau-1} (\mathbf{J}(s_t))^T \mathbf{M}(s_t) \mathbf{J}(s_t) \right) \cdot v$$

$$\text{or (faster) by } \widehat{F}_{\pi_{\theta_k}} \cdot v = \nabla_{\theta_{k+1}} \left(\frac{1}{\sum_{\tau \in D_k} N_\tau} \left(\sum_{\tau \in D_k} \sum_{t=0}^{N_\tau-1} \left(\frac{\partial(\text{kl}(\mu_{\theta_k}, \mu_{\theta_{k+1}}(s_t)))}{\partial \theta_{k+1}} \right)^T \right) \cdot v \right)$$

run the Conjugate Gradient algorithm with that function to find update direction vector

$$x_k = \widehat{F}_{\pi_{\theta_k}}^{-1} \widehat{g}_k$$

$$\text{compute the step size upper limit } \eta = \sqrt{\frac{2\beta}{x_k^T \widehat{F}_{\pi_{\theta_k}} x_k}}$$

update the policy network by backtracking line search, with $j \in \{0, 1, 2 \dots U\}$ being the smallest number that satisfies the Kullback-Leibler divergence and that

improves the average surrogate objective $\frac{\pi_{\theta_{k+1}}(a|s)}{\pi_{\theta_k}(a|s)} \hat{A}_t$ on the samples:

$$\theta_{k+1} \leftarrow \theta_k + b^j \eta x_k$$

(no policy update if conditions are not satisfied)

For $z = 0, 1, 2, \dots, Z-1$ do a value function gradient descent iteration:

estimate the value function gradients as:

$$\widehat{h}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \frac{1}{T} \sum_{t=0}^{T-1} \nabla_{\psi_k} (V_{\psi_k}(s_t) - G(\tau_t))^2$$

update the value function with a gradient descent step (or other method):

$$\psi_k \leftarrow \psi_k + \omega \widehat{h}_k$$

end of for

$$\psi_{k+1} \leftarrow \psi_k$$

end of for

7.10 Vine Sampling Scheme

The original authors [Schulman et al. 2015] experimented also a sampling scheme named "Vine", applicable on any environment in which it is possible to restart from a certain well-defined state (as in a simulated environment). The Vine sampling consists of trying different

actions from the same starting point, generating different trajectories to compute an average Q estimate, which in this way will have a lower variance.

8. Proximal Policy Optimization

8.1 Base Idea

The Proximal Policy Optimization algorithm [Schulman et al. 2017] builds on the ideas and theoretical framework of Trust Region Policy Optimization. To understand this chapter it is necessary to have read the TRPO chapter (Ch. 7) until equation 7.41 included. In fact, PPO aims at optimizing the same eq. 7.41:

$$\begin{aligned} \underset{\theta'}{\text{maximize}} \quad & L_{\theta}(\theta') = E_{\tau \sim \pi_{\theta}} \left[\frac{\pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)} A^{\pi_{\theta}}(s, a) \right] \\ \text{s.t.} \quad & E_{\tau \sim \pi_{\theta}} [D_{KL}(\pi_{\theta}(\cdot | s) || \pi_{\theta'}(\cdot | s))] \leq \beta \end{aligned}$$

Where π_{θ} is the old policy and $\pi_{\theta'}$ is the new one.

While in TRPO there was only one policy gradient ascent iteration in each optimization round (in which the step length was computed so to not make the new policy too different from the old), in PPO the algorithm does more than one policy gradient ascent iteration in each optimization round using the same set of trajectories: the trajectories are re-used as long as the new policy is not too different from the old.

What distinguishes Proximal Policy Optimization from TRPO is that it uses a different strategy to constraint the KL divergence: (1) it uses ratio clipping in the objective function as a way to never have a too big gradient update, and (2) it checks when the old and new policies are too divergent so to stop reusing the samples generated by the old policy.

8.2 Surrogate Objective Clipping

The new surrogate objective function L_{θ}^{CLIP} uses clipping to limit how big the objective value can be, and is defined as:

$$L_{\theta}^{CLIP}(\theta') = E_{\tau \sim \pi_{\theta}} \left[\min \left(\frac{\pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)} A^{\pi_{\theta}}(s, a), \text{clip} \left(\frac{\pi_{\theta'}(a|s)}{\pi_{\theta}(a|s)}, 1 - \epsilon, 1 + \epsilon \right) A^{\pi_{\theta}}(s, a) \right) \right]$$

(8.1)

For ease of reading, let us call the ratio between the two probabilities $d(\theta) = \frac{\pi_{\theta'}(a|S)}{\pi_{\theta}(a|S)}$. So the surrogate objective can be written in a simplified way:

$$L_{\theta}^{CLIP}(\theta') = E_{\tau \sim \pi_{\theta}} [\min(d(\theta)A^{\pi_{\theta}} , \quad clip(d(\theta), 1 - \epsilon, 1 + \epsilon) A^{\pi_{\theta}})] \quad (8.2)$$

That means that, generating actions with policy θ , the objective is the expectation of the minimum among $d(\theta)A^{\pi_{\theta}}$ and $clip(d(\theta), 1 - \epsilon, 1 + \epsilon) A^{\pi_{\theta}}$.

The *clip* function imposes both a ceiling and a floor to $d(\theta)$: if $d(\theta) > 1 + \epsilon$ it returns $1 + \epsilon$, if $d(\theta) < 1 - \epsilon$ it returns $1 - \epsilon$, otherwise it returns $d(\theta)$.

The min function then has the result of returning a lower bound on the unclipped objective. In this way, when the probability ratio would make the objective improve, it is bounded to be $< 1 + \epsilon$, so to not have too big steps. While when it would make the objective worse it is unbounded towards negative infinity. When the objective is negative, it is bounded towards zero so to always have some negative value that has an impact on the policy update.

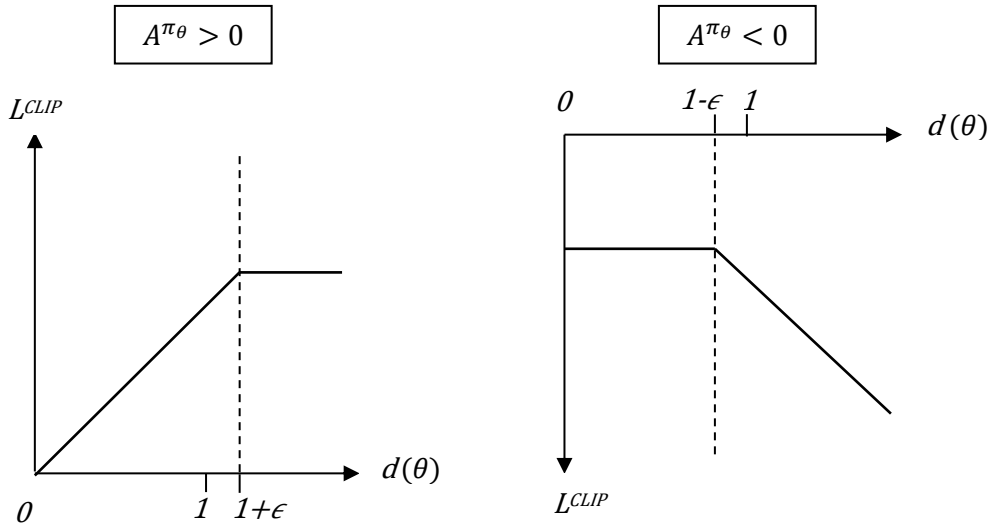


Figure 8.1 Surrogate objective with policy ratio clipping in Proximal Policy Optimization

To practically use it, we have to do gradient ascent with respect to θ' on equation 8.2, using an automatic differentiation library that supports clipping (e.g. PyTorch).

Since the gradient is with respect to θ' , when using automatic differentiation libraries it is necessary to first compute $\pi_\theta(a|s)$ without accumulating gradients, and only accumulating gradients when computing $\pi_{\theta'}(a|s)$.

This is necessary when using the same policy neural network both with the old parameters θ to compute the policy actions $\pi_\theta(a|s)$, and with the new parameters θ' to compute the updated policy $\pi_{\theta'}(a|s)$, because when computing the gradient wrt θ' it is not correct to accumulate also gradients related to previous operations with old parameters. In the surrogate objective $L_\theta^{CLIP}(\theta')$ we consider $\pi_\theta(a|s)$ as a constant (the unknown is θ' , not θ).

8.3 Multiple Steps of Gradient Ascent

When doing the policy gradient ascent round with a batch of trajectories, in PPO we do not do only one step of gradient ascent: we do an iteration of many steps of gradient ascent while the Kullback-Leibler divergence of the old policy from the new policy stays under a certain threshold (see next section 8.4 for the KL divergence check). During all those steps of gradient ascent the batch of trajectory is retained and used at each step, as well as all related computations of $\pi_\theta(a|s)$. What changes at each step is that the parameters θ' will be updated so at each new step it is necessary to compute new values for $\pi_{\theta'}(a|s)$ and new gradients of $L_\theta^{CLIP}(\theta')$ with respect to θ' . Note that in the first step of the iteration the parameters θ' have the same value as previous parameters θ . Once the Kullback-Leibler divergence reveals that the new policy would be too different from the old, the last update is discarded, the gradient ascent iteration ends, and a new batch of trajectories is generated, to feed a new round of policy gradient ascent. The algorithm 8.1 will clarify this.

8.4 Kullback-Leibler divergence check

This clipping does not guarantee that the new policy is not too different from the old policy, so it is necessary to check the Kullback-Leibler divergence of the old policy from the new policy, using the taken actions and computing their probability both under the old policy and under the new policy. Then stop the policy gradient ascent iterations in case the KL divergence is over a certain threshold (“early stopping”). This avoids using samples generated from a policy

that has a distribution too different from the one that gets optimized. The KL divergence computed in this way is just an approximation because is the average of the KL divergence for each sample. Since it is sampled on the distribution of actions from the old policy, it is already the empirical expectation over the old policy (so there is no need to multiply for $\pi_\theta(a_t|s_t)$ in the divergence formula):

$$\widehat{KL} = \frac{1}{T} \sum_t^T \log(\pi_\theta(a_t|s_t)) - \log(\pi_{\theta'}(a_t|s_t)) \quad (8.3)$$

Until now we used π_θ to name the old policy (parametrized by θ) and $\pi_{\theta'}$ to name the new one (parametrized by θ'). For clarity's sake in the pseudocode algorithm, we are going to use only one policy neural network which is parametrized by θ_k (with k the number of the optimization round). Hence the old policy π_θ is now π_{θ_k} and the new policy $\pi_{\theta'}$ is now $\pi_{\theta_{k+1}}$.

Algorithm 8.1 Proximal Policy Optimization with ratio clipping and early stopping

Require: Policy network step size η

Require: Value network step size ω

Require: Threshold for Kullback-Leibler divergence based early stopping ζ

Require: Initialize parameters θ of network π_θ with small random values

Require: Initialize parameters ψ of network V_ψ with small random values

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using policy π_{θ_k}

 compute rewards-to-go $G(\tau_{i,t})$, for each t of each τ_i

 compute advantage estimates \hat{A}_t using current estimate of value function V_{ψ_k} :

$$\hat{A}_t = G(\tau_t) - V_{\psi_k}(s_t)$$

 (alternatively we can use GAE):

$$\hat{A}_t^{GAE(\gamma, \lambda)} = \sum_{l=t}^{|\tau|} (\lambda \gamma)^{l-t} \delta_l^\gamma \text{ with } \delta_l^\gamma = r_l + \gamma V_{\psi_k}(s_{l+1}) - V_{\psi_k}(s_l)$$

 compute $\pi_{\theta_k}(a_t|s_t)$ without accumulating gradients

$$\theta_{k+1} \leftarrow \theta_k$$

 For $m = 0, 1, 2, \dots, M-1$ do a policy gradient ascent iteration:

 compute policy ratios:

$$d_t(\theta) = \frac{\pi_{\theta_{k+1}}(a_t|s_t)}{\pi_{\theta_k}(a_t|s_t)}$$

compute clipped surrogate losses:

$$L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_t = \min (d_t(\theta) \hat{A}_t, \quad clip(d_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t)$$

compute KL divergence of $\pi_{\theta_k}(a_t|s_t)$ from $\pi_{\theta_{k+1}}(a_t|s_t)$ on taken actions:

$$\widehat{KL} = \frac{1}{|D_k|} \sum_{\tau \in D_k} \frac{1}{T} \sum_{t=0}^{T-1} \log (\pi_{\theta_k}(a_t|s_t)) - \log (\pi_{\theta_{k+1}}(a_t|s_t))$$

If $\widehat{KL} \geq \zeta$:

stop doing policy gradient ascent and exit this internal For-cycle

end of If

estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \nabla_{\theta_{k+1}} L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_t$$

update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_{k+1} + \eta \widehat{g}_k$$

end of for

For $n = 0, 1, 2, \dots, N - 1$ do a value function gradient descent iteration:

estimate the value function gradients as:

$$\widehat{h}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \frac{1}{T} \sum_{t=0}^{T-1} \nabla_{\psi_k} (V_{\psi_k}(s_t) - G(\tau_t))^2$$

update the value function with a gradient descent step (or other method):

$$\psi_k \leftarrow \psi_k + \omega \widehat{h}_k$$

end of for

$$\psi_{k+1} \leftarrow \psi_k$$

end of for

PPO is considered on-policy because it aims at optimizing the same policy function that generates the training trajectories (and it is not valid to use an arbitrary policy function to generate trajectories). But actually, it is also *slightly* off-policy because it does possibly more than one step of optimization with the same trajectories, and at each step further than the first it is optimizing a policy that is not *exactly* the same that generated trajectory: it is very similar (small KL-divergence) but not the same.

Discussions on implementation details can be found in [Engstrom et al. 2020], [Andrychowicz et al. 2020], [Huang et al. 2022].

9. Deep Deterministic Policy Gradient

9.1 Base Idea

Deterministic Policy Gradient is an algorithm invented by [Silver et al. 2014] (with similarities to [Hafner and Riedmiller 2011]), and then adapted to deep networks by [Lillicrap et al. 2016], that allows us to do an effective off-policy policy gradient optimization for tasks with continuous actions. The basic ideas are (a) to have a target policy that is deterministic, which means that for each state the policy chooses only one well-determined action, and (b) to improve the policy greedily, that here is done by moving the policy in the direction of the gradient of $Q^\pi(s, a)$. This is considered to be a policy gradient algorithm that is similar in spirit to Q-Learning because of the greediness of learning and determinism of the target policy. The full algorithm is off-policy, but we will present an on-policy version at first, before describing the changes to make it off-policy.

Let us digress a little to prepare field for the full picture.

If you recall, the “*Expected return*” as used until now is the one expressed in eq. 1.5 , in section 1.2 :

$$J(\pi) = E_{\tau \sim \pi}[G(\tau)] = E_{a \sim \pi, s_0 \sim \rho_0, s \sim P}[G(\tau)]$$

where $G(\tau)$ is the total return $G(\tau) = \sum_{t=0}^{\infty} \gamma^t r_t$.

In that formula, $J(\pi)$ is an expectation over the trajectory τ , that means that it is actually an expectation over actions and over states, with actions a distributed by policy π , initial state s_0 distributed by ρ_0 , and the other states s distributed by the transition function P . Until now, the policy π has been considered stochastic: we selected an action basing on a random sample that followed that distribution. For discrete actions that means that every possible discrete action occupies a subset of $[0,1]$, right next to the subset of the previous action, as big as the probability assigned by the policy function. Then a random number between 0 and 1 is sampled, which permits us to select one of those subsets and hence the corresponding action. For continuous actions instead, for each action dimension the policy π gives the average of the value and the variance (which may or not be fixed, a priori), and those parameters are used within a certain distribution (usually a normal distribution) to sample each action dimension.

This kind of stochastic behaviour is responsible for some variety of exploration, and for a minimum level of robustness in adversarial games (but beware: for some optimal policy the level of exploration and variety of action may be very little).

It is also the reason why the Expected Return $J(\pi)$ is an expectation also over actions and not only over states. We could make the computation of its gradient simpler if it was only over states. How to do it ? We could decide that our target policy π is deterministic instead of stochastic, and hence instead of sampling the action, we just choose the action with maximum probability. So, we would not have any expectation over actions anymore. As we will see below, this will also allow us to derive a policy gradient ascent that is “greedy”: in the direction of the gradient of $Q^\pi(s, a)$. That is the basic idea of Deterministic Policy Gradient.

Now, for discrete actions there is already an algorithm that includes a similar idea: Q-Learning (see Ch.2), which is not a policy based algorithm but essentially applies the same idea of having the target policy choosing only the action with maximum probability (which in Q-Learning is the action with the maximum action-state value, since Q-Learning is a value-based algorithm), hence being both deterministic and greedy at the same time.

As we already saw, applying Q-Learning to continuous actions instead is complicated, since it would entail finding the maximum of a function approximator (the one used for the action-value function) which is very problematic to do if the function approximator is a deep neural network. But it is still possible to apply the idea of deterministic policy to policy-gradient systems with continuous actions: instead of sampling the action using the average and the variance given by the policy function, we can just choose always that average as action. This will make the expectation over the actions collapse to just one action, and simplify the computation of the gradient of $J(\pi)$. A second benefit of having a deterministic policy is that it is applicable to environments where stochastic policies cannot be used, such as in robots with a differentiable control policy where it is not possible to inject noise.

Since a deterministic policy does not offer any level of exploration per se, a different exploratory policy (“*behavior policy*”) must be used which guarantees some variety of actions, so the whole algorithm will be off-policy.

9.2 Action-Value Gradient

Let us rewrite the expected return $J(\pi)$ in a way that can be handy with deterministic policy. The first thing to note is that the expectation of $\sum_{t=0}^{\infty} \gamma^t r_t$, over actions a distributed by π and

states s distributed by ρ_0 (if initial) and P , is equal to the expectation of $Q^\pi(s, a)$, as already showed in eq. 1.8.

$$E_{\tau \sim \pi}[G(\tau)] = E_{\tau \sim \pi}[Q^\pi(s, a)] = E_{a \sim \pi, s_0 \sim \rho_0, s \sim P}[Q^\pi(s, a)] \quad (9.1)$$

Then, analogously to what we did in section 7.2, we can define the probability of being in state s at time t , depending on policy π as $Prob^\pi(s_t = s)$.

And we can define the frequency $\xi_\pi(s)$ of a state s as the number of times that s is expected to be visited, computed as unnormalized (it is a frequency, not a probability) and time-discounted, under the policy π :

$$\xi_\pi(s) = Prob^\pi(s_0 = s) + \gamma Prob^\pi(s_1 = s) + \gamma^2 Prob^\pi(s_2 = s) + \dots \quad (9.2)$$

To be noted that $\xi_\pi(s)$ it is not a probability distribution, but it is a time-discounted frequency (where γ is the discount factor), even if in the original paper [Silver et al. 2014] that was defined “the (improper) discounted state distribution” and was denoted as $\rho^\pi(s')$.

So, the expected return $J(\pi)$ can be rewritten as:

$$\begin{aligned} J(\pi) &= \sum_t \sum_s Prob^\pi(s_t = s) \sum_a \pi(a_t | s_t) \gamma^t Q^\pi(s_t, a_t) \\ J(\pi) &= \sum_s \sum_t \gamma^t Prob^\pi(s_t = s) \sum_a \pi(a_t | s_t) Q^\pi(s_t, a_t) \\ J(\pi) &= \sum_s \xi_\pi(s) \sum_a \pi(a_t | s_t) Q^\pi(s_t, a_t) \end{aligned} \quad (9.3)$$

Which can be written more generally also with integral notation as:

$$J(\pi) = E_{\tau \sim \pi}[Q^\pi(s, a)] = \int_s \xi_\pi(s) \int_a \pi(a | s) Q^\pi(s, a) \quad (9.4)$$

At this point if the policy π is deterministic, we have that the integral over actions disappears because we have only one possible action a to choose from for each state, or in other terms the policy tells us that the probability of choosing that action a is $= 1$. We can name μ_θ the vector of outputs by the policy neural network which computes the averages for each action dimension, parametrized by weights θ . We can also make explicit that the policy π depends on θ , writing it as π_θ , and writing $J(\mu_\theta)$ instead of $J(\pi)$. The expected return then is:

$$J(\mu_\theta) = \int_s \xi_{\pi_\theta}(s) Q^{\mu_\theta}(s, \mu_\theta(s)) \quad (9.5)$$

Now, in that form it is evident that we can do gradient ascent on $J(\mu_\theta)$ with respect to θ , sampling the states that appear while the agent is following policy π (those states will appear with discounted frequency $\xi_\pi(s)$).

$$\nabla_\theta J(\mu_\theta) = \nabla_\theta \int_s \xi_{\pi_\theta}(s) Q^{\mu_\theta}(s, \mu_\theta(s)) \quad (9.6)$$

Naming η the step size of parameters' updates, the update rule is:

$$\theta' \leftarrow \theta + \eta \nabla_\theta \int_s \xi_{\pi_\theta}(s) Q^{\mu_\theta}(s, \mu_\theta(s)) \quad (9.7)$$

Applying the chain rule of calculus, we find the gradient composed by the gradient of μ_θ wrt θ , multiplied by the gradient of $Q^{\mu_\theta}(s, a)$, with respect to a , having a valued at point $a = \mu_\theta(s)$, that is the a chosen by the deterministic policy.

$$\theta' \leftarrow \theta + \eta \nabla_\theta \int_s \xi_{\pi_\theta}(s) \nabla_\theta \mu_\theta(s) \cdot \nabla_a Q^{\mu_\theta}(s, a) \Big|_{a = \mu_\theta(s)} \quad (9.8)$$

Now, when we sample a process where a state has (time-discounted) frequency $\xi_{\pi_\theta}(s)$, it means that we are sampling a state which is distributed by the probability $Prob^{\pi_\theta}(s_t = s)$.

That means:

$$\nabla_{\theta} \int_s \xi_{\pi_{\theta}}(s) Q^{\mu_{\theta}}(s, \mu_{\theta}(s)) \propto E_{s \sim \text{Prob}^{\mu_{\theta}}} \left[\nabla_{\theta} \mu_{\theta}(s) \cdot \nabla_a Q^{\mu_{\theta}}(s, a) \Big|_{a = \mu_{\theta}(s)} \right] \quad (9.9)$$

With a little abuse of notation, we can say that from eq. 9.7 we can deduce eq. 9.10 below (the abuse of notation consists in the fact that η of eq. 9.7 is not the same η of eq. 9.10).

$$\theta' \leftarrow \theta + \eta E_{s \sim \text{Prob}^{\mu_{\theta}}} \left[\nabla_{\theta} \mu_{\theta}(s) \cdot \nabla_a Q^{\mu_{\theta}}(s, a) \Big|_{a = \mu_{\theta}(s)} \right] \quad (9.10)$$

To be noted that, following original paper [Silver et al. 2014] description, $\nabla_{\theta} \mu_{\theta}(s)$ is a transposed Jacobian matrix where each column d is the gradient $\nabla_{\theta} [\mu_{\theta}(s)]_d$ of the d th action dimension with respect to the policy parameters θ . It becomes clear then that $\nabla_a Q^{\mu_{\theta}}(s, a)$ is a column vector, and both θ and θ' are column vectors as well. (In usual notation, $\nabla_{\theta} \mu_{\theta}(s)$ would have made the reader think that it was a Jacobian matrix, and not a transposed Jacobian matrix, but original authors specified the orientation of columns and rows, and following it makes it a transposed Jacobian).

Eq. 9.9 shows that it is irrelevant the fact that changing the policy μ_{θ} by the gradient ascent step in eq. 9.10 could change in turn the distribution of the states: the gradient of the expected return does not depend on the gradient of the state distribution (which we do not know), similarly to what we saw with stochastic on-policy policy gradient in section 3.2. In other words: we are already optimizing the performance $J(\mu_{\theta})$, hence the optimization step in 9.10 will indeed improve the policy and we do not need to care about change in state distribution. Original authors [Silver et al. 2014, supplement material] use a slightly different proof that originates as a deterministic version of Policy Gradient Theorem.

One remarkable thing is that the deterministic policy gradient represents the limit of the stochastic policy gradient, when the variance of the stochastic policy tends to zero, for a family of stochastic policies. This is important because it means that many techniques applicable to (stochastic) policy gradient algorithms may be applied also to deterministic policy gradient. We are not writing the proof here because it is quite long, and it is already written in a detailed and clear way in the supplementary material of [Silver et al. 2014]. We will just outline the theorem definition.

Suppose that $\mu_\theta(s)$ is a deterministic policy function, parametrized by θ , which outputs a vector of deterministic action dimensions (i.e. $\mu_\theta(s)$ has the same meaning as it had above in this chapter). Consider a stochastic policy function $\pi_{\mu_\theta, \sigma}(a|s) = v_\sigma(\mu_\theta(s), a)$ satisfying some conditions called “regular delta-approximation” (detailed in [Silver et. al 2014]), where σ is the variance. An example of such a stochastic policy could be a gaussian policy function with mean $\mu_\theta(s)$ and variance σ .

Then:

$$\lim_{\sigma \rightarrow 0^+} \nabla_\theta J(\pi_{\mu_\theta, \sigma}) = \nabla_\theta J(\mu_\theta) \quad (9.11)$$

On the left-hand side there is the gradient of the stochastic policy $\pi_{\mu_\theta, \sigma}(a|s)$, and on the right-hand side the gradient of the deterministic policy $\mu_\theta(s)$, so as we anticipated it means that the deterministic policy gradient is the limit of the stochastic policy gradient when the variance of the stochastic policy tends to zero.

9.3 On-Policy Deterministic Actor-Critic

Now, we have all elements to start build a learning algorithm.

If the agent follows the same policy that we optimize, that means that we are doing on-policy optimization, and we can use an approximation function parametrized by ψ as critic $Q_\psi(s, a)$ to estimate $Q^{\mu_\theta}(s, \mu_\theta(s))$ with a Sarsa algorithm (see section 2.10). This on-policy version is slower than the off-policy one because it needs to spend more time generating trajectories, and it is meant only as a preliminary example to later introduce the off-line algorithm.

In the following pseudo-code algorithms, the variable *is_terminal* is a Boolean which indicates if reached state $s_{m, t+1}$ is a terminal state, i.e. the trajectory from which the transition is sampled ends there.

Algorithm 9.1 On-Policy Deterministic Actor-Critic

Require: Policy network step size η

Require: Q-Value network step size ω

Require: Initialize parameters θ of network μ_{θ_k} with small random values

Require: Initialize parameters ψ of network Q_ψ with small random values

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using policy μ_{θ_k}

 For $n = 0, 1, 2, \dots, N - 1$ do a q-value function gradient descent iteration:

 For $m = 0, 1, 2, \dots, M - 1$ do:

 from D_k sample a transition $\langle s_{m,t}, a_{m,t}, s_{m,t+1}, r_{m,t}, is_terminal \rangle$

 if not *is_terminal* then:

 compute $a_{m,t+1} = \mu_{\theta_k}(s_{m,t+1})$

$y_{m,t} = r_{m,t} + \gamma Q_\psi(s_{m,t+1}, a_{m,t+1})$

 else:

$y_{m,t} = r_{m,t}$

 endif

 End of for

 estimate the q-value network gradients as:

$$\widehat{h}_k = \frac{1}{M} \sum_{m=0}^{M-1} (y_{m,t} - Q_\psi(s_{m,t}, a_{m,t})) \nabla_\psi Q_\psi(s_{m,t}, a_{m,t})$$

 update the q-value network with a gradient descent step (or other method):

$$\psi \leftarrow \psi + \omega \widehat{h}_k$$

 end of for

 estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \frac{1}{T} \sum_{t=0}^{T-1} \nabla_\theta \mu_\theta(s_t) \cdot \nabla_a Q_\psi(s_t, a) \Big|_{a = \mu_{\theta_k}(s_t)}$$

 update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g}_k$$

end of for

9.4 Off-Policy Deterministic Actor-Critic

If instead we want to build an off-policy algorithm, we have a behavior policy parametrized by β , so under that policy we denote $\xi_\beta(s)$ the time-discounted frequency of a state s . The performance to be optimized, following [Silver et al. 2014], then is the total expected value where the frequency of the visited states is the one under the behavior policy, while the q-value of each state is computed as if we followed the target policy:

$$J_\beta(\mu_\theta) = \int_s \xi_\beta(s) Q^{\pi_\theta}(s, \mu_\theta(s)) \quad (9.11)$$

We should be aware that this goal performance is not theoretically justified, because analogously as we discussed in section 4.3 about off-policy policy gradients we know that when using function approximators:

$$J_\beta(\mu_{\theta'}) \geq J_\beta(\mu_\theta) \text{ does not imply } J(\mu_{\theta'}) \geq J(\mu_\theta) \quad (9.12)$$

As with the on-policy algorithm, we can use an approximation function parametrized by ψ as critic $Q_\psi(s, a)$ to estimate $Q^{\pi_\theta}(s, \mu_\theta(s))$ with a Sarsa algorithm (see section 2.10). The algorithm then is very similar to On-Policy Deterministic Actor-Critic (algorithm 9.1), with the difference that the trajectories are generated following the behavior policy and not the target policy. Since the policy is deterministic, this Sarsa evaluation coincides with an off-policy Expected Sarsa evaluation.

Algorithm 9.2 Off-Policy Deterministic Actor-Critic

Require: Policy network step size η

Require: Q-Value network step size ω

Require: Behavior policy β

Require: an empty replay buffer B

Require: Initialize parameters θ of network μ_{θ_k} with small random values

Require: Initialize parameters ψ of network Q_ψ with small random values

For $k = 0, 1, 2, \dots$ do:

collect trajectories $D_k = \{ \tau_i \}$ using behavior policy β

put all transitions $\langle s, a, s', r \rangle$ of all trajectories in replay buffer B

For $n = 0, 1, 2, \dots, N - 1$ do a q-value function gradient descent iteration:

For $m = 0, 1, 2, \dots, M - 1$ do:

from buffer B sample a transition $\langle s_{m,t}, a_{m,t}, s_{m,t+1}, r_{m,t}, is_terminal \rangle$

if not $is_terminal$ then:

compute $a_{m,t+1} = \mu_{\theta_k}(s_{m,t+1})$

$y_{m,t} = r_{m,t} + \gamma Q_{\psi}(s_{m,t+1}, a_{m,t+1})$

else:

$y_{m,t} = r_{m,t}$

endif

End of for

estimate the q-value network gradients as:

$$\widehat{h}_k = \frac{1}{M} \sum_{m=0}^{M-1} (y_{m,t} - Q_{\psi}(s_{m,t}, a_{m,t})) \nabla_{\psi} Q_{\psi}(s_{m,t}, a_{m,t})$$

update the q-value network with a gradient descent step (or other method):

$$\psi \leftarrow \psi + \omega \widehat{h}_k$$

end of for

estimate policy gradient as:

From replay buffer B sample U transitions (you need only the state s_u)

$$\widehat{g}_k = \frac{1}{U} \sum_{u=0}^{U-1} \nabla_{\theta} \mu_{\theta}(s_u) \cdot \nabla_a Q_{\psi}(s_u, a) \Big|_{a = \mu_{\theta_k}(s_u)}$$

update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g}_k$$

end of for

9.5 Compatible Function Approximator

The gradient of the approximated q-value function is used in the policy gradient update formula instead of the real, unknown, q-value gradient, and this implies that the resulting policy gradient may not follow the true gradient (in the off-policy algorithm this is even worsened by the presence of the three conditions “function approximation”, “bootstrapping”, “off-policy training” that may lead the function approximator to diverge instead of converge, as discussed in section 2.4).

For this reason, we should find a way to recognize if a certain function approximator will have an approximated q-value gradient that follows the true q-value gradient. The original authors call it “Compatible Function Approximator”, and discover that there is a class of such approximators characterized by two well-defined properties enunciated below in the theorem. Please note that the following theorem is valid both for on-policy and off-policy deterministic actor critic, and that with the symbol E we denote both the expectation $E_{s \sim \text{Prob} \mu_\theta}$ over states distributed by the target policy μ_θ (on-policy case) and the expectation $E_{s \sim \text{Prob} \beta}$ over states distributed by the behavior policy β (off-policy case). In the same way, in the theorem when we write the performance gradient $\nabla_\theta J(\mu_\theta)$, we mean both the on-policy gradient (properly: $\nabla_\theta J(\mu_\theta)$) and the off-policy gradient (it would be properly denoted as $J_\beta(\mu_\theta)$ as previously seen).

Compatible function approximation theorem: saying that a function approximator $Q_\psi(s, a)$ parametrized by ψ is *compatible* with a deterministic policy $\mu_\theta(s)$ means that:

$$\nabla_\theta J(\mu_\theta) = E \left[\nabla_\theta \mu_\theta(s) \cdot \nabla_a Q_\psi(s, a) \Big|_{a = \mu_\theta(s)} \right] \quad (9.12)$$

Two conditions imply the compatibility:

$$1. \quad \nabla_a Q_\psi(s, a) \Big|_{a = \mu_\theta(s)} = (\nabla_\theta \mu_\theta(s))^T \psi$$

$$\text{where } \psi \text{ is a column vector of same length of } \theta \quad (9.13)$$

$$2. \quad \psi \text{ minimizes the mean - squared error } MSE(\theta, \psi) = E[\epsilon(s, \theta, \psi)^T \epsilon(s, \theta, \psi)] \quad (9.14)$$

$$\text{where } \epsilon(s, \theta, \psi) = \left. \nabla_a Q_\psi(s, a) \right|_{a = \mu_\theta(s)} - \left. \nabla_a Q^{\mu_\theta}(s, a) \right|_{a = \mu_\theta(s)} \quad (9.15)$$

Proof:

By condition 2 (eq. 9.14), ψ minimizes $MSE(\theta, \psi)$, then the gradient of ϵ^2 w.r.t. ψ must be zero.

$$\nabla_\psi MSE(\theta, \psi) = 0$$

$$E[\nabla_\psi (\epsilon(s, \theta, \psi)^T \epsilon(s, \theta, \psi))] = 0$$

applying derivative product rule:

$$E\left[\left(\nabla_\psi \epsilon(s, \theta, \psi)\right)^T \epsilon(s, \theta, \psi) + \left(\epsilon(s, \theta, \psi)^T \nabla_\psi \epsilon(s, \theta, \psi)\right)^T\right] = 0$$

$$E\left[2\left(\nabla_\psi \epsilon(s, \theta, \psi)\right)^T \epsilon(s, \theta, \psi)\right] = 0$$

$$E\left[\left(\nabla_\psi \epsilon(s, \theta, \psi)\right)^T \epsilon(s, \theta, \psi)\right] = 0$$

(9.16)

By condition 1 (eq. 9.13), we find that $\nabla_\psi \epsilon(s, \theta, \psi) = \nabla_\theta \mu_\theta(s)$. Let us show it starting from eq. 9.15

$$\epsilon(s, \theta, \psi) = \left. \nabla_a Q_\psi(s, a) \right|_{a = \mu_\theta(s)} - \left. \nabla_a Q^{\mu_\theta}(s, a) \right|_{a = \mu_\theta(s)}$$

$$\nabla_\psi \epsilon(s, \theta, \psi) = \nabla_\psi \left(\left. \nabla_a Q_\psi(s, a) \right|_{a = \mu_\theta(s)} - \left. \nabla_a Q^{\mu_\theta}(s, a) \right|_{a = \mu_\theta(s)} \right)$$

$$\nabla_\psi \epsilon(s, \theta, \psi) = \nabla_\psi \left. \nabla_a Q_\psi(s, a) \right|_{a = \mu_\theta(s)}$$

plugging-in eq. 9.13:

$$\nabla_{\psi} \epsilon(s, \theta, \psi) = \nabla_{\psi} \left((\nabla_{\theta} \mu_{\theta}(s))^T \psi \right)$$

$$\nabla_{\psi} \epsilon(s, \theta, \psi) = \nabla_{\theta} \mu_{\theta}(s)$$

(9.17)

plugging eq. 9.17 into eq. 9.16:

$$E \left[(\nabla_{\theta} \mu_{\theta}(s))^T \epsilon(s, \theta, \psi) \right] = 0$$

plugging-in eq. 9.15:

$$E \left[(\nabla_{\theta} \mu_{\theta}(s))^T \left(\nabla_a Q_{\psi}(s, a) \Big|_{a = \mu_{\theta}(s)} - \nabla_a Q^{\mu_{\theta}}(s, a) \Big|_{a = \mu_{\theta}(s)} \right) \right] = 0$$

$$E \left[(\nabla_{\theta} \mu_{\theta}(s))^T \nabla_a Q_{\psi}(s, a) \Big|_{a = \mu_{\theta}(s)} \right] = E \left[(\nabla_{\theta} \mu_{\theta}(s))^T \nabla_a Q^{\mu_{\theta}}(s, a) \Big|_{a = \mu_{\theta}(s)} \right]$$

$$E \left[(\nabla_{\theta} \mu_{\theta}(s))^T \nabla_a Q_{\psi}(s, a) \Big|_{a = \mu_{\theta}(s)} \right] = \nabla_{\theta} J(\mu_{\theta})$$

(9.18)

Which proves the theorem (statement 9.12).

■

It turns out that for any policy $\mu_{\theta}(s)$ it is possible to find a function approximator that satisfies the condition 1 (eq. 9.13) of compatible function approximators, with the form:

$$Q_{\psi}(s, a) = (a - \mu_{\theta}(s))^T (\nabla_{\theta} \mu_{\theta}(s))^T \psi + V_v(s)$$

(9.18)

Where $V_v(s)$ is any differentiable baseline independent from the action a , parametrized by v . A way to see this is to think of $V_v(s)$ as a value function, for instance a linear one, where, if called the state features as $\phi(s)$, we can have $V_v(s) = v^T \phi(s)$.

Then, if we consider that the q-value function is equivalent to the state value function plus the advantage function (eq. 1.13), we can think of $(a - \mu_{\theta}(s))^T (\nabla_{\theta} \mu_{\theta}(s))^T \psi$ as a linear advantage

function. Such advantage function would have state-action features $\phi(s, a) = (\nabla_{\theta} \mu_{\theta}(s))(a - \mu_{\theta}(s))$ and parameters ψ , so to have $A^{\psi}(s, a) = \phi(s, a)^T \psi$.

We can also denote the deviation of the action from current policy as $\Delta_a = a - \mu_{\theta}(s)$, in this way the advantage function may be written as $A^{\psi}(s, \mu_{\theta}(s) + \Delta_a) = \Delta_a^T \psi$.

To indicate that $Q_{\psi}(s, a)$ is now parametrized not only by ψ but also by v , because it is the sum of the advantage function A^{ψ} and the value function V_v , we can write it as $Q_{\psi, v}(s, a)$.

Linear approximation is not very precise to predict action-value globally (it can diverge towards infinity), but it can work well as local critic. Here it is used in an advantage function which indicates the local advantage with respect to the current policy: a linear approximator may be enough to find the right direction in which the policy parameters should be updated.

Let us examine the second condition (eq. 9.14 and 9.15): ψ should minimize the mean squared error between the approximated gradient of Q_{ψ} and the true gradient of $Q^{\mu_{\theta}}$. To fulfil this condition with a linear approximator, we should build a linear regression to approximate

$\nabla_a Q^{\mu_{\theta}}(s, a) \Big|_{a = \mu_{\theta}(s)}$, with $\phi(s, a)$ as features, and ψ as parameters. But we are not doing exactly that, because we are learning parameters ψ through Sarsa/Q-Learning evaluation, so ψ is not guaranteed to actually minimize $MSE(\theta, \psi)$, hence the condition 2 is not satisfied (also, it is hard to obtain unbiased samples of the real gradient of $Q^{\mu_{\theta}}$). We can only say that despite not having the condition 2 satisfied, in practice our policy evaluation may be good enough to make Q_{ψ} an acceptable approximation.

This algorithm with linear approximation and quasi-compliance with compatible function approximation is named “Compatible Off-Policy Deterministic Actor-Critic”, COPDAC in short.

Algorithm 9.3 Compatible Off-Policy Deterministic Actor-Critic

Require: Policy network step size η

Require: Advantage linear approximator step size ω

Require: Value linear approximator step size ς

Require: Behavior policy β

Require: an empty replay buffer B

Require: Init parameters θ of Policy network μ_{θ_k} with small random values

Require: Init parameters ψ of Advantage linear approximator A_ψ with small random values

Require: Init parameters v of Value linear approximator V_v with small random values

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using behavior policy β

 put all transitions $\langle s, a, s', r, is_terminal \rangle$ of all trajectories D_k in replay buffer B

 For $n = 0, 1, 2, \dots, N - 1$ do a minibatch iteration:

 For $m = 0, 1, 2, \dots, M - 1$ do:

 From buffer B sample a transition $\langle s_{m,t}, a_{m,t}, s_{m,t+1}, r_{m,t}, is_terminal \rangle$

$$\phi(s_{m,t}, a_{m,t}) = \left(\nabla_{\theta} \mu_{\theta_k}(s_{m,t}) \right) \left(a_{m,t} - \mu_{\theta_k}(s_{m,t}) \right)$$

$$A^{\psi}(s_{m,t}, a_{m,t}) = \phi(s_{m,t}, a_{m,t})^T \psi$$

 extract state features $\phi(s_{m,t})$

$$V_v(s_{m,t}) = v^T \phi(s_{m,t})$$

$$Q_{\psi,v}(s_{m,t}, a_{m,t}) = A^{\psi}(s_{m,t}, a_{m,t}) + V_v(s_{m,t})$$

 if not $is_terminal$ then:

$$\text{compute } a_{m,t+1} = \mu_{\theta_k}(s_{m,t+1})$$

$$\phi(s_{m,t+1}, a_{m,t+1}) = \left(\nabla_{\theta} \mu_{\theta_k}(s_{m,t+1}) \right) \left(a_{m,t+1} - \mu_{\theta_k}(s_{m,t+1}) \right)$$

$$A^{\psi}(s_{m,t+1}, a_{m,t+1}) = \phi(s_{m,t+1}, a_{m,t+1})^T \psi$$

 extract state features $\phi(s_{m,t+1})$

$$V_v(s_{m,t+1}) = v^T \phi(s_{m,t+1})$$

$$Q_{\psi,v}(s_{m,t+1}, a_{m,t+1}) = A^{\psi}(s_{m,t+1}, a_{m,t+1}) + V_v(s_{m,t+1})$$

$$y_{m,t} = r_{m,t} + \gamma Q_{\psi,v}(s_{m,t+1}, a_{m,t+1})$$

 else:

$$y_{m,t} = r_{m,t}$$

 endif

$$\delta_{m,t} = y_{m,t} - Q_{\psi}(s_{m,t}, a_{m,t})$$

End of for

estimate the advantage and value functions gradients as:

$$\widehat{h}_k = \frac{1}{M} \sum_{m=0}^{M-1} \delta_{m,t} \phi(s_{m,t}, a_{m,t})$$

$$\widehat{d}_k = \frac{1}{M} \sum_{m=0}^{M-1} \delta_{m,t} \phi(s_{m,t})$$

update advantage and value functions by gradient descent step (or other):

$$\psi \leftarrow \psi + \omega \widehat{h}_k$$

$$v \leftarrow v + \varsigma \widehat{d}_k$$

end of for

estimate policy gradient as:

From replay buffer B sample U transitions (you only need the state s_u)

$$\widehat{g}_k = \frac{1}{U} \sum_{u=0}^{U-1} \nabla_{\theta} \mu_{\theta_k}(s_u) \cdot \left(\nabla_{\theta} \mu_{\theta_k}(s_u) \right)^T \psi$$

update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g}_k$$

end of for

9.6 Deep Deterministic Policy Gradient

Compatible Function Approximator provides a way to have a q-value gradient closer to the true one, but at the expense of having only a linear approximator of the q-value function. If we want to have a richer and more complex capacity of processing the input state, we need to use a large neural network as function approximator for the q-value, but this is known to be unstable, which precisely motivated the usage of COPDAC.

[Lillicrap et al. 2016] found an effective way to use big neural networks as function approximators for the q-value, using the following expedients:

- Using a replay buffer (see section 2.3). We already did this in previous DPG algorithms.
- Target Network for Q-Value function (see section 2.5) with Polyak averaging (see section 2.6). This means to use a different Q-Value network to compute target Q-

Values, which gets updated slowly towards the real Q-Value network, in order to give more stable results.

- Target policy network with Polyak averaging, to be used to compute $a_{t+1} = \mu_{\theta}(s_{t+1})$ in the target Q-Value computation.
- Batch normalization [Ioffe and Szegedy 2015]
- The Behavior policy is the Target policy plus gaussian noise on actions (i.e. is a stochastic gaussian policy, whose average is the deterministic target policy). In the original paper [Lillicrap et al. 2016], authors used a Ornstein-Uhlenbeck process to generate the noise, but gaussian noise seems to work even better [Achiam & OpenAI 2020E].

Algorithm 9.4 Deep Deterministic Policy Gradient (Actor-Critic)

Require: Policy network step size η

Require: Q-Value network step size ω

Require: Behavior policy β , usually $= \mu_{\theta_k}(s) + \text{Gaussian}$

Require: an empty replay buffer B

Require: Initialize parameters θ of network μ_{θ_k} with small random values

Require: Initialize parameters ψ of network Q_{ψ} with small random values

Require: Polyak Q-Value Step size ς

Require: Polyak Policy Step size ζ

copy Q-Value target network parameters $\psi' \leftarrow \psi$

copy Policy target network parameters $\theta' \leftarrow \theta$

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using behavior policy β

 put all transitions $\langle s, a, s', r, is_terminal \rangle$ of all trajectories D_k in replay buffer B

 For $n = 0, 1, 2, \dots, N - 1$ do a minibatch iteration:

 For $m = 0, 1, 2, \dots, M - 1$ do:

 From buffer B sample a transition $\langle s_{m,t}, a_{m,t}, s_{m,t+1}, r_{m,t}, is_terminal \rangle$

 if not $is_terminal$ then:

 compute $a_{m,t+1} = \mu_{\theta'_{k+1}}(s_{m,t+1})$

$y_{m,t} = r_{m,t} + \gamma Q_{\psi'}(s_{m,t+1}, a_{m,t+1})$

 else:

$y_{m,t} = r_{m,t}$

 endif

end of for

estimate the q-value network gradients as:

$$\widehat{h}_k = \frac{1}{M} \sum_{m=0}^{M-1} (y_{m,t} - Q_{\psi}(s_{m,t}, a_{m,t})) \nabla_{\psi} Q_{\psi}(s_{m,t}, a_{m,t})$$

update the q-value network with a gradient descent step (or other method):

$$\psi \leftarrow \psi + \omega \widehat{h}_k$$

update the q-value target network

$$\psi' \leftarrow \varsigma \psi' + (1 - \varsigma) \psi$$

estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{M} \sum_{m=0}^{M-1} \nabla_{\theta} \mu_{\theta_k}(s_{m,t}) \cdot \nabla_a Q_{\psi}(s_{m,t}, a) \Big|_{a = \mu_{\theta_k}(s_{m,t})}$$

update the policy network with gradient ascent (or other methods like Adam):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g}_k$$

update the policy target network

$$\theta'_{k+1} \leftarrow \zeta \theta'_k + (1 - \zeta) \theta_{k+1}$$

end of for

end of for

Note that we wrote the policy gradient increment as $\nabla_{\theta} \mu_{\theta_k}(s_{m,t}) \nabla_a Q_{\psi}(s_{m,t}, a) \Big|_{a = \mu_{\theta_k}(s_{m,t})}$ but since it is usually computed through automatic differentiation libraries, it could be just have been written as $\nabla_{\theta} Q_{\psi}(s_{m,t}, \mu_{\theta_k}(s_{m,t}))$.

10. TD3 (Twin Delayed DDPG)

10.1 Base Idea

Twin Delayed Deep Deterministic Policy Gradient [Fujimoto et al. 2018], also known as TD3 because of the three consecutive Ds in the title, is an algorithm based on DDPG, with some modifications. Just like DDPG it is off-policy and for continuous actions.

A frequent problem with DDPG is that it may overestimate the q-value (see section 2.7), and as a consequence of it the policy learning will be distorted. The authors introduced some improvements to contrast the tendency of q-value overestimation:

- Twin Networks for Q-Value estimation
- Noise Injection into Target actions, with Action Clipping
- Delayed Policy updates

10.2 Twin Networks for Q-Value estimation

[Fujimoto et al. 2018] introduced two different q-value networks, with separated parameters. When computing the target value, both networks are used to calculate their respective q-value, the minor of which is used to bootstrap the target value. This contrasts the overestimation bias. Then both networks go through a q-learning step using the same target value just computed. The word “Twin” in the algorithm name comes from the two networks, that are called “Twin Networks”. Note that this mechanism is different from “Double Q-Learning” of section 2.8 .

10.3 Computing Target Actions

When using the two q-value networks to compute the target value, a further improvement is to inject noise into the action. In DDPG the action would be just the output of the policy

network, i.e. $a_{m,t+1} = \mu_{\theta'_{k+1}}(s_{m,t+1})$, but here a Gaussian noise is added to it. Given a normal \mathcal{N} with zero average and standard deviation σ , and given noise clipping boundary c :

$$a_{m,t+1} = \mu_{\theta'_{k+1}}(s_{m,t+1}) + \text{clip}(\mathcal{N}(0, \sigma), -c, +c). \quad (10.1)$$

In this way if some unjustified peak in the q-value estimate forms because of q-value overestimation, the action taken is not exactly the one in the middle of the peak, so the target value computation will be more robust. This trick induces a regularization, because it forces the policy learning mechanics to not exploit actions that coincide with sharp q-value peaks. The injected noise is clipped to not have too big deviations from the policy.

10.4 Delayed Policy Updates

The last improvement consists in updating the policy less frequently than the q-value networks. This improves learning stability because each change in policy may correspond to a change in the computation of the target value, resulting in too much target volatility. The original authors suggest to have one round of policy updating after 2 rounds of q-networks updating. With the same frequency update with Polyak averaging the q-value target networks and policy target network.

Also, note that only the first q-network is used to update the policy parameters through deterministic policy gradient.

Algorithm 10.1 TD3 (Twin Delayed DDPG)

Require: Policy network step size η

Require: Q-Value network step size ω

Require: Behavior policy β , usually $= \mu_{\theta_k}(s) + \text{Gaussian}$

Require: an empty replay buffer B

Require: Initialize parameters θ of network μ_{θ_k} with small random values

Require: Initialize parameters ψ_1 of network Q_{ψ_1} with small random values

Require: Initialize parameters ψ_2 of network Q_{ψ_2} with small random values

Require: Polyak Q-Value Step size ς

Require: Polyak Policy Step size ζ

Require: Action noise clipping boundary c

Require: Action noise standard deviation σ

Require: number of steps to wait before doing policy (and target networks) update d

copy Q-Value target network parameters $\psi_1' \leftarrow \psi_1$

copy Q-Value target network parameters $\psi_2' \leftarrow \psi_2$

copy Policy target network parameters $\theta' \leftarrow \theta$

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using behavior policy β

 put all transitions $\langle s, a, s', r, is_terminal \rangle$ of all trajectories D_k in replay buffer B

 For $n = 0, 1, 2, \dots, N - 1$ do a minibatch iteration:

 For $m = 0, 1, 2, \dots, M - 1$ do:

 From buffer B sample a transition $\langle s_{m,t}, a_{m,t}, s_{m,t+1}, r_{m,t}, is_terminal \rangle$

 if not $is_terminal$ then:

 sample action noise from Gaussian and clip it:

$$\epsilon \sim \text{clip}(\mathcal{N}(0, \sigma), -c, +c)$$

 compute $a_{m,t+1} = \mu_{\theta'_{k+1}}(s_{m,t+1}) + \epsilon$

$$Q_{m,t+1}^{min} = \min(Q_{\psi_1'}(s_{m,t+1}, a_{m,t+1}), Q_{\psi_2'}(s_{m,t+1}, a_{m,t+1}))$$

$$y_{m,t} = r_{m,t} + \gamma Q_{m,t+1}^{min}$$

 else:

$$y_{m,t} = r_{m,t}$$

 endif

 end of for

 estimate the q-value network gradients as:

$$\widehat{h_{1,k}} = \frac{1}{M} \sum_{m=0}^{M-1} (y_{m,t} - Q_{\psi_1}(s_{m,t}, a_{m,t})) \nabla_{\psi} Q_{\psi_1}(s_{m,t}, a_{m,t})$$

$$\widehat{h_{2,k}} = \frac{1}{M} \sum_{m=0}^{M-1} (y_{m,t} - Q_{\psi_2}(s_{m,t}, a_{m,t})) \nabla_{\psi} Q_{\psi_2}(s_{m,t}, a_{m,t})$$

update the q-value networks with a gradient descent step (or other method):

$$\psi_1 \leftarrow \psi_1 + \omega \widehat{h_{1,k}}$$

$$\psi_2 \leftarrow \psi_2 + \omega \widehat{h_{2,k}}$$

if $n \bmod d$ then:

estimate policy gradient as:

$$\widehat{g_k} = \frac{1}{M} \sum_{m=0}^{M-1} \nabla_{\theta} \mu_{\theta_k}(s_{m,t}) \cdot \nabla_a Q_{\psi_1}(s_{m,t}, a) \Big|_{a = \mu_{\theta_k}(s_{m,t})}$$

update the policy network with gradient ascent (or other methods):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g_k}$$

update the policy target network

$$\theta'_{k+1} \leftarrow \zeta \theta'_k + (1 - \zeta) \theta_{k+1}$$

update the q-value target networks

$$\psi_1' \leftarrow \varsigma \psi_1' + (1 - \varsigma) \psi_1$$

$$\psi_2' \leftarrow \varsigma \psi_2' + (1 - \varsigma) \psi_2$$

endif:

end of for

end of for

11. Soft Actor-Critic

11.1 Base Idea

Soft Actor-Critic (“SAC” in short) [Haarnoja et al. 2018A], [Haarnoja et al. 2018B], is another algorithm derived from DDPG that implements some techniques to improve it, since, as we already wrote, DDPG may overestimate the Q-Values (see section 2.7) and consequently jeopardize learning. SAC also tries to improve the fact that DDPG is brittle with respect to hyperparameter settings: it is not easy in DDPG to find a configuration of hyperparameters that leads to a satisfying learning. Like DDPG, SAC is off-policy, but while DDPG is only for continuous actions space, SAC has versions both for continuous and discrete actions.

To improve over DDPG it uses 2 networks for Q-Value estimation with some similarity with the Twin Networks of TD3 (to which it is contemporary) and some difference:

- It uses two different q-networks for Q-Value estimation, and learns them on the same target value, a thing that also TD3 does.
- The target value is computed using two target networks that are a Polyak-averaged version of the two learnt Q-Value networks, as in TD3
- The target value is computed using the minor Q-Value of the two, also as in TD3
- Similarly to TD3, the action sampled to be used in the q-function for the target computation has a noise component and it is clipped. Differently from TD3, here the clipping is done by a hyperbolic tangent function. Also, while TD3 is considered to have a deterministic policy with noise injection during the computation of target value, here in SAC the policy is considered to be stochastic, so the noise is part of the policy itself.
- Differently from TD3, the policy used to compute the action for the target is not a target policy network but it is directly the policy network itself.
- Differently from TD3, the q-network used to compute the q-value for the policy update is not always the first one, instead the minor q-value of the two is used
- Differently from TD3, the variance of the gaussian used to create noise in the action is not fixed, but learnt by the policy network itself.
- Differently from TD3, in SAC there are not delayed policy updates.
- Differently from TD3, in SAC the target value includes also an entropy term.

Let us see the algorithm in detail, starting from the last point.

11.2 Entropy Regularization

To understand this section it is necessary to know the concept of Information Entropy [Shannon 1948], the reader who is unaware of it may read appendices A.9 and A.10 . In brief, “*information entropy*” is a functional of a random variable that assesses how much random or surprising is its probability distribution, in the sense of how it would be easy to guess its value. Mathematically, the information entropy of a random variable X is defined as:

$$H(X) = E_{p(x)}[-\log_b p(x)] \quad (11.1)$$

With b base of the logarithm (it can be 2, or Napier’s e , or 10).

In SAC the random variable of which we want to compute the entropy is the action, whose probability distribution $p(x)$ is the policy π_θ . So, a high entropy of the action means that the policy is quite diversified, i.e. oriented to exploration, while a low entropy means that the policy is dictating quite a regular and stable behavior, i.e. oriented to exploitation.

When the random variable is continuous, the entropy it is named “differential Entropy” and it is denoted with lowercase h instead of uppercase. As described in appendix A.8, differential Entropy in general has different behavior and meaning than discrete variable’s information Entropy. But since in SAC with continuous actions the probability density function $p(x)$ is a normal, we know analytically what its differential entropy is (see eq. a.48 in appendix A.9), and we known that it is greater when the variance is greater, that means that even in this case a bigger differential entropy means more exploration and less exploitation.

In SAC, a term including the entropy of the action is added to the expected reward to obtain the objective function that must be maximized. This has the effect of regularizing the state value and the q-value, to help avoiding getting stuck in local maxima: it creates a trade-off between exploration and exploitation.

Moreover, promoting a certain level of information entropy of the action distribution fosters the learning of multimodal nearly-optimal policies.

If we introduce $u > 0$ as a coefficient that multiplies the information entropy, and that is bigger if we want to encourage exploration, and smaller if we want to encourage exploitation, the objective function is:

$$J_H(\pi) = E_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t (r_t + uH(\pi(\cdot | s_t))) \right] \quad (11.2)$$

Consequently, we have updated definitions of value and q-value that take the entropy in account:

$$V_H^\pi(s) = E_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t (r_t + uH(\pi(\cdot | s_t))) \mid s_0 = s \right] \quad (11.3)$$

$$Q_H^\pi(s, a) = E_{\tau \sim \pi} \left[\sum_{t=0}^{\infty} \gamma^t r_t + u \sum_{t=1}^{\infty} \gamma^t H(\pi(\cdot | s_t)) \mid s_0 = s, a_0 = a \right] \quad (11.4)$$

Here for the q-value we started adding the entropy only for $t = 1$, because the first action at $t = 0$ is the one already chosen as the argument of the q-value function. Anyway this is not the only possible way of doing it, and some papers adopt different choices.

Applying the relationship of the Bellman Equation we have:

$$\begin{aligned} Q_H^\pi(s, a) &= E_{s' \sim P, a' \sim \pi} [R(s, a, s') + \gamma(Q_H^\pi(s', a') + uH(\pi(\cdot | s')))] \\ &= E_{s' \sim P, a' \sim \pi} [R(s, a, s') + \gamma(Q_H^\pi(s', a') - u \log_b(\pi(a' | s')))] \end{aligned} \quad (11.5)$$

Since eq. 11.5 is an expectation, we can just sample that, and it will be an approximation. To do that we will obtain the transition $\langle s, a, s', r \rangle$ from the trajectory of the agent (or more precisely from the replay buffer containing the transitions from the trajectories), and we must sample a new a' each time we use that transition.

To sample a' we must compute:

$$a' = \tanh(\mu_\theta(s') + \sigma_\theta(s') \odot \mathcal{N}(0, 1)) \quad (11.6)$$

The symbol \odot denotes element-wise product.

The term $\sigma_\theta(s') \odot \mathcal{N}(0, 1)$ is added as a noise to the mean $\mu_\theta(s')$. This is different from TD3 where standard deviation of the gaussian noise was fixed. Here instead the neural network for

the policy will output both the mean μ_θ and the standard deviation σ_θ , depending on the state s . Since we want to learn the optimal standard deviation for the noise, depending on each state s , the right way to do that is to model the gaussian noise as standard deviation $\sigma_\theta(s)$ term that has an element-wise product with the sample from the normal $\mathcal{N}(0,1)$. Sampling operations are non-differentiable, but the sample of the normal $\mathcal{N}(0,1)$ can be considered just as a constant, so the backpropagation algorithm will compute the right gradient for $\sigma_\theta(s')$. This is commonly named “reparameterization trick”, and has been used also in other deep learning algorithms such as Variational Auto Encoders [Kingma and Welling 2014], [Doersch 2016]. If instead the noise was modelled directly as $\mathcal{N}(0, \sigma_\theta(s'))$, there would be no way to compute the gradients for it because that is only a sample from a random sampling process, which, as already said, it is non differentiable. Finally, the effect of hyperbolic tangent \tanh is to clip the value without doing a hard clipping.

11.3 Policy Update

Now, since we added the entropy term in the expectation inside the objective function, we must modify accordingly the policy network updating, inserting the entropy term also in it.

So if the policy gradient increment for DDPG was $\nabla_\theta Q_\psi(s_{m,t}, \mu_{\theta_k}(s_{m,t}))$, now it must be:

$$\hat{g} = \nabla_\theta (\tilde{Q}^{min}(s, \tilde{a}) - u \log(\pi(\tilde{a}|s))) \quad (11.7)$$

where \tilde{Q}^{min} is the minimum q-value among the two possible q-values estimations of the two q-values target networks, and \tilde{a} is an action sampled from the policy function from the current state with the same mechanism in which a' was sampled from the next state in eq. 11.6: sampling from a gaussian whose standard deviation σ_θ depends on the state and has learned parameters, and soft-clipping the resulting value through a \tanh .

$$\tilde{a} = \tanh(\mu_\theta(s) + \sigma_\theta(s) \odot \mathcal{N}(0,1)) \quad (11.8)$$

Let us now see a version of the algorithm with fixed entropy multiplier u . We follow the second version of the algorithm [Haarnoja et al. 2018B] in which the q-value function was

approximated but not the value function, while in the first version of the algorithm [Haarnoja et al. 2018A] also the value function approximation was used.

Algorithm 11.1 SAC (Soft Actor-Critic)

Require: Policy network step size η

Require: Q-Value network step size ω

Require: Behavior policy β , usually $= \mu_{\theta_k}(s) + \text{Gaussian}$

Require: an empty replay buffer B

Require: Initialize parameters θ of network μ_{θ_k} with small random values

Require: Initialize parameters ψ_1 of network Q_{ψ_1} with small random values

Require: Initialize parameters ψ_2 of network Q_{ψ_2} with small random values

Require: Polyak Q-Value Step size ς

Require: Action noise clipping boundary c

Require: Action noise standard deviation σ

Require: number of steps to wait before doing policy (and target networks) update d

copy Q-Value target network parameters $\psi_1' \leftarrow \psi_1$

copy Q-Value target network parameters $\psi_2' \leftarrow \psi_2$

For $k = 0, 1, 2, \dots$ do:

 collect trajectories $D_k = \{ \tau_i \}$ using behavior policy β

 put all transitions $\langle s, a, s', r, is_terminal \rangle$ of all trajectories D_k in replay buffer B

 For $n = 0, 1, 2, \dots, N - 1$ do a minibatch iteration:

 For $m = 0, 1, 2, \dots, M - 1$ do:

 From buffer B sample a transition $\langle s_{m,t}, a_{m,t}, s_{m,t+1}, r_{m,t}, is_terminal \rangle$

 if not $is_terminal$ then:

 sample action from Gaussian:

$$a_{m,t+1} = \tanh \left(\mu_{\theta}(s_{m,t+1}) + \sigma_{\theta}(s_{m,t+1}) \odot \mathcal{N}(0,1) \right)$$

$$Q_{m,t+1}^{min} = \min \left(Q_{\psi_1'}(s_{m,t+1}, a_{m,t+1}), Q_{\psi_2'}(s_{m,t+1}, a_{m,t+1}) \right)$$

$$y_{m,t} = r_{m,t} + \gamma Q_{m,t+1}^{min} + u \log(\pi(a_{m,t+1}|s_{m,t+1}))$$

 else:

$$y_{m,t} = r_{m,t}$$

 endif

 sample action for policy gradient descent:

$$\tilde{a}_m = \tanh\left(\mu_\theta(s_{m,t}) + \sigma_\theta(s_{m,t}) \odot \mathcal{N}(0,1)\right)$$

Compute Q-values for policy gradient descent

$$\tilde{Q}_m^{min} = \min\left(Q_{\psi_1'}(s_{m,t}, \tilde{a}_m), Q_{\psi_2'}(s_{m,t}, \tilde{a}_m)\right)$$

end of for

estimate the q-value network gradients as:

$$\widehat{h}_{1,k} = \frac{1}{M} \sum_{m=0}^{M-1} (y_{m,t} - Q_{\psi_1}(s_{m,t}, a_{m,t})) \nabla_{\psi} Q_{\psi_1}(s_{m,t}, a_{m,t})$$

$$\widehat{h}_{2,k} = \frac{1}{M} \sum_{m=0}^{M-1} (y_{m,t} - Q_{\psi_2}(s_{m,t}, a_{m,t})) \nabla_{\psi} Q_{\psi_2}(s_{m,t}, a_{m,t})$$

update the q-value networks with a gradient descent step (or other method):

$$\psi_1 \leftarrow \psi_1 + \omega \widehat{h}_{1,k}$$

$$\psi_2 \leftarrow \psi_2 + \omega \widehat{h}_{2,k}$$

update the q-value target networks

$$\psi_1' \leftarrow \varsigma \psi_1' + (1 - \varsigma) \psi_1$$

$$\psi_2' \leftarrow \varsigma \psi_2' + (1 - \varsigma) \psi_2$$

estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{M} \sum_{m=0}^{M-1} \nabla_{\theta} (\tilde{Q}_m^{min} - u \log(\pi(\tilde{a}_m | s_{m,t})))$$

update the policy network with gradient ascent (or other methods):

$$\theta_{k+1} \leftarrow \theta_k + \eta \widehat{g}_k$$

end of for

end of for

The second version of the algorithm [Haarnoja et al. 2018B] also describes a way to compute an adaptive entropy multiplier u , to counter the fact that it can be too sensitive to reward scaling.

12. Reward Shaping and Curriculum Learning

Reinforcement Learning may be defined as a method to automatically learn from experience without human help. In fact, this is not completely correct. One of the crucial aspects of actual Reinforcement Learning algorithms is the modelling of the rewards: it is up to the human designer to decide when a reward should be given and how high (or low) the reward is, and that must be done in a way that reflects the actual goal that the agent is supposed to learn. Reward modelling can be made in a simple way if the problem has a single, unique goal: a positive reward is given if the goal is reached, no other rewards are given while the goal is not reached. While this seems to make sense, it may work badly: the agent may reach a state very close to goal without fulfilling the goal, so it would not receive any reward for it and it would not learn that the reached state was good and desirable to be reached again. Since RL algorithms “propagate back” the information obtained from a state that obtained reward to the previous states that lead to it, agents are not able to learn when a state is ideally good if there is not a reward deriving from that state or future reached states. So, it is clear that for some tasks, having only a final reward when the goal is reached is not the best way to assign rewards, and the reward system should instead be designed and engineered. This means that the human choices about how to model the rewards may change a lot the performance of the same algorithm.

Designing the reward system is called “Reward Shaping” and it must be done in a way that “guides” the agent into intermediate subgoals. It must be noted that Reward Shaping is somehow like “cheating”: the more you shape rewards, the more you are inserting human knowledge in a system that was meant to learn automatically, without prior knowledge. That does not mean it is a bad thing, it just means that we have to be aware that we are inserting external knowledge, and potentially any bias or flaw that comes with it.

In more complicated settings, where the goal may not be unique and there are different “good things to do”, there may be rewards for minor goals, and sometimes this may lead to the fact that the agent learns only to solve the minor goals because the obtained rewards distract it from solve the bigger goal (small rewards may be much smaller than big rewards but much easier to get).

Another method used to make agents learn in complicated task is the “curriculum learning”: a simplified version of the task (or the environment) is used at the beginning, and when the

agent has learnt to reach the goal in that setting, a progressively more complete version is used.

Curriculum learning may also be intended as making the agent learn a set of certain skills or behaviours that are preparatory for the complete goal.

Also curriculum learning, like reward shaping, is injecting human knowledge into the system, because a human is deciding which simplified task or skill is necessary to be learnt before learning the true task, and that opens the possibility to bias the learning (i.e. without a certain curriculum learning maybe the agent would learn a better policy that is not based on what the curriculum designer thought were the necessary skills).

13. Imitation Learning

Imitation Learning [Schaal 1999] is a term used to describe methods that learn a behavior from examples: any observed behavior that is considered optimal or at least satisfactory may be used to obtain a policy. This opens the possibility of using human experts to produce demonstrations of the wanted behavior.

Some authors use “Imitation Learning” as a synonym of “Behavioral Cloning”, which is a supervised method, but it is more common to consider “Imitation Learning” as a general term for methods that learn from examples. Those include also “Inverse Reinforcement Learning” [Ng and Russell 2000] [Abbeel and Ng 2004], in which the examples are used to learn an approximate reward function, which in turn is then used inside a Reinforcement Learning algorithm.

13.1 Behavioral Cloning

Behavioral Cloning is a method to learn a policy in a supervised learning manner, using examples of supposedly optimal or almost-optimal behavior.

Similarly to Reinforcement Learning, it aims at selecting the best action depending on the state. Differently from RL, Behavioral Cloning does not use reinforcement signals such as rewards, but rather it is a supervised learning method in which examples reproduce the behaviour of an expert (that usually is a human). The input data X represent the state, and the label Y is the action taken by the expert in that state. It has the valuable characteristic of being able to use directly human knowledge in form of examples. Unfortunately, one of the issues of Behavioral Cloning is that training examples are usually created in a limited subset of the state space, and a trained agent may, during his functioning, exit from that subspace, because even small differences in action responses or in starting states may accumulate and lead to very different states. Once the agent is outside the subspace of states in which training examples have been produced (outside the training set distribution), it likely will not be able to generalize the new states, and the actions selected will not be optimal, or may even be disastrous. For this reason, it is important to cover a big part of the state space with training samples, and that could be difficult, if not impossible.




Figure 13.1

When the agent acts following a cloned behavior, small differences from the example behavior may add, so that the trajectory diverges, and the agent may find itself in a state space out of the training distribution, with unpredictable consequences. The figure is inspired by the similar one in [Levine 2021d]

13.2 Dataset Aggregation (DAgger)

A way to fix this is to learn a first policy through supervised learning as in normal Behavioral Cloning. Then, an agent following that policy collects new trajectories. The states reached in these new trajectories may be different from the states in the training set, so the human experts are asked to create new optimal examples starting from these states. Then a new supervised learning cycle on the first policy using those new examples will create a more refined second policy. This way of alternating between refining the policy with supervised learning, make the agent sample new states, and creating new examples from new states may go on for various cycles until the behavior is considered satisfactory. This method, named “DAgger” (“Dataset Aggregation”) was pioneered by [Ross et al. 2010].

Also having a training set with intentional errors and their respective corrections can improve learning, because making an error will include in the trajectory a state subspace that would not likely be included in a perfect-demonstration trajectory, but it could potentially be reached by an imperfect agent trained with behavioral cloning. Hence in this way the training data will include the right demonstration of how to behave also in those states.

13.3 Implementation

A Behavioral Cloning system may be built as a neural network of any kind of architecture, with a number of input neurons equal to the dimension of a sample, and a last layer whose output neurons are the ones that indicate the action to take (so if the actions are discrete and there are K different actions there will be K final neurons with a Softmax applied to them, if the actions are continuous there will be other K output neurons outputting the average of the computed distributions on the continuous values, etc.). In other words, it will be a neural network just like the one that you would build for a policy network in RL, except that this network will be trained as a supervised network with usual (minibatch) gradient descent algorithm, and not with any policy gradient algorithm.

A first thing to note is that the fact that the policy network may have the same architecture both in Behavioral Cloning and in policy gradient RL makes possible to train the same network both with IL and RL, so it is theoretically possible to integrate human knowledge with the RL process. For instance, the human-made examples may be used to learn an initial policy network in a supervised learning way, and then this policy network can be refined through Reinforcement Learning.

An interesting characteristic of Behavioral Cloning training is its mathematical relationship with Reinforcement Learning policy gradient methods, as we will see it below.

Let us start with the example of discrete actions. To use the same notation of Reinforcement Learning, we call $\pi_{\theta}(a_t|s_t)$ the Behavioral Cloning neural network that outputs a probability distribution over actions taking the state as input, and we imagine of having N example trajectories of length T_i each.

Now, the last layer of the network is a Softmax which implies that the gradient of the Loss is the negative of the log of the probability of the correct action. So, it turns out that this Behavioral Cloning minibatch gradient (of Loss function) will be computed as:

$$\nabla_{\theta} J_{IL}(\pi_{\theta}) = -\frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \quad (13.1)$$

While in case of policy gradient Reinforcement Learning the minibatch gradient (of expected return) is (from eq. 3.9 and eq. 3.29):

$$\nabla_{\theta} J_{RL}(\pi_{\theta}) = \frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} \log \pi_{\theta}(a_{i,t} | s_{i,t}) \Phi_{i,t} \quad (13.2)$$

Remarkably, they are similar (except for the initial minus in IL because in IL we minimize the objective while in RL we maximize the objective), the Reinforcement Learning gradient is just like the Behavioral Cloning gradient in which each sample gradient is also multiplied by $\Phi_{i,t}$, that is related to the rewards (it may be the rewards-to-go, the Advantage function etc.).

The same happens if the actions are not discrete but continuous, represented by a real number for each action parameter, so the output of the network is not a Softmax operator but real numbers, intended to be the means of a distribution, usually a Gaussian. For simplicity let us deal with the case of just one action parameter. In that configuration the Loss function of the Behavioral Cloning system would be the mean square error between the network output and the parameter chosen by the expert (from the sample). The RL formula for the log probability of the policy network will again be similar to the one of IL. To see that, for the RL policy network let us call σ^2 the variance of the gaussian of the action parameter and call $\pi_{\theta}(s_t)$ the output of the policy network, whose value is to be intended as the mean of the gaussian distribution of the parameter of the action. It is necessary to carefully not confuse the policy output $\pi_{\theta}(s_t)$ with π , the transcendental number π that is necessary for the gaussian formula and appearing in the equation. The agent will select the value of action a_t by sampling from a gaussian distribution with mean $\pi_{\theta}(s_t)$ and variance σ^2 . So, the probability of the action taken by the agent will be the probability of the action a_t considering that it is distributed by that gaussian.

$$\begin{aligned}
\nabla_{\theta} J_{RL}(\pi_{\theta}) &= \frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} \log \frac{1}{\sqrt{2\pi\sigma^2}} e^{-(\pi_{\theta}(s_{i,t}) - a_{i,t})^2 / (2\pi\sigma^2)} \Phi_{i,t} = \\
&\frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} \left(-\frac{(\pi_{\theta}(s_{i,t}) - a_{i,t})^2}{2\pi\sigma^2} - \log \sqrt{2\pi\sigma^2} \right) \Phi_{i,t} = \\
&-\frac{1}{2\pi\sigma^2} \frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} (\pi_{\theta}(s_{i,t}) - a_{i,t})^2 \Phi_{i,t}
\end{aligned} \tag{13.3}$$

Now, $1/(2\pi\sigma^2)$ can be cancelled from the equation because it is just a constant that may be incorporated with the learning rate. The same would apply to the Behavioral Cloning gradient if we derived it within the maximum log-likelihood framework: a $1/(2\pi\sigma^2)$ constant would appear, and we would incorporate it with the learning rate. So, the gradient for Reinforcement Learning policy is:

$$\nabla_{\theta} J_{RL}(\pi_{\theta}) = -\frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} (\pi_{\theta}(s_{i,t}) - a_{i,t})^2 \Phi_{i,t} \tag{13.4}$$

While the gradient for the Behavioral Cloning is given by the usual supervised regression loss gradient, the gradient of the square error:

$$\nabla_{\theta} J_{IL}(\pi_{\theta}) = \frac{1}{N} \sum_{i=0}^N \sum_{t=0}^{T_i-1} \nabla_{\theta} (\pi_{\theta}(s_{i,t}) - a_{i,t})^2 \tag{13.5}$$

Again, the two formulas are almost the same, apart from the minus sign (depending on ascending or descending the gradient) and the multiplication by $\Phi_{i,t}$.

So, intuitively, Reinforcement Learning policy gradient is like a Behavioral Cloning gradient that also uses the information of the rewards $\Phi_{i,t}$ to inform “how good” or “how bad” the action is.

I think it is surprising to find out that two different optimization problems related to two different underlying tasks end up with such similar computations.

14. Reinforcement Learning from Human Feedback

14.1 Introduction

Reinforcement Learning from Human Feedback [Christiano et al. 2017], is a method to obtain a reward function from preferences expressed by a human jury.

Imagine not having a sure and robust way to express the reward function for the goal that the agent must learn. This may happen in case there are not domain experts available to design a reward function, or because there is not a clear explicit reward in the process, or because it is not completely understandable how to analytically or algorithmically model the reward structure. Or maybe you are afraid that given a certain designed reward function, the RL algorithm would exploit it without learning the actual goal (i.e. maximizing the expected returns would not completely coincide with learning the goal).

In those cases, you could ask people to evaluate some examples of behavior, and use those evaluations to build a reward function. The evaluators are not needed to be domain experts, they must just be able to rank correctly different behaviors, i.e. they must recognize if a behavior is better, equal or worse than another.

14.2 Method

1. The first step of the method consists in using the trajectories that the behavior policy produced, and selecting small part of them called “trajectory segments”. Ideally, the trajectory segments should be long enough to contain an underlying evident behavior which the human judge may recognize, and also short enough to make it possible to evaluate a certain particular behavior but not a collection of too many unrelated different behaviors.

Two of such trajectory segments will be compared to decide which is better. They are not required to begin at the same starting state, hence there is no need to set a specific initial state in a simulated environment, nor to have the complete control of initial configuration in a real-world environment.

2. The second step consists in presenting to a human evaluator two different trajectory segments (for instance showing video clips of the agent), and the human evaluator will

assess which one is better, or if they are equally good, or if they are not comparable. Even if the two trajectory segments are in general beginning from different states, the human evaluator should be able to compare the goodness/badness of the behavior of each segment. If they think that it is not possible to judge, they will signal them as “not comparable”. Pairs which are judged not comparable will be discarded from the subsequent steps.

To speed up, this process of human evaluation is usually conducted by many persons (e.g. 40), to whom many trajectory segments pairs are given to be compared.

3. In the third step the comparison data is used to fit an approximate reward function (for instance a neural network) through a supervised learning algorithm.

At the end of the learning process, this approximate reward function can be used as reward function in any Reinforcement Learning algorithm.

The approximate reward function may be trained asynchronously with respect to the Reinforcement Learning algorithm. This means that it is possible to update online the approximate reward function during the RL learning, but it is also possible to use one fixed version of the approximate reward function to optimize the policy, while a new approximate version of the approximate reward is being computed using new comparison data from the most recent trajectories.

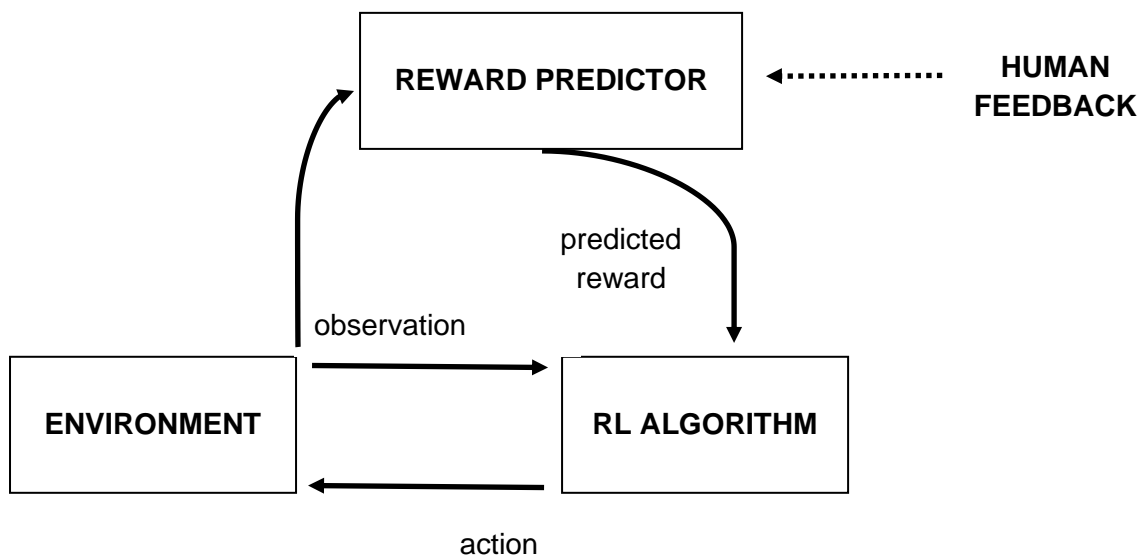


Figure 14.1

The approximated reward function (“Reward Predictor”) is trained asynchronously from the generation of the trajectories and from the human evaluation process. The figure is a reproduction from an analogous one in [Christiano et al. 2017].

14.3 Implementation

Before detailing the implementation, it is necessary to introduce a slightly different concept of reward function. Until now (see section 1.2) we used the following reward function $r_t = R(s_t, a_t, s_{t+1})$, which means that the obtained reward is a stochastic function of the current state s_t , the action a_t and the next state s_{t+1} reached as a consequence of the action. The next state s_{t+1} itself is modelled after a stochastic Transition function $P(s_{t+1}|s_t, a_t)$. Now to make our computation easier we should put together the stochasticity of s_{t+1} and r_t , considering a different form of reward function \bar{r}_t which depends only on the current state s_t , the action a_t , and is computed as an expectation with respect to the next state s_{t+1} .

$$\bar{r}_t = \bar{r}(s_t, a_t) = E_{s_{t+1} \sim P}[R(s_t, a_t, s_{t+1})] \quad (14.1)$$

Consider two different trajectory segments σ^1 and σ^2 , of length k , extracted either both from the same trajectory in different positions of the trajectory, or from two different trajectories:

$$\begin{aligned} \sigma^1 &= (s_0^1, a_0^1, s_1^1, a_1^1, \dots, s_{k-1}^1, a_{k-1}^1) \\ \sigma^2 &= (s_0^2, a_0^2, s_1^2, a_1^2, \dots, s_{k-1}^2, a_{k-1}^2) \end{aligned} \quad (14.2)$$

To denote that a person prefers segment σ^1 over segment σ^2 we use the notation $\sigma^1 > \sigma^2$, while in the opposite case we write $\sigma^1 < \sigma^2$.

When we express a preference like $\sigma^1 > \sigma^2$, it implies that there exists an underlying (and unknow to us) reward function \bar{r}_t such that:

$$\bar{r}(s_0^1, a_0^1) + \bar{r}(s_1^1, a_1^1) + \dots + \bar{r}(s_{k-1}^1, a_{k-1}^1) > \bar{r}(s_0^2, a_0^2) + \bar{r}(s_1^2, a_1^2) + \dots + \bar{r}(s_{k-1}^2, a_{k-1}^2) \quad (14.3)$$

Now, each segment pair preference can be associated to a discrete probability distribution μ over $\{1, 2\}$, assigning all the probability mass to the preferred choice, or half probability mass in case of equality:

$$\sigma^1 > \sigma^2 \Leftrightarrow \mu(1) = 1.0, \quad \mu(2) = 0.0$$

$$\sigma^1 < \sigma^2 \Leftrightarrow \mu(1) = 0.0, \quad \mu(2) = 1.0$$

$$\sigma^1 = \sigma^2 \Leftrightarrow \mu(1) = 0.5, \quad \mu(2) = 0.5$$

(14.4)

Each triplet $(\sigma^1, \sigma^2, \mu)$ can then be put in a database \mathcal{D} to be used for training.

Now, we assume that the probability that a human prefers σ^1 over σ^2 depends exponentially on the sum of all rewards in each segment, similarly to a Gibb's distribution (or Softmax):

$$Prob(\sigma^1 > \sigma^2) = \frac{e^{\sum_t \bar{r}(s_t^1, a_t^1)}}{e^{\sum_t \bar{r}(s_t^1, a_t^1)} + e^{\sum_t \bar{r}(s_t^2, a_t^2)}}$$

$$Prob(\sigma^2 > \sigma^1) = \frac{e^{\sum_t \bar{r}(s_t^2, a_t^2)}}{e^{\sum_t \bar{r}(s_t^1, a_t^1)} + e^{\sum_t \bar{r}(s_t^2, a_t^2)}}$$

(14.5)

Now, we want to learn an approximate reward function $\hat{r}(s_t, a_t)$ which is as close as possible to the real $\bar{r}(s_t, a_t)$. Please notice that we use the spiked hat to denote the approximate version, which can be implemented as a neural network. Hence in the equation 14.5 we substitute the reward function with the approximate reward function, and that will produce a predicted probability $\widehat{Prob}(\sigma^1 > \sigma^2)$ or $\widehat{Prob}(\sigma^2 > \sigma^1)$.

$$\widehat{Prob}(\sigma^1 > \sigma^2) = \frac{e^{\sum_t \hat{r}(s_t^1, a_t^1)}}{e^{\sum_t \hat{r}(s_t^1, a_t^1)} + e^{\sum_t \hat{r}(s_t^2, a_t^2)}}$$

$$\widehat{Prob}(\sigma^2 > \sigma^1) = \frac{e^{\sum_t \hat{r}(s_t^2, a_t^2)}}{e^{\sum_t \hat{r}(s_t^1, a_t^1)} + e^{\sum_t \hat{r}(s_t^2, a_t^2)}}$$

(14.6)

In order to learn $\hat{r}(s_t, a_t)$, we want to minimize (with respect to the parameters of \hat{r} , i.e. the weights of its neural network) the difference between the predictions and the human preferences, i.e. we want to bring $\widehat{Prob}(\sigma^1 > \sigma^2)$ as close as possible to $\mu(1)$, and $\widehat{Prob}(\sigma^2 > \sigma^1)$ as close as possible to $\mu(2)$.

Using cross entropy loss we obtain:

$$Loss(\hat{r}) = - \sum_{(\sigma^1, \sigma^2, \mu) \in \mathcal{D}} \mu(1) \log(\widehat{Prob}(\sigma^1 \succ \sigma^2)) + \mu(2) \log(\widehat{Prob}(\sigma^2 \succ \sigma^1)) \quad (14.7)$$

With that loss function we can train the neural network \hat{r} in a supervised learning way (for instance with stochastic gradient descent), using the triplets $(\sigma^1, \sigma^2, \mu) \in \mathcal{D}$ as samples.

The approximated reward function \hat{r} then can be used in the Reinforcement Learning algorithm. Since \hat{r} can be improved asynchronously through further training cycles based on new samples, it is wise to use a RL algorithm that is robust with respect to changes in the reward function, such as policy gradient methods.

This method has been used by original authors [Christiano et al. 2017], under the name of “Deep Reinforcement Learning from Human Preferences”, in simulated robotics environments and to play Atari games. In their experiments, trajectory segments were presented to the human evaluators in form of short video clips. They added the following refinements to the general algorithm described above:

- They used an ensemble of predictors to estimate \hat{r} . Each predictor was trained on a number of $|\mathcal{D}|$ triplets $(\sigma^1, \sigma^2, \mu)$ sampled from \mathcal{D} with replacement. Then each predictor was normalized (zero mean and constant variance) and the average of them was the final estimate of \hat{r} .
- A part of the data was used as validation set for each predictor. The supervised training algorithm used ℓ_2 regularization, and the regularization coefficient was adjusted to keep the regularization loss between 1 and 1.5 times the training loss.
- Other regularization techniques used in the predictor are batch normalization and dropout.
- Original authors declare to take in account the hypothesis of having 10% of the user evaluation to be considered random, uniformly distributed. It turns out that this anyway does not affect learning, see [Rahtz 2021], so this hypothesis can be ignored.

Moreover, to decide which trajectory segments among many should be presented to human evaluators to be assessed, the criterium was to choose the segments that appear to be predicted with more “uncertainty” by the ensemble predictors of \hat{r} .

It concretely means the following:

1. Sample a large number of segments into pairs.

2. For each pair, using each predictor of the ensemble compute a different predicted return (which is the sum of all the predicted rewards for all t).
3. Now, for each pair, for each predictor we have a segment that has a greater predicted return and is hence “preferred”.
4. The pairs in which more ensembles disagree on which is the preferred segment are the ones which more “variance”, or more “uncertainty”. These are the pairs that are selected to be presented to human judgement

Is to be noted that in some cases this way of selecting trajectory segments was found to worsen performance, so depending on the case it may be better to select at random the pairs to be presented to humans.

Another application of Reinforcement Learning from Human Feedback to play Atari games is described in [Ibarz et al. 2018].

RLHF has been used also to refine Large Language Models in [Ziegler et al. 2019], [Stiennon et al 2020], [Wu et al. 2021], [Ouyang et al. 2022]: the next chapter will deal with it.

15. Reinforcement Learning in Large Language Models

15.1 Introduction

Reinforcement Learning has been used to improve the performance of Large Language Models for different purposes, for instance at OpenAI [Ziegler et al. 2019] used RL to improve both summarization and style continuation, [Stiennon et al 2020] for summarization, while [Ouyang et al. 2022] used RL to align a model to human intentions, to follow instructions and to be less toxic and untruthful. Also [Bai et al. 2022a] at Anthropic used RL to align a LLM towards being “helpful, honest, and harmless”, while [Askell et al. 2021] to the same purpose experimented with “Preference Models” that are similar to the reward approximators of RLHF, but without the RL algorithm. It is quite common nowadays to train a language model in three sequential steps, using a different methodology at each step:

1. Self-supervised pretraining, using the text corpus.
2. Supervised fine-tuning, using a training set of tailored examples that teach some skill or discipline. The supervised examples are sometimes in the form of question/answer, especially if the skill to be learnt is to answer questions on a certain topic or solve problems in a certain field.
3. Reinforcement Learning fine-tuning. This may use an approximate reward function learnt through Reinforcement Learning from Human Feedback (see previous chapter). The RL algorithm to be used is commonly Proximal Policy Optimization, or a modification of it named Group Relative Policy Optimization [Shao et al. 2024]. As anticipated, RL may be used to align the model to the intended usage (e.g. focus on a particular topic/functionality avoiding general answers), or to make the model harmless and honest, or to teach the model to follow instructions. RL may be used also to train the model on particular skills as a further refinement after the supervised learning phase, or in substitution of it (e.g. when there is scarcity of training examples for the supervised fine-tuning phase).

We will not go through steps 1 and 2, which can be found in many other books and papers dedicated to LLMs, and we will detail only the Reinforcement Learning part. Note that not all LLM may have been trained in step 2 or 3, and some LLM may have a more complicate sequence of training steps: the 3-steps process is a simple generalization (for instance [Bai et

al. 2022a] do not use the Supervised Fine-tuning step, but they have a context distillation step).

15.2 RLHF in LLMs

15.2.1 Analogy between LLM and RL

Since Reinforcement Learning from Human Feedback has already been described in the previous chapter, here we will describe its application to LLMs.

We must define what, in a LLM, is the equivalent of a RL system's policy, environment, and action.

The LLMs is a probabilistic language model which, given an initial sequence of tokens, defines a distribution over the next tokens.

Hence, we can say that the initial sequence of tokens is equivalent to a RL observation/state, and the next tokens are the equivalent of the RL actions to be taken. The distribution over the tokens is the RL policy.

Typically in a LLMs query you provide an input sequence (also called "prompt") of n tokens $x = \{w_0, w_1, \dots, w_{n-1}\}$ and you obtain an output result of m tokens $y = \{w_n, w_{n+1}, \dots, w_{n+m-1}\}$. Since autoregressive/causal/decoder-only LLMs emit the output tokens sequentially, one after the previous, we can consider the state s_t as the concatenation of x and all the output tokens already emitted before step t : $s_t = \{w_0, w_1, \dots, w_{n-1+t}\}$ with t from 0 to $m - 1$.

The action a_t is the token emitted at time t , that is w_{n+t} .

It is common to use the notation π_ϕ^{RL} for the LLM policy currently being learnt through Reinforcement Learning, parametrized by ϕ . Note that the notation we used for the RL policy in other chapters was slightly different, we used π_θ instead.

So

$$\pi_\theta(a_t, s_t)$$

is now corresponding to:

$$\pi_\phi^{RL}(w_{n+t}, \{w_0, w_1, \dots, w_{n+t-1}\}) \quad (15.1)$$

The first thing to do is learning an approximate reward function through RLHF, and once we have one, we can optimize the policy/LLMs using a RL algorithm, such as Proximal Policy

Optimization. As already written, RLHF may be learnt in an asynchronous way, so once that we have a first version of the approximate reward function, we can further refine it without blocking the RL optimization process.

To do RLHF we need to collect a good number of trajectories (i.e. input x and output y), and we want them to be in the state space related to the purpose for which we are finetuning the LLM. This means that if we want to finetune for harmlessness we must use input sequences of tokens x that may potentially trigger toxic or dangerous answers, so to choose the less harmful behavior. Or, if we want to finetune the LLM to learn following instructions, we must use input sequences of tokens x that contain such instructions; if we want to improve the LLM in some problem-solving skill we must use input sequences that propose the kind of problems that we want to solve.

The samples to be sent to human for comparison can be collected from different sources: they can be generated by the pretrained or supervised finetuned policy, they can be generated by other models, they can come from existing baselines, and they can be also be generated by the current RL learnt policy (this is the case when we just learnt a RL policy and we want to update the reward function to make further improvements to the RL policy).

15.2.2 Differences from standard RLHF: Human Evaluation

There are some differences between the RLHF algorithm described in the previous chapter and how it is applied to LLMs.

The first difference is that when comparing two “trajectory segments” to choose the best one, in standard RLHF the two segments could potentially begin from different initial states (and usually they do), while in LLM they start from the same initial state, i.e. the same textual input x has been provided to generate both outputs.

The second difference is that in LLM application usually the segments have not the options to be evaluated as “equal” but there is a preference towards one of them.

The third difference is that in the reward function the whole output phrase is considered as a unique action (and not as a sequence of actions).

So the reward function will be denoted as:

$$\hat{r}(x, y) \tag{15.2}$$

That is equivalent to evaluate the whole concatenation of input and output:

$$\begin{aligned} \hat{r}(\{w_0, w_1, \dots, w_{n-1}\}, \{w_n, w_{n+1}, \dots, w_{n+m-1}\}) = \\ \hat{r}(\{w_0, w_1, \dots, w_{n+m-1}\}) \end{aligned} \quad (15.3)$$

This is an important point to understand and may create a little confusion. Let us explain it in detail. In the previous section, in the analogy between RL and LLM, we considered each generated token as an action. And this is how it will be considered also inside the RL algorithm. But the computation of the reward in the RL algorithm will be done only on the complete output, at the end of the phrase. It is like we had a RL trajectory with only one reward at the end of the trajectory, without rewards in the intermediate steps. This is because it makes sense to judge the whole phrase to understand if it is good or bad, and human evaluators will evaluate the complete phrase as well. Actually, in the RL algorithm there will be some additional reward for each intermediate token/action but it will not be computed through the reward function as it will be just a Kullback-Leibler divergence penalty (more on this later). This is the reason why (1) we make human evaluate only the complete output, and (2) we use the reward function only on the complete output.

Knowing that, let us see how the RLHF formulas of previous chapter change for LLMs usage. If we call x the initial sequence, and y_0, y_1 , the 2 different outputs generated by the LLM, the preferred output has index b (that means, it is y_b), and \mathcal{D} is the dataset of human evaluations. We have:

$$\begin{aligned} \widehat{Prob}(y_0 > y_1) &= \frac{e^{\hat{r}(x, y_0)}}{e^{\hat{r}(x, y_0)} + e^{\hat{r}(x, y_1)}} \\ \widehat{Prob}(y_1 > y_0) &= \frac{e^{\hat{r}(x, y_1)}}{e^{\hat{r}(x, y_0)} + e^{\hat{r}(x, y_1)}} \\ Loss(\hat{r}) &= - \sum_{(x, y_0, y_1, b) \in \mathcal{D}} (1 - b) \log(\widehat{Prob}(y_0 > y_1)) + b \log(\widehat{Prob}(y_1 > y_0)) \end{aligned} \quad (15.4)$$

The fourth difference of RLHF in its LLMs application is about the number of segments to compare at once.

While in [Stiennon et al 2020] the human evaluation is still among a pair of LLM outputs (so eq. 15.4 is to be used in that case), in [Ziegler et al. 2019] to make the humans consume less effort in the evaluation, they are not asked to compare a pair of segment/texts but instead to choose the favourite one among four of them at once. In this way the human evaluator must read only one initial sequence every four evaluated segments.

Since in this case the choice is among 4 segments and not 2, the loss function for the computation of the reward predictor changes a little.

If we call x the initial sequence, and y_0, y_1, y_3, y_4 the 4 different outputs generated by the LLM, and the preferred output has index b (that means, it is y_b) we have the softmax:

$$\widehat{Prob}(y_b) = \frac{e^{\hat{r}(x, y_b)}}{\sum_i e^{\hat{r}(x, y_i)}}$$

$$Loss(\hat{r}) = - \sum_{(x, y_0, y_1, y_2, y_3, b) \in \mathcal{D}} \log \left(\frac{e^{\hat{r}(x, y_b)}}{\sum_i e^{\hat{r}(x, y_i)}} \right)$$

(15.5)

A different method is used by [Ouyang et al. 2022]: for each input prompt they produce K outputs, with K between 4 and 9, to be ranked. Then, they make the human evaluators compare all the possible $\binom{K}{2}$ pairs from the combinations of those K outputs. Each of such pair comparison does not become a separate sample, because that would create overfitting since those pairs are correlated (each output would be used for $K - 1$ gradient updates). Instead, a single gradient update is done for all the $\binom{K}{2}$ output pairs for a certain input X , with that gradient equal to the total of all the gradients from the $\binom{K}{2}$ output pairs, divided by $\binom{K}{2}$. In other terms, using the indices i and j for two outputs in a pair among those of the combination $\binom{K}{2}$, and $b_{i,j}$ expressing the preferences between the outputs i and j :

$$\widehat{Prob}(y_i > y_j) = \frac{e^{\hat{r}(x, y_i)}}{e^{\hat{r}(x, y_i)} + e^{\hat{r}(x, y_j)}}$$

$$\widehat{Prob}(y_i > y_j) = \frac{e^{\hat{r}(x, y_j)}}{e^{\hat{r}(x, y_i)} + e^{\hat{r}(x, y_j)}}$$

$$Loss(\hat{r}) = - \sum_{(x, y_0, \dots, y_{K-1}) \in \mathcal{D}} \frac{1}{\binom{K}{2}} \sum_{(i, j) \in \binom{K}{2}} (1 - b_{i, j}) \log(\widehat{Prob}(y_i > y_j)) + b_{i, j} \log(\widehat{Prob}(y_j > y_i)) \quad (15.6)$$

15.2.3 Reward Model

Another remarkable thing is that, since the reward approximators must “understand” the language, they are built as neural network layers over the last layer of the language model that outputs the final embeddings. The reward approximators use the output of the unoptimized language model as input (i.e. the supervised-finetuned model if there has been that training phase, otherwise the pretrained model), and not of the RL-optimized model to avoid overfitting (see section 4.2 of [Ziegler et al. 2019]).

It became a common practice to use a smaller LLM for the reward approximator: [Ouyang et al. 2022] used a LLM with 6 billion of parameters as input for their reward approximator, while the LLM that they were optimizing with RL had 175 billion parameters: the smaller LLM seemed to be more stable and saved computation. When using RLHF with Large Language Models, it is common to call the reward approximator as “Reward Model”.

To summarize: the lower architecture of a Reward Model is copied from a LLM which is smaller than the LLMs that we want to RL-optimize, and on top of that smaller LLM we have one or more network layers which learn to output the reward, trained through a gradient descent optimization process based on the loss function of eq. 15.6 (or eq. 15.5, or eq. 15.4) .

15.2.4 Value function

The Reward Model, from a mathematical/algorithmic standpoint, could obviously be used also to evaluate incomplete outputs, i.e. the results of intermediate steps, because it takes a sequence of tokens as input. But inside the PPO algorithm applied to LLM, it is used only on the complete output.

On the other hand, here the Value Function is a function that is used to evaluate intermediate states, i.e. incomplete outputs, in the same way it is done in common PPO to evaluate the value of every state in the trajectory. The Value Function here in LLMs takes a sequence of tokens as input, and those tokens represent the current state. That implies that the tokens used as input for the Value Function are all the token of the prompt x , concatenated with all tokens emitted until time t .

$$V^\pi(s_t) = V^\pi(\{w_0, w_1, \dots, w_{n+t-1}\}) \quad (15.7)$$

Given the similarity of the type of input with the Reward Model, usually the Value Function estimate is initialized as a copy of the Reward Model, and then it is optimized via gradient descent on Least Squares, as in common PPO, where the measured error is the difference between the Value Function and the Rewards-to-Go $G(\tau_t)$, i.e. $(V^\pi(s_t) - G(\tau_t))^2$.

15.2.5 Kullback-Leibler Penalty for the Reward in RL Stage

Another difference between standard RLHF and its application to LLM is the usage during the Reinforcement Learning optimization of a reward function which adds a per-token Kullback-Leibler divergence penalty to the approximated reward function.

The penalty is proportional to the Kullback-Leibler divergence of the new policy from the old one, i.e. the Kullback-Leibler divergence of the token distribution of the RL-optimized LLM from the token distribution of the unoptimized LLM (the “unoptimized LLM” is the LLM as it was in the previous stage, that is supervised finetuned version, or in case there was not supervised fine-tuning it is the self-supervised pretrained version).

This KL penalty discourages outputs that are too different from the ones for which the approximated reward function has been learnt, because presumably on those outputs the approximated reward is less precise. In doing so it makes the LLM learn outputs that are not too different from the ones learnt in previous stages (so in some sense it encourages a refinement rather than a complete change in outputs). This is a “guard rail” against the possibility that the RL model finds ways of exploiting the reward model to maximize the

expected reward without actually improving the quality of the text but just emitting the right “gibberish” that has high reward.

It must be noted that this KL divergence penalty is for each generated token: that means that while we have only one “true” reward given by the Reward Model on the complete output, we have also a KL divergence penalty for each action/token.

Using the same notation as in sections above, we have input of n tokens $x = \{w_0, w_1, \dots, w_{n-1}\}$ and output of m tokens $y = \{w_n, w_{n+1}, \dots, w_{n+m-1}\}$.

Actions a_t are coinciding with token w_{n+t} , with t from 0 to $m-1$.

If we call β the multiplicative factor for the KL penalty, and π^{SFT} the LLM policy as it was at previous stage (SFT= supervised finetuned), the per-token penalty related to token w_{n+t} , i.e. action a_t is :

$$\begin{aligned} kl\ penalty(t) &= D_{KL}^t(\pi_\phi^{RL} || \pi^{SFT}) = -\log\left(\frac{\pi_\phi^{RL}(a_t, s_t)}{\pi^{SFT}(a_t, s_t)}\right) \\ &= -\log\left(\frac{\pi_\phi^{RL}(w_{n+t}, \{w_0, w_1, \dots, w_{n+t-1}\})}{\pi^{SFT}(w_{n+t}, \{w_0, w_1, \dots, w_{n+t-1}\})}\right) \end{aligned} \quad (15.8)$$

We see that the KL term $\log\left(\frac{\pi_\phi^{RL}(\dots)}{\pi^{SFT}(\dots)}\right)$, is not, as in common KL formula, inside an expectation with respect to π_ϕ^{RL} , because that term it is approximated by samples (the actual output of the LLM) and those samples are obtained following π_ϕ^{RL} , so it already converges to that expectation.

The multiplicative factor β can be either a constant or change dinamically.

15.2.6 Reward-to-go in RL Stage

We can now compute the LLM equivalent of reward-to-go, $G(\tau_t) = \sum_{k=t}^{T-1} \gamma^k r_k$ (see eq. 3.5), that is the total reward associated with the emission of output token at time t (i.e. token w_{n+t}). The output sequence has of m tokens, so $T = m$. The multiplier γ is usually set a 1, so that the reward does not decrease with the length of the phrase (but it is possible to experiment with other values).

The reward r_k consists of the sole KL penalty for all tokens except for the final one where in addition to the KL penalty we have the computation of the Reward Model $\hat{r}(x, y)$.

$$r_k = \begin{cases} \hat{r}(x, y) - \log \left(\frac{\pi_{\phi}^{RL}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})}{\pi^{SFT}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})} \right) & \text{if } k = m - 1 \\ -\log \left(\frac{\pi_{\phi}^{RL}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})}{\pi^{SFT}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})} \right) & \text{if } k < m - 1 \end{cases} \quad (15.9)$$

Which can be written also as:

$$r_k = \begin{cases} \hat{r}(\{w_0, w_1, \dots, w_{n+m-1}\}) - \log \left(\frac{\pi_{\phi}^{RL}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})}{\pi^{SFT}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})} \right) & \text{if } k = m - 1 \\ -\log \left(\frac{\pi_{\phi}^{RL}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})}{\pi^{SFT}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})} \right) & \text{if } k < m - 1 \end{cases} \quad (15.10)$$

Knowing that, we can write $G(\tau_t)$:

$$\begin{aligned} G(\tau_t) &= \sum_{k=t}^{T-1} \gamma^k r_k \\ &= \gamma^{m-1} \hat{r}(x, y) - \sum_{k=t}^{m-1} \gamma^k \log \left(\frac{\pi_{\phi}^{RL}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})}{\pi^{SFT}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})} \right) \end{aligned} \quad (15.10)$$

In case of $\gamma = 1$ (as commonly set):

$$G(\tau_t) = \hat{r}(x, y) - \sum_{k=t}^{m-1} \log \left(\frac{\pi_{\phi}^{RL}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})}{\pi^{SFT}(w_{n+k}, \{w_0, w_1, \dots, w_{n+k-1}\})} \right) \quad (15.11)$$

15.2.7 Training process

The training process as a whole is made by the following steps:

1. Collect the input texts, i.e. the prompts x that are related to the topic/purpose of your fine-tuning, and put them in the dataset \mathcal{D} .
2. Collect/Produce the K outputs $y_0 \dots y_{K-1}$ for each input x and put them in the dataset \mathcal{D} .

As we saw before, we have many options here. We can have just 2 outputs for each input to put them in a pair to be evaluated as in [Stiennon et al 2020], or we can have more and we can have them evaluated all together as [Ziegler et al. 2019], or pairwise in $\binom{K}{2}$ combinations as in [Ouyang et al. 2022].

To collect/produce the outputs, initially it is common to use the LLM (that is not RL-optimized) π^{SFT} to generate them. Some authors like [Stiennon et al 2020] use also original human-authored texts and other baselines. To avoid a distribution shift, during the RL-Optimization it could be helpful to generate new outputs using the current policy $\pi_{\phi}^{RL}(y|x)$, to retrain/adapt the reward approximator to the new distribution.

3. Ask human evaluators to assess their preference b for the tuples, i.e. choose the best output given the input.
4. Using the data from human evaluation train the Reward Model $\hat{r}(x, y)$ through the “Reinforcement Learning from Human Feedback” method, using the formula corresponding to your method among equations 15.4 , 15.5 , 15.6 .
5. Now that we have a reward approximator (also called “Reward Model”). We can use it to finetune our policy through RL learning. We use a new LLM/Policy π_{ϕ}^{RL} which is initialized with the weights of π^{SFT} . It is common to use PPO, but some authors use a derived algorithm called GRPO [Shao et al. 2024] (it will be explained in a next section). During the RL training it is common to use new input prompts x which are different from the ones used to train the reward model $\hat{r}(x, y)$, but it is not strictly necessary to avoid the ones that have been used for that.
6. During the running of the RL algorithm the token distribution of π_{ϕ}^{RL} changes, so it is necessary to continuously make it produce outputs to be submitted to human evaluation in order to update the reward approximator $\hat{r}(x, y)$ through RLHF (which is asynchronous and can be executed without blocking PPO).

A Collect human feedback


1 A text input is sampled from the dataset



2 Various policies are used to produce k outputs




3 The outputs are evaluated by humans




B Train reward model

4a Each output, preceded by the input is fed to the reward model together with the human



4b The reward model calculates a reward \hat{r} for each output



\hat{r}_b

\hat{r}_j

4c The loss of the reward model is calculated based on the reward and human evaluation

$$loss = -\log \left(\frac{e^{\hat{r}_b}}{\sum_i e^{\hat{r}_i}} \right)$$

C Train policy with PPO


5a A new text input is sampled from the dataset



5b The policy π generates a new output



5c The reward model calculates a reward for the output



5d The reward is used to update the policy via PPO

\hat{r}



Figure 15.1

The process of Reinforcement Learning from Human Feedback applied to Large Language Models training. The figure is inspired by the similar one in [Stiennon et al. 2020].

15.2.8 Pretrain Likelihood Term in RL Objective

Since the RL optimization may worsen the performance on the pretrain distribution (i.e. on the texts used for the pretraining phase), [Ouyang et al. 2022] inserted in the final RL objective a term proportional to the log-likelihood under the current policy/LLM of texts used in the pretrained phase. This should encourage the process to not forget the pretrain distribution. These models are called “PPO-ptx”.

15.2.9 PPO implementation details

Some implementation details to be aware of:

- The reward is batch-normalized to a $\mathcal{N}(0,1)$
- No dropout is used for the policy network
- Policy logits are temperature-scaled as usual in LLMs

Since optimizing one aspect of the LLM with RLHF may worsen another aspect of the LLM [Bai et al. 2022a], some authors used more than one reward model, each one specialized in rewarding a particular characteristic of the LLM. For instance [Touvron et al. 2023] at Meta, after noticing that optimizing for safety would sacrifice helpfulness, used one reward model for safety and one for helpfulness.

Also [Glaese et al. 2022] at DeepMind used two reward models to be summed to obtain the total reward. One reward model was trained with user preference data (which output text the human liked most), and the other was trained with rule-compliance judgments: the human evaluators checked if the output text was violating a list of rules that established guidelines for harmlessness and correctness. Moreover, that study used the RL algorithm A2C (see section 3.6) instead of PPO, as well as other studies at DeepMind like [Menick et al 2022].

For further practical details about RLHF using PPO implementation, see [Huang et al. 2023], [Huang et al. 2024], [Liu 2023], [Hu et al. 2024].

15.3 Group Relative Policy Optimization

Group Relative Policy Optimization (in short “GRPO”) is a modification of the PPO algorithm applied to Large Language Models made by [Shao et al. 2024] and used in the popular model Deepseek-R1 [DeepSeek-AI 2024].

Its differences with respect to RLHF PPO are the following:

1. It does not use the Value Function. Instead, to compute a surrogate of the Advantage, it samples many outputs for each input (commonly named a “group” of outputs), computes their rewards and normalizes them (subtracts the average and divides by the standard deviation).
2. It does not use per-token Kullback-Leibler divergence in the computation of the rewards (or returns). It just sums the Kullback-Leibler divergence to the loss, without complicating the computation of the Advantage.

Let us see those points in detail.

15.3.1 Advantage surrogate without Value Function: Outcome Supervision

In order to spare memory, authors devised a way to not use a Value Function. Since we are not using a Value Function, we need a different way to compute the Advantage. Authors in [Shao et al. 2024] describe two alternative ways to do so: “Outcome Supervision” and “Process Supervision”. Let us begin with the former.

For each input text x we sample a group of Z different output texts y_j (with $j = 0, 1, \dots, Z - 1$) and compute the reward of the complete text of each y_j :

$$\hat{r}_j = \hat{r}(x, y_j) \tag{15.12}$$

We do not add any per-token Kullback-Leibler divergence.

We normalize the rewards:

$$mean(\hat{r}) = \frac{1}{Z} \sum_{i=0}^{Z-1} \hat{r}_i$$

$$std(\hat{r}) = \sqrt{\frac{1}{Z-1} \sum_{j=0}^{Z-1} (\hat{r}_j - mean(\hat{r}))^2}$$

$$\tilde{r}_j = \frac{\hat{r}_j - mean(\hat{r})}{std(\hat{r})}$$
(15.13)

Note that this is equivalent to using the mean as baseline and then scaling the result. This will be used as Advantage surrogate, instead of the usual Advantage in PPO.

$$\hat{A}_{j,t} = \tilde{r}_j$$
(15.14)

It must be noted that since we sampled Z different output texts, we are using all of them to compute the policy update (i.e. all y_i are used as sample trajectories), but because each sampled output is meant to potentially give the same magnitude of contribution to the gradient, the gradient of each single token in each output is scaled proportionally to the length of the output.

Hence, if we refer to the inner loop in PPO algorithm 8.1 where we had the estimate policy gradient as:

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{\tau \in D_k} \sum_{t=0}^{T-1} \nabla_{\theta_{k+1}} L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_t$$

if we want to rewrite it considering that for each input text x , we will have Z different sampled trajectories τ :

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{x \in D_k} \frac{1}{Z} \sum_{j=0}^{Z-1} \frac{1}{T_{x,j}} \sum_{t=0}^{T_{x,j}-1} \nabla_{\theta_{k+1}} L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_{j,t}$$
(15.15)

Where $T_{x,j}$ is the length of the trajectory j for the input prompt x , and the clipped surrogate loss $L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_{j,t}$ is equal to:

$$L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_{j,t} = \min \left(\frac{\pi_{\theta_{k+1}}(a_{j,t}|s_{j,t})}{\pi_{\theta_k}(a_{j,t}|s_{j,t})} \hat{A}_{j,t}, \text{clip} \left(\frac{\pi_{\theta_{k+1}}(a_{j,t}|s_{j,t})}{\pi_{\theta_k}(a_{j,t}|s_{j,t})}, 1 - \epsilon, 1 + \epsilon \right) \hat{A}_{j,t} \right) \quad (15.16)$$

We can notice that the term $\frac{1}{T_{x,j}}$ is what scales the contribution of the single token proportionally to the length of the output.

15.3.2 Process Supervision

The previous method only computes the reward of the whole output, and this may not be well suited for Large Reasoning Models and math problem solvers in which intermediate tokens are of fundamental importance in the step-by-step reasoning.

This motivates the experimentation with a second method, named “Process Supervision” in which the reward model is used to approximate a reward for each token.

Like before, we sample a group of Z different output texts y_j (with $j = 0, 1, \dots, Z - 1$), but now we compute a reward r_k^j for each token k emitted in output y_j .

If x is the input text composed of n tokens $\{w_0, w_1, \dots, w_{n-1}\}$, and y_j is the j th output text composed of m_j tokens $\{w_n^j, w_{n+1}^j, \dots, w_{n+m_j-1}^j\}$:

:

$$r_{j,k} = \hat{r}(x, \{w_n^j, \dots, w_k^j\})$$

$$r_{j,k} = \hat{r}(\{w_0, w_1, \dots, w_{n-1}\}, \{w_n^j, \dots, w_k^j\}) \quad (15.17)$$

In RL terminology, this can be thought as:

$$r_{j,t} = \hat{r}(s_{j,t}, a_{j,t})$$

with $a_{j,t}$ = new token generated at time t for output j
 $s_{j,t}$ = input tokens + tokens generated until time t in output j

(15.18)

Now, similarly to before in “Outcome Supervision”, we normalize the rewards, but this time we sum for all j and all k .

$$\begin{aligned}
 c &= \sum_{j=0}^{Z-1} m_j \\
 \text{mean}(\hat{r}) &= \frac{1}{c} \sum_{j=0}^{Z-1} \sum_{k=0}^{m_j-1} \hat{r}_{j,k} \\
 \text{std}(\hat{r}) &= \sqrt{\frac{1}{c-1} \sum_{j=0}^{Z-1} \sum_{k=0}^{m_j-1} (\hat{r}_{j,k} - \text{mean}(\hat{r}))^2} \\
 \tilde{r}_{j,k} &= \frac{\hat{r}_{j,k} - \text{mean}(\hat{r})}{\text{std}(\hat{r})}
 \end{aligned}$$

(15.19)

And then to obtain the surrogate Advantage we sum the $\tilde{r}_{j,k}$ for all the tokens k generated from the current timestamp t until the end of the output i .

$$\hat{A}_{j,t} = \sum_{k=t}^{m_j-1} \tilde{r}_{j,k}$$

(15.20)

15.3.3 Kullback-Leibler Divergence for the Loss

As anticipated, in GRPO the per-token Kullback-Leibler divergence is not summed to the reward to compute the total returns. Instead, it is summed to the Loss, and it is computed with the approximation described in [Schulman 2020] which is slightly biased but with a low variance.

$$\begin{aligned}
D_{KL}^{i,t}(\pi_{\phi}^{RL} || \pi^{SFT}) &= \frac{\pi^{SFT}(a_{i,t}, s_{i,t})}{\pi_{\phi}^{RL}(a_{i,t}, s_{i,t})} - \log \frac{\pi^{SFT}(a_{i,t}, s_{i,t})}{\pi_{\phi}^{RL}(a_{i,t}, s_{i,t})} - 1 \\
&= \frac{\pi^{SFT}(w_{i,n+k}, \{w_{i,0}, \dots, w_{i,n+k-1}\})}{\pi_{\phi}^{RL}(w_{i,n+k}, \{w_{i,0}, \dots, w_{i,n+k-1}\})} - \log \frac{\pi^{SFT}(w_{i,n+k}, \{w_{i,0}, \dots, w_{i,n+k-1}\})}{\pi_{\phi}^{RL}(w_{i,n+k}, \{w_{i,0}, \dots, w_{i,n+k-1}\})} - 1
\end{aligned} \tag{15.21}$$

15.4 Decoupled Clip and Dynamic sAmpling Policy Optimization (DAPO)

Decoupled Clip and Dynamic sAmpling Policy Optimization (DAPO) [Yu et al. 2025] is a variation of GRPO, made by researchers from ByteDance (TikTok). The difference with GRPO are the following:

- No KL divergence penalty
- A higher clipping threshold
- Dynamic Sampling, i.e. excluding samples that would give very small gradients
- Token-Level Policy Gradient Loss
- Overlong Reward Shaping, i.e. setting a soft and progressive penalty for very long outputs

Let us discuss them more in detail.

15.4.1 No Kullback-Leibler Divergence penalty

When we are dealing with reasoning models (e.g. math solvers, code generators) the distribution of the refined model can differ much from the distribution of the initial model because of the usage of long Chain-of-Thoughts, and it is something that we may expect to not being a bad thing, so we do not need to put the KL difference to curb it. This is reason why in DAPO authors decided to no add any penalty based on the KL divergence from the current RL LLM policy to the previous pretrained or supervise-finetuned LLM policy.

15.4.2 Higher Clipping Threshold

Authors noted that the entropy of the RL policy collapsed quickly, this means a lesser variety of actions/tokens with consequently less exploration. Their analysis hinted that the upper clipping threshold restricts the probability increase of low-probability tokens. The solution is to use different ϵ for the high and low clipping threshold: ϵ_{high} and ϵ_{low} , with ϵ_{high} relatively bigger than ϵ_{low} :

$$L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_{j,t} = \min \left(\frac{\pi_{\theta_{k+1}}(a_{j,t}|s_{j,t})}{\pi_{\theta_k}(a_{j,t}|s_{j,t})} \hat{A}_{j,t}, \text{clip} \left(\frac{\pi_{\theta_{k+1}}(a_{j,t}|s_{j,t})}{\pi_{\theta_k}(a_{j,t}|s_{j,t})}, 1 - \epsilon_{low}, 1 + \epsilon_{high} \right) \hat{A}_{j,t} \right) \quad (15.22)$$

15.4.3 Dynamic Sampling

DAPO has been used to train a LLM for math problem solving. It does not use another LLM as reward model: the reward is of Rule-Based type. If the answer is equal of the ground truth, then the reward is 1, otherwise it is -1.

When we sample the group of Z different outputs, if all outputs are equivalent to the ground truth for the problem solution, the advantage will be zero. This will lead to having a zero gradient. With the progression of the learning, the number of outputs equivalent to the ground truth will increase, that means that there will be less prompt in each batch that will give an effective gradient, which in turn will increase the variance in gradient. To contrast that, authors decided to discard the outputs with accuracy 1 or 0 and substitute them with additional samples until there are Z of them with the accuracy in the right interval.

15.4.3 Token-Level Policy Gradient Loss

In GRPO the gradient loss of each token was scaled proportionally to the length of the output. But this makes any good action inside a long output less relevant for the computation of the total gradient. Authors wanted instead to make each token count the same towards the gradient computation, without depending on the length of the final generated text, so instead of scaling each token proportionally to the length of the output to which it belongs, each token is scaled proportionally to the sum of the lengths of all the outputs in the same group.

$$\widehat{g}_k = \frac{1}{|D_k|} \sum_{x \in D_k} \frac{1}{\sum_{j=0}^{Z-1} T_{x,j}} \sum_{j=0}^{Z-1} \sum_{t=0}^{T_{x,j}-1} \nabla_{\theta_{k+1}} L_{\pi_{\theta_k}}^{CLIP}(\theta_{k+1})_{j,t}$$

(15.23)

15.4.4 Overlong Reward Shaping

In math problem solving, which is the setting where DAPO was applied, it is not uncommon to have generated answers that are longer than usual. Setting a maximum length for the generated output, as commonly done, would truncate longer answers, resulting in the evaluation of an incomplete answer and hence a lower reward bias. Authors verified such hypothesis, comparing the results of usual training with the ones where an “Overlong filtering” was applied to mask the loss of truncated samples.

Then they invented a way to make the LLM learn to generate shorter answers, named “Soft Overlong Punishment”: they introduced a soft penalty whenever the answer was bigger than a certain threshold, and such a penalty is increasing proportionally to how much the threshold is surpassed.

If we name L_{max} the maximum length of the generated output, and we name L_{cache} the length of the space before reaching L_{max} where there is the soft punishment, then the threshold length for which there is not penalty is equal to $L_{max} - L_{cache}$.

The penalty is defined as:

$$R_{length}(y) = \begin{cases} 0 & \text{when } |y| \leq L_{max} - L_{cache} \\ \frac{(L_{max} - L_{cache}) - |y|}{L_{cache}} & \text{when } L_{max} - L_{cache} < |y| \leq L_{max} \\ 1 & \text{when } L_{max} < |y| \end{cases}$$

(15.24)

Where $|y|$ is the length of the output y .

With such a penalty during training the LLM is softly pushed towards generating answers that do not exceed the length threshold.

15.5 RLAI (Constitutional AI)

Reinforcement Learning from Artificial Intelligence Feedback [Bai et al. 2022b], also named “Constitutional AI”, is a LLM refinement method derived from RLHF which aims at using a Large Language Model as evaluator instead of humans, making easier to scale the process.

The experiment described in the paper actually uses the LLM evaluator only to judge the “harmfulness” aspect, while the “helpfulness” aspect is still judged by human evaluators.

The underlying idea is to have a list of rules or principles, i.e. a “Constitution”, compiled by humans, which describe the criteria to evaluate and revision an answer from the standpoint of harmfulness.

There is a first supervised fine-tuning phase in which a LLM is asked to criticize and revise answers to potentially harmful prompts, using the constitutional rules as criteria, and then use the revisions to supervise fine-tune the target LLM.

Then there is a second phase in which the LLM evaluator will be asked, for each prompt, to judge two answers and decide which is better according to the rules. This preference data will then be used to train a Reward Model as in standard RLHF, and finally the target LLM will be refined by Reinforcement Learning using that Reward Model.

Let us explain that in more detail.

15.5.1 Initial Setup

At the beginning of the process it is necessary to have:

- a **list of rules or principles** (a “**Constitution**”), to critique, revise and evaluate texts with respect to the harmfulness aspects. The principles are created by humans.
- a “**purely helpful assistant**” which is a LLM that has been previously improved in its helpfulness by RLHF. We also have the dataset from that process, i.e. the prompts, the answers, and the human preference data about the answers. The original author did not have a supervised learning phase for this LLM, but since they claim to use the same method as in [Bai et al. 2022a] we can assume they had a context distillation step before RLHF (context distillation is described in [Askell et al. 2021]).
- a “**pretrained LLM**” which has not undergone any RLHF (and supposedly, has not undergone any supervised finetuning). This is the LLM that we are going to refine.
- another pretrained LLM named “**Feedback Model**” which will be the one used to evaluate the answers in the Reinforcement Learning phase.

- **a Dataset of potentially harmful prompts**, that are likely to trigger a harmful answer, commonly named “red teaming” prompts.

15.5.2 Supervised Phase

The first stage is about doing a supervised finetuning with answers that are revised following the constitutional principles.

First of all, we need to create revised versions of harmful answers.

Let us describe each step:

1. Use the initial model, the "purely helpful assistant" to create an output answering a harmful prompt.
2. Ask the model to critique its own answer, basing of one of the rules/principles (randomly drawn), and then revise the answer based on the critique.
This is realized by appending to the prompt+answer text a phrase instructing to critique it basing on the principle drawn from the constitutional list.
Then, to the context obtained by this answer (prompt+answer+critique), another phrase is appended, instructing the LLM to rewrite the answer in a way that would make it harmless.
3. Repeat the critique/revision procedure many times, using each time a new context consisting in the initial harmful prompt followed by the final revised answer of the previous revision step, and randomly drawing a rule from the constitutional list.
To improve the process, the instructions asking for critique and revision may be followed by some few-shot examples.
4. Save the final revision to be used supervised fine-tuning.

Repeat the all process for all the potentially harmful prompts in the dataset, so to have a dataset of good examples to use for supervised finetuning.

Then, we use those pairs of prompt+final revision to do the supervised finetuning on the pretrained LLM.

Actually, while in section 1.2 of [Bai et al. 2022b] authors claim to use "final revised responses", in section 3.1 they report using "revisions (from all revisional steps)", but it is clear that it is the final revised response that must be used, since it is the most revised one.

Please note that the pretrained LLM is the one that is subject to finetuning here, not to be confused with the "purely helpful assistant" which instead was not only pretrained but also already trained by RLHF for helpfulness and remains untouched.

To preserve helpfulness, some prompt used previously during the RLHF training of the "purely helpful assistant" are used to make the "purely helpful assistant" answer, and those pairs of prompt+answer for helpfulness are used in the supervised fine-tuning together with the prompt+revision data for harmlessness.

The obtained model is named SL-CAI.

15.5.3 Reinforcement Learning Phase

The second stage is about asking to a pretrained LLM to use constitutional principles to evaluate answers to a prompt (choosing an answer among a pair), then use those preference data to train a reward model, and finally use Reinforcement Learning to refine the target LLM using that reward model.

Note that the pretrained LLM that will be used to judge the answer basing on the principles instead of using humans evaluators is named "Feedback Model", and it is not the same model that we are optimizing (which at this point is SL-CAI).

Let us describe each step:

1. Sample from the supervised fine tuned model (SL-CAI): for each harmful prompt in the dataset the supervised finetuned model is used to answer two outputs.

For each prompt build a multiple-choice question in which it is given the prompt and it is asked which of the two answers is better, with respect to one of the principles of the constitution (randomly draw the principle).

Prepend some examples of the multiple-choice decision before the actual question, to have few-shot examples. Those few-shots examples are pre-written by humans and their principles are drawn independently.

Make the "Feedback Model" answer that preference question.

Please note that the principles for the RL phase are worded differently than the ones for the supervised phase, because the former ask the LLM to choose, the latter ask the LLM to critique and revise.

2. Use this preference data to train a reward model in the common RLHF way: this preference data about the harmfulness are mixed with the existing preference data generated by Humans about helpfulness to train the reward model.
3. use the reward model to train the SL-CAI model with RL (that is in the same way it would be in RLHF, but this time it is named RLAIF since some preferences have been generated by AI).

The final model obtained in this way is named “RL-CAI”.

Authors experimented also a variation of this method in which the preference between two answers is generated through "Chain-of-thought", that is appending "Let's think step-by-step:" to the final prompt. In this case they don't use the "Feedback Model" to answer, but the "purely helpful assistant", which is supposed to elaborate better the steps.

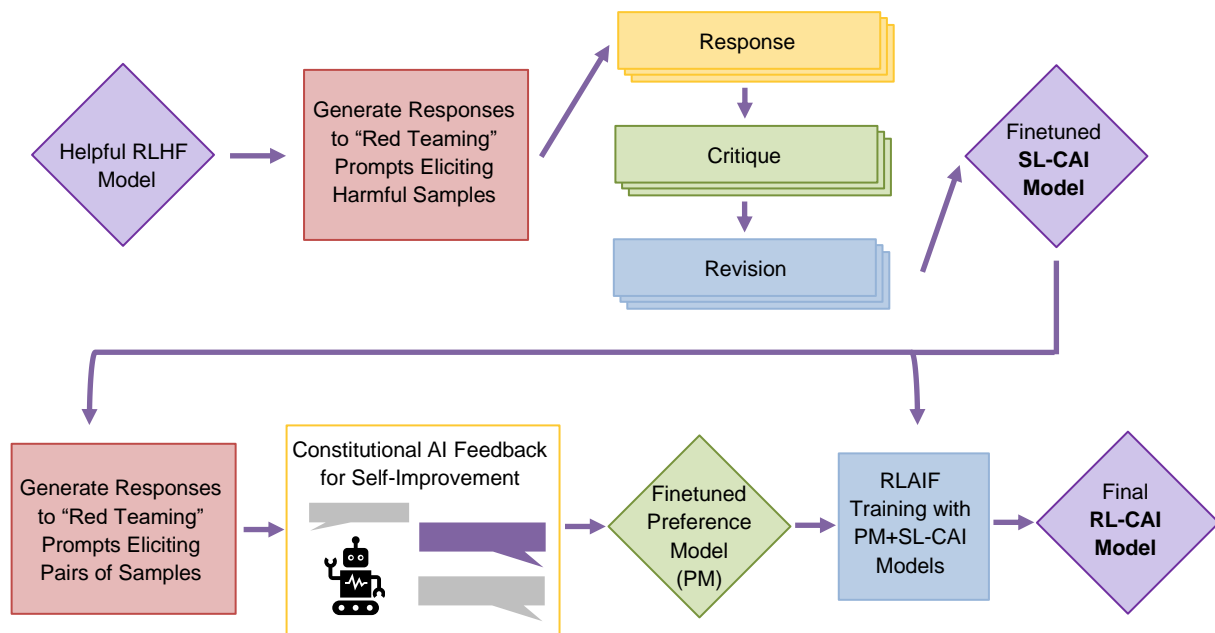


Figure 15.2

The process of RLAIF. The steps at the top are the ones of the supervised fine-tuning stage, while the ones at the bottom are the ones of the reinforcement learning stage. The figure is inspired by the similar one in [Bai et al. 2022b].

15.6 Other RL without Human Feedback

In the previous section we saw how it is possible to build a reward model without a team of human evaluators, having humans only writing down the principles while the LLMs follow them to both do a supervised finetuned process and build the reward model for the RL process.

That was feasible because it was used to refine the harmfulness aspect, that is something that may be defined by a series of principles.

But is it possible to use Reinforcement Learning to improve Large Language Models without human feedback also in other aspects or skills ?

It all goes down to the possibility of having available a reward model for that particular skill (or aspect), an example of which are the Math-Shepherd model [Wang et al. 2023] which can evaluate math solutions, and similarly the reward model for math problems in [Shao et al. 2024]. Reward Models trained to evaluate the solutions of math problems are sometimes named “Verifiers”.

Also classical algorithms and programs can be used as reward mechanism (sometimes named “Rule-Based RM”) as is the case of [DeepSeek-AI 2024] where authors used the LeetCode compiler to test coding solutions and a rule-based system to verify the final answer of mathematical problems, or in [Yu et al. 2025] (which we saw in a previous paragraph) where the reward for math problem solving questions was computed confronting the result with the ground truth.

16. Afterword

This has been just an introduction to the theory of Deep Reinforcement Learning. To not make it too boring, some mathematical passages and proofs have been skipped, some others have been simplified, but still some have been necessarily reported in detail, to create a consistent presentation that follows a principled thread.

This introduction has focused on “Deep” RL, that means that I have not dealt in detail with tabular methods, even if there has been a great amount of research on them and they cannot be ignored. Even among Deep RL methods, since this is just an introduction, I did not examine all algorithms, but only a didactically representative small subset of them.

The reader who wants the complete math and proofs, as well as the reader that wants to know more about tabular methods or about other deep RL methods, is left to the reference literature.

Also, while algorithms are described in detail, I acknowledge that it may not be easy for the primer to understand exactly how to implement a RL system after reading this introduction. To that purpose I suggest the reader to follow one of the many introductory courses on the internet (there are excellent ones both from MOOCs platforms and from famous brick and mortar universities), that focus on the practical side. A valid help to understand RL algorithms may come from checking the open source implementations of the algorithms available online, some of which are made by the same authors of the reference literature.

A topic that I did not discuss but it is worth a final mention is the fact the Reinforcement Learning is not “sample efficient”, that means that it needs a lot of samples (a lot of experience, or trajectories, or “trial and error”) to learn good policies in complex environments. Often training time may be much longer than expected, for instance longer than (usually) with supervised learning systems. This and other difficulties of Reinforcement Learning practice are well detailed in an article written by [Irpan 2018], which, even if not recent, I suggest every RL practitioner to read. Another useful disquisition on the RL practice, regarding debugging RL algorithms, can be found in [Jones 2021].

Appendix A: Information Theory Refresh

A.1 Kullback-Leibler Divergence

It is a way to compute how different two distributions are. Given two distributions P and Q , the Kullback-Leibler divergence (also named *Relative Entropy*) of P from Q is defined as:

$$\text{continuous} \quad D_{KL}(P||Q) = \int_x P(x) \log \frac{P(x)}{Q(x)}$$

$$\text{discrete} \quad D_{KL}(P||Q) = \sum_x P(x) \log \frac{P(x)}{Q(x)}$$

(a.1)

Where “log” is the natural logarithm (even if it is possible to compute it also using different bases for the logarithm, such as base 2).

- KL divergence is always non-negative (as a consequence of *Gibb's Inequality*): $D_{KL}(P||Q) \geq 0$.
- When the two distributions are the same, KL divergence is zero $D_{KL}(P||P) = 0$.
- KL divergence is not symmetrical: it is easy to see that in general $D_{KL}(P||Q) \neq D_{KL}(Q||P)$, hence it is not a metric. This is the reason why it is more precise to say “KL divergence of P from Q ” than just “KL divergence between P and Q ”.
- Anyway a symmetric computation may be devised: $Symmetric_{KL} = D_{KL}(P||Q) + D_{KL}(Q||P)$

In Deep Reinforcement Learning, when using policy gradient methods, you may have either discrete (categorical) actions or continuous actions. If actions are categorical you have as many output neurons as the number of possible actions within which to choose. Each output neuron represents the probability of choosing the related action. In this case the distribution of actions is discrete, and the KL divergence may be computed with the formula for discrete distributions (the second formula in eq. a.1), where x is the action.

If instead actions are continuous, the policy neural network outputs a value for each dimension of the action. This output value usually represents the mean of a normal distribution from which

it is possible to sample the action to take. The standard deviation of such a normal distribution may be a fixed value (decided ex ante), or it may be another output of the neural network (one for each dimension of the action). So it may be useful to recall an analytical computation of Kullback-Leibler divergence for normal distributions.

For univariate normal distributions (see proof in [Statproofbook 2021A]) where variable x is unidimensional and distributed normally with mean μ and variance σ^2 :

$$P: x \sim \mathcal{N}(\mu_1, \sigma_1^2)$$

$$Q: x \sim \mathcal{N}(\mu_2, \sigma_2^2)$$

$$KL(P||Q) = \frac{1}{2} \left[\frac{(\mu_2 - \mu_1)^2}{\sigma_2^2} + \frac{\sigma_1^2}{\sigma_2^2} - \log \frac{\sigma_1^2}{\sigma_2^2} - 1 \right] \quad (\text{a.2})$$

Now let us see the version for multivariate normal distributions (see proof in [Statproofbook 2021B]), using the notation $|\mathbf{M}|$ for the determinant of a matrix \mathbf{M} , and the notation $tr(\mathbf{M})$ for the trace of a matrix \mathbf{M} . The variable x is a vector of dimension n distributed normally with mean vector μ (analogously of dimension n) and covariance matrix Σ (of dimension $n \times n$).:

$$P: x \sim \mathcal{N}(\mu_1, \Sigma_1)$$

$$Q: x \sim \mathcal{N}(\mu_2, \Sigma_2)$$

$$KL(P||Q) = \frac{1}{2} \left[(\mu_2 - \mu_1)^T \Sigma_2^{-1} (\mu_2 - \mu_1) + tr(\Sigma_2^{-1} \Sigma_1) - \log \frac{|\Sigma_1|}{|\Sigma_2|} - n \right] \quad (\text{a.3})$$

When the two covariance matrices Σ_1 and Σ_2 are diagonal, i.e. the covariances between variables are zero, the computation of eq. a.3 is equivalent to compute eq. a.2 for each variable and sum all the results together. In Deep Reinforcement Learning usually the covariances between action dimensions are supposed to be zero, so this simplification may be applied (either the variances for each action are fixed, with covariances equal to zero, or the variances are outputted by an output neuron of the policy neural network that outputs only a standard deviation -or a log-standard deviation- for each action dimensions and no covariances).

A.2 Total Variation Distance

It is another way to compute how different two distributions are. Given two distributions P and Q , the Total Variation distance $D_{TV}(P||Q)$ is defined as:

$$\text{continuous } D_{TV}(P||Q) = \frac{1}{2} \int_x |P(x) - Q(x)|$$

$$\text{discrete } D_{TV}(P||Q) = \frac{1}{2} \sum_x |P(x) - Q(x)|$$

(a.4)




Figure a.1

Given two distributions P and Q , here depicted in their probability density functions, we notice that by definition the Total Variation Distance is equal to the sum of all areas of P above the area of Q , plus the sum of all areas of Q above the area of P , all divided by 2. That is, in the figure, the sum of the red-only area and the two blue-only areas, excluding the purple area, all divided by two.

If you watch Figure a.1 you can see that the Total Variation Distance is equal to the sum of all areas in which either distribution overwhelms the other (without the overwhelmed areas), divided by two. This is by definition.

But, since the area of every distribution is always equal to one, if a distribution overwhelms the other for a certain area, it must be overwhelmed for an equivalent area in a different subset of the domain. This means that the Total Variation Distance is also equal to the area of any of the distributions that is overwhelming the other, without summing the other, and without dividing by two.

Hence, since the Total Variation Distance may be computed measuring only the area of a distribution that overwhelms the area of the other distribution, another definition is possible.

Watching Figure a.1 if we name the red distribution P and the blue distribution Q , we can see that where the red distribution is above the blue one, we named the support as A (the subset of the domain where $P > Q$), and where the blue distribution is above the red one we named the support as B (the subset of the domain where $Q > P$). We just noticed that the Total Variation Distance may be computed measuring the area where the probability density function of a distribution is greater than the pdf of the other.

Hence, another definition of Total Variation distance may be based on finding the subset A where $P > Q$, and then computing the difference between the probabilities of P and Q on that subset.

That is equivalent to say that we want to find the subset A that maximizes the difference in probabilities between P and Q . If we want to use statistical language, we may call the domain of P as “*sample space*”, hence the subset A is an “*event*”, and any $x \in A$ is an “*outcome*” included in that event. So, to compute the Total Variation distance we find the event for which the difference between the probabilities of the two distributions is greater, and then compute that difference.

$$D_{TV}(P||Q) = \max_{A \subseteq S} (Prob_P(A) - Prob_Q(A))$$

with S support of P

(a.5)

That is equivalent to:

$$\begin{aligned}
\text{continuous } D_{TV}(P||Q) &= \max_{A \subseteq S} \left(\int_{x \in A} P(x) - \int_{x \in A} Q(x) \right) \\
\text{discrete } D_{TV}(P||Q) &= \max_{A \subseteq S} \left(\sum_{x \in A} P(x) - \sum_{x \in A} Q(x) \right) \\
&\text{with } S \text{ support of } P
\end{aligned} \tag{a.6}$$

Those can be equivalently written as:

$$\begin{aligned}
\text{continuous } D_{TV}(P||Q) &= \max_{A \subseteq S} \left(\int_{x \in A} (P(x) - Q(x)) \right) \\
\text{discrete } D_{TV}(P||Q) &= \max_{A \subseteq S} \left(\sum_{x \in A} (P(x) - Q(x)) \right) \\
&\text{with } S \text{ support of } P
\end{aligned} \tag{a.7}$$

The Total Variation distance has the following relationship with Kullback-Leibler divergence (see [Pollard 2000] Ch.3):

$$D_{TV}(P||Q)^2 \leq D_{KL}(P||Q) \tag{a.8}$$

A.3 Score (of the Log-Likelihood)

It is the gradient of the log-likelihood function with respect to the parameters vector. Hence, given a certain parameters vector, the score denotes the steepness of the log-likelihood at that point in parameters space, or in other terms how much the log-likelihood would change for an infinitesimal change in the parameters.

Naming θ the parameters vector:

$$s(\theta) = \frac{\partial \log \mathcal{L}(\theta)}{\partial \theta} \quad (\text{a.9})$$

Since the log-likelihood is a function of the samples X , we can make that explicit:

$$s(\theta; x) = \frac{\partial \log \mathcal{L}(\theta; x)}{\partial \theta} \quad (\text{a.10})$$

If θ are the true parameters of the distribution $P_\theta(x)$ of samples X , the expectation of the score with respect of the distribution of X is the zero vector:

$$E_{x \sim P_\theta(x)}[s(\theta; x)] = \mathbf{0} \quad (\text{a.11})$$

To show it:

$$E_{x \sim P_\theta(x)}[s(\theta; x)] = \int_x P_\theta(x) \frac{\partial \log \mathcal{L}(\theta; x)}{\partial \theta}$$

*since θ are the true parameters,
 $P_\theta(x)$ is $\mathcal{L}(\theta; x)$. Substitute it, and apply log – derivative trick (eq. 3.3)*

$$\begin{aligned} &= \int_x P_\theta(x) \frac{1}{P_\theta(x)} \frac{\partial P_\theta(x)}{\partial \theta} \\ &= \int_x \frac{\partial P_\theta(x)}{\partial \theta} \end{aligned}$$

*under regularity conditions it is possible to interchange the derivative and the integral
 (Leibniz Integral Rule)*

$$= \frac{\partial}{\partial \theta} \int_x P_\theta(x) = \frac{\partial}{\partial \theta} 1 = 0$$

■
 (a.12)

A.4 Fisher Information Matrix

The Fisher Information is the variance of the score with respect to the distribution of the samples. If the parameters θ of the distribution are more than one, the Fisher Information is a matrix of covariances of the elements of the score vector, and we name it \mathbf{F}_{P_θ} .

Assuming that θ are the true parameters of the distribution of X :

$$\begin{aligned}\mathbf{F}_{P_\theta} &= \text{Var}_{x \sim P_\theta(x)}[s(\theta; x)] \\ &= E_{x \sim P_\theta(x)} \left[\left(s(\theta; x) - E_{x \sim P_\theta(x)}[s(\theta; x)] \right) \left(s(\theta; x) - E_{x \sim P_\theta(x)}[s(\theta; x)] \right)^T \right]\end{aligned}$$

from equation a.11 we know that $E_{x \sim P_\theta(x)}[s(\theta; x)] = \mathbf{0}$, hence:

$$\mathbf{F}_{P_\theta} = E_{x \sim P_\theta(x)}[s(\theta; x) s(\theta; x)^T] \quad (\text{a.13})$$

That may also be written as (recall that $s(\theta; x) = \frac{\partial \log \mathcal{L}(\theta; x)}{\partial \theta} = \nabla_\theta \log P_\theta(x)$ by definition):

$$\mathbf{F}_{P_\theta} = E_{x \sim P_\theta(x)}[\nabla_\theta \log P_\theta(x) \nabla_\theta \log P_\theta(x)^T] \quad (\text{a.14})$$

Hence, each matrix element $\mathbf{F}_{P_\theta}[i, j]$ has the following form:

$$\mathbf{F}_{P_\theta}[i, j] = E_{x \sim P_\theta(x)} \left[\left(\frac{\partial}{\partial \theta_i} \log P_\theta(x) \right) \left(\frac{\partial}{\partial \theta_j} \log P_\theta(x) \right) \right] \quad (\text{a.15})$$

It is clearly a square symmetric matrix, and it is also positive semi-definite [Watanabe 2009]. It becomes evident that such a matrix may be estimated by samples x_i which follow the P_θ distribution :

$$\widehat{\mathbf{F}}_{P_\theta} = \frac{1}{N} \sum_{i=1}^N \nabla_\theta \log P_\theta(x_i) \nabla_\theta \log P_\theta(x_i)^T \quad (\text{a.16})$$

(That is actually what we do when approximating the Fisher Information Matrix for the Natural Policy Gradient in Ch.6, equation 6.3, where instead of P_θ there is the policy function π_θ .).

Intuitively and informally speaking, the Fisher Information Matrix measures how much information the random variable x carries about any parameter in θ .

A.5 Fisher Information Matrix equivalence to negative expectation of Hessian Matrix of log-probability

Now, for the next part we need to use the concept of the Hessian Matrix. Recall that the Hessian is the square matrix of second order partial derivatives of a scalar function (shortly, the Hessian is the Jacobian of the gradient). In our case the Hessian would be with respect to the parameters θ of a probability function $P_\theta(x)$, that is $\frac{\partial^2 P_\theta(x)}{\partial_i \partial_j}$. To simplify reading I will use the symbol $H_{P_\theta(x)}$ for it, while for the Jacobian instead I will use the symbol $\mathbf{J}()$.

If the log-likelihood function is twice differentiable with respect to θ , and under certain regularity conditions it can be shown that the Fisher Information Matrix is equivalent to the negative expected value of the Hessian matrix of the log-likelihood.

$$\mathbf{F}_{P_\theta} = -E_{x \sim P_\theta(x)} [H_{\log P_\theta(x)}] \quad (\text{a.17})$$

So, each matrix element $F_{P_\theta}[i, j]$ is:

$$F_{P_\theta}[i, j] = -E_{x \sim P_\theta(x)} \left[\left(\frac{\partial^2}{\partial \theta_i \partial_j} \log P_\theta(x) \right) \right] \quad (\text{a.18})$$

To show it let us start from the Hessian of the log-likelihood, following [Kristiadi 2018]:

$$H_{\log P_\theta(x)} = \mathbf{J}(\nabla_\theta \log P_\theta(x))$$

$$= \mathbf{J} \left(\frac{\nabla_{\theta} P_{\theta}(x)}{P_{\theta}(x)} \right)$$

applying the quotient rule for derivatives

$$\begin{aligned} &= \frac{H_{P_{\theta}(x)} P_{\theta}(x) - \nabla_{\theta} P_{\theta}(x) \nabla_{\theta} P_{\theta}(x)^T}{P_{\theta}(x) P_{\theta}(x)} \\ &= \frac{H_{P_{\theta}(x)}}{P_{\theta}(x)} - \frac{\nabla_{\theta} P_{\theta}(x)}{P_{\theta}(x)} \frac{\nabla_{\theta} P_{\theta}(x)^T}{P_{\theta}(x)} \end{aligned}$$

let us take the expectation with respect to x distributed by $P_{\theta}(x)$

$$\begin{aligned} E_{x \sim P_{\theta}(x)} [H_{\log P_{\theta}(x)}] &= E_{x \sim P_{\theta}(x)} \left[\frac{H_{P_{\theta}(x)}}{P_{\theta}(x)} - \frac{\nabla_{\theta} P_{\theta}(x)}{P_{\theta}(x)} \frac{\nabla_{\theta} P_{\theta}(x)^T}{P_{\theta}(x)} \right] \\ &= \int_x \frac{H_{P_{\theta}(x)}}{P_{\theta}(x)} P_{\theta}(x) - E_{x \sim P_{\theta}(x)} \left[\frac{\nabla_{\theta} P_{\theta}(x)}{P_{\theta}(x)} \frac{\nabla_{\theta} P_{\theta}(x)^T}{P_{\theta}(x)} \right] \end{aligned}$$

under regularity conditions it is possible to interchange the derivative and the integral

(Leibniz Integral Rule)

$$= H_{\int_x P_{\theta}(x)} - E_{x \sim P_{\theta}(x)} \left[\frac{\nabla_{\theta} P_{\theta}(x)}{P_{\theta}(x)} \frac{\nabla_{\theta} P_{\theta}(x)^T}{P_{\theta}(x)} \right]$$

the integral of a distribuion is $= 1$ and its derivative is $= 0$, hence $H_{\int_x P_{\theta}(x)} = \mathbf{0}$

$$\begin{aligned} &= -E_{x \sim P_{\theta}(x)} \left[\frac{\nabla_{\theta} P_{\theta}(x)}{P_{\theta}(x)} \frac{\nabla_{\theta} P_{\theta}(x)^T}{P_{\theta}(x)} \right] \\ &= -E_{x \sim P_{\theta}(x)} [\nabla_{\theta} \log P_{\theta}(x) \nabla_{\theta} \log P_{\theta}(x)^T] \\ &= -\mathbf{F}_{P_{\theta}} \end{aligned}$$

■

(a.19)

Since the Hessian Matrix is commonly used to study the curvature of a function at a critical point, the equivalence above tells us that the Fisher Information Matrix contains information about the curvature of the expectation of the log-likelihood.

A.6 Kullback-Leibler divergence approximation by second order Taylor expansion using Fisher Information Matrix

The equivalence of the Fisher Information Matrix with the negative expected value of the Hessian matrix of the log-likelihood (see appendix A.5) allows us to compute an approximation of the Kullback-Leibler divergence between two distributions where one is the “perturbed” version of the other (i.e. they share the same parametric form, and the parameters vector of the second distribution is obtained adding a small vector to the parameters of the first), using second order Taylor expansion for a part of the KL divergence formula.

Recall that second order Taylor series expansion for a function $f(\theta)$ is an approximation used to evaluate the function $f(\theta)$ around a certain point θ_0 , at the point $\theta_0 + \delta$:

$$f(\theta_0 + \delta) \approx f(\theta_0) + \nabla_{\theta} f(\theta_0)^T \delta + \frac{1}{2} \delta^T (\nabla_{\theta}^2 f(\theta_0)) \delta \quad (\text{a.20})$$

Consider having the distributions $P_{\theta}(x) = f(x; \theta)$ and $Q_{\theta, \delta}(x) = f(x; \theta + \delta)$, that means that both share the same function f , and the parameters vector of P is θ , while the parameters vector of Q is $(\theta + \delta)$, with small δ . In other terms, Q is a perturbed version of P . For instance in Reinforcement Learning P may be the old policy function and Q the new policy function obtained by an optimization step on P .

I follow the proof by [Ratliff 2013]:

$$\begin{aligned} D_{KL}(P||Q) &= D_{KL}(P_{\theta}(x)||Q_{\theta, \delta}(x)) = D_{KL}(f(x; \theta)||f(x; \theta + \delta)) = \\ &= \int_x f(x; \theta) \log \frac{f(x; \theta)}{f(x; \theta + \delta)} \\ &= \int_x f(x; \theta) \log f(x; \theta) - \int_x f(x; \theta) \log f(x; \theta + \delta) \end{aligned} \quad (\text{a.21})$$

Now we apply the second order Taylor expansion only to $\log f(x; \theta + \delta)$:

$$\log f(x; \theta + \delta) \approx \log f(x; \theta) + \nabla_{\theta} \log f(x; \theta)^T \delta + \frac{1}{2} \delta^T (\nabla_{\theta}^2 \log f(x; \theta)) \delta$$

$$= \log f(x; \theta) + \left(\frac{\nabla_{\theta} f(x; \theta)}{f(x; \theta)} \right)^T \delta + \frac{1}{2} \delta^T (\nabla_{\theta}^2 \log f(x; \theta)) \delta \quad (\text{a.22})$$

Now we plug a.22 into a.21:

$$\begin{aligned} D_{KL}(f(x; \theta) || f(x; \theta + \delta)) &\approx \\ &\approx \int_x f(x; \theta) \log f(x; \theta) - \int_x f(x; \theta) \left(\log f(x; \theta) + \left(\frac{\nabla_{\theta} f(x; \theta)}{f(x; \theta)} \right)^T \delta + \frac{1}{2} \delta^T (\nabla_{\theta}^2 \log f(x; \theta)) \delta \right) \\ &= \int_x f(x; \theta) \log \frac{f(x; \theta)}{f(x; \theta)} - \left(\int_x \nabla_{\theta} f(x; \theta) \right)^T \delta - \frac{1}{2} \delta^T \left(\int_x f(x; \theta) \nabla_{\theta}^2 \log f(x; \theta) \right) \delta \end{aligned} \quad (\text{a.23})$$

Now: $\int_x f(x; \theta) \log \frac{f(x; \theta)}{f(x; \theta)}$ is equal to 0 (easy to do the math or to see it as a KL divergence of two equal distributions).

Also $\left(\int_x \nabla_{\theta} f(x; \theta) \right)^T$ is equal to 0: it is possible to swap the gradient and derivative signs (under regularities) to obtain $\left(\nabla_{\theta} \int_x f(x; \theta) \right)^T$.

The integral of any distribution is 1, so $\nabla_{\theta} 1 = 0$.

So it turns out:

$$\begin{aligned} D_{KL}(f(x; \theta) || f(x; \theta + \delta)) &\approx -\frac{1}{2} \delta^T \left(\int_x f(x; \theta) \nabla_{\theta}^2 \log f(x; \theta) \right) \delta \\ &= -\frac{1}{2} \delta^T \left(\int_x f(x; \theta) H_{\log f(x; \theta)} \right) \delta \\ &= -\frac{1}{2} \delta^T E_{x \sim f(x; \theta)} [H_{\log f(x; \theta)}] \delta \end{aligned} \quad (\text{a.24})$$

Now, we already know that the expected value of the Hessian matrix of the log-likelihood is equal to the negative of Fisher Information Matrix (eq. a.17). So we substitute it:

$$D_{KL}(f(x; \theta) || f(x; \theta + \delta)) \approx \frac{1}{2} \delta^T \mathbf{F}_{P_\theta} \delta \quad (\text{a.25})$$

A.7 Relationship between the Hessian of Kullback-Leibler divergence and Fisher Information Matrix

Again, consider having two distributions with the same function but different parameters $P_\theta(x)$ and $P_{\theta'}(x)$.

Then the gradient of Kullback-Leibler divergence between P_θ and $P_{\theta'}$ with respect to θ' would be:

$$\nabla_{\theta'} D_{KL}(P_\theta || P_{\theta'}) = \nabla_{\theta'} \int_x P_\theta(x) \log \frac{P_\theta(x)}{P_{\theta'}(x)} \quad (\text{a.26})$$

That gradient is a vector of partial derivatives with respect to all parameters of θ' . To denote the partial derivative with respect to parameter j of parametrization θ' I use the symbol $\partial_{\theta'j} d$.

We have that:

$$\partial_{\theta'j} d = \partial_{\theta'j} \int_x P_\theta(x) \log \frac{P_\theta(x)}{P_{\theta'}(x)}$$

under regularity conditions the derivative and integral symbol can be exchanged

$$\begin{aligned} &= \int_x P_\theta(x) \partial_{\theta'j} \left(\log \frac{P_\theta(x)}{P_{\theta'}(x)} \right) \\ &= \int_x P_\theta(x) \frac{P_{\theta'}(x)}{P_\theta(x)} P_\theta(x) \frac{-1}{P_{\theta'}(x)^2} \partial_{\theta'j} P_{\theta'}(x) \\ &= - \int_x \frac{P_\theta(x)}{P_{\theta'}(x)} \partial_{\theta'j} P_{\theta'}(x) \end{aligned} \quad (\text{a.27})$$

Now let us take the derivative of eq. a.27, that is the second derivative of the KL divergence (computing at each parameter i, j): $\partial_{\theta'i} \partial_{\theta'j} d$.

$$\partial_{\theta' i} \partial_{\theta' j} d = \partial_{\theta' i} \int_x -\frac{P_{\theta}(x)}{P_{\theta'}(x)} \partial_{\theta' j} P_{\theta'}(x)$$

under regularity conditions the derivative and integral symbol can be exchanged

$$\begin{aligned} &= \int_x -P_{\theta}(x) \partial_{\theta' i} \left(\frac{\partial_{\theta' j} P_{\theta'}(x)}{P_{\theta'}(x)} \right) \\ &= \int_x -P_{\theta}(x) \left(\frac{\partial_{\theta' i} \partial_{\theta' j} P_{\theta'}(x) P_{\theta'}(x) - (\partial_{\theta' i} P_{\theta'}(x)) (\partial_{\theta' j} P_{\theta'}(x))}{P_{\theta'}(x)^2} \right) \\ &= \int_x P_{\theta}(x) \left(\frac{(\partial_{\theta' i} P_{\theta'}(x)) (\partial_{\theta' j} P_{\theta'}(x))}{P_{\theta'}(x)^2} - \frac{\partial_{\theta' i} \partial_{\theta' j} P_{\theta'}(x)}{P_{\theta'}(x)} \right) \end{aligned}$$

now, if we evaluate that integral at the parameter point $\theta' = \theta$ we obtain :

$$\begin{aligned} \partial_{\theta' i} \partial_{\theta' j} d \Big|_{\theta' = \theta} &= \int_x P_{\theta}(x) \left(\frac{(\partial_{\theta' i} P_{\theta'}(x)) (\partial_{\theta' j} P_{\theta'}(x))}{P_{\theta'}(x)^2} - \frac{\partial_{\theta' i} \partial_{\theta' j} P_{\theta'}(x)}{P_{\theta'}(x)} \right) \Big|_{\theta' = \theta} \\ &= \int_x \frac{(\partial_{\theta' i} P_{\theta'}(x)) (\partial_{\theta' j} P_{\theta'}(x))}{P_{\theta'}(x)} - \partial_{\theta' i} \partial_{\theta' j} P_{\theta'}(x) \Big|_{\theta' = \theta} \end{aligned}$$

it is easy to see that $\int_x \partial_{\theta' i} \partial_{\theta' j} P_{\theta'}(x) = \partial_{\theta' i} \partial_{\theta' j} \int_x P_{\theta'}(x) = \partial_{\theta' i} \partial_{\theta' j} 1 = 0$

$$\begin{aligned} &= \int_x \frac{(\partial_{\theta' i} P_{\theta'}(x)) (\partial_{\theta' j} P_{\theta'}(x))}{P_{\theta'}(x)} \Big|_{\theta' = \theta} \\ &= \int_x \frac{P_{\theta'}(x)}{P_{\theta'}(x)} \frac{(\partial_{\theta' i} P_{\theta'}(x)) (\partial_{\theta' j} P_{\theta'}(x))}{P_{\theta'}(x)} \Big|_{\theta' = \theta} \\ &= \int_x P_{\theta'}(x) (\partial_{\theta' i} \log P_{\theta'}(x)) (\partial_{\theta' j} \log P_{\theta'}(x)) \Big|_{\theta' = \theta} \\ &= E_{x \sim P_{\theta}(x)} [(\partial_{\theta i} \log P_{\theta}(x)) (\partial_{\theta j} \log P_{\theta}(x))] \end{aligned}$$

$$= \mathbf{F}_{\mathbf{P}_{\theta}}[i, j] \quad (\text{see eq. a. 15})$$

$$\Rightarrow H(D_{KL}(P_{\theta}||P_{\theta'}))_{\theta'}|_{\theta' = \theta} = \mathbf{F}_{P_{\theta}} \quad (\text{a.28})$$

That means that given two distributions of the same family $P_{\theta}(x)$ and $P_{\theta'}(x)$, the Hessian with respect to θ' of Kullback-Leibler divergence from P_{θ} to $P_{\theta'}$, evaluated at $\theta' = \theta$ is equal to the Fisher Information Matrix of $P_{\theta}(x)$.

A.8 Information Entropy

Information Entropy, ideated by [Shannon 1948], is a measure of the “randomness” or “uncertainty” of a discrete random variable, in the sense of how much “surprise” we expect when trying to guess the value of a sample from it.

Formally, Information Entropy of a discrete random variable X with probability mass function $p(x)$, is the functional:

$$H(X) = - \sum_{x \in \mathcal{X}} p(x) \log_b p(x) \quad (\text{a.29})$$

Where b is the base of logarithm, and can assume different values: when $b = 2$ the units of Entropy are named “bits”, when $b = e$ (Euler’s number) the units of Entropy are named “nats”, when $b = 10$ the units of Entropy are named “Dits” or “Hartleys” or “Bans”.

It is easy to see that the Information Entropy $H(X)$ cannot be negative. It is also easy to see that in the simplest case, when we have only one possible value for variable X , the Entropy $H(X) = 0$, and it grows if we have a random variable with more possible values. For any fixed number of values of the random variable, the entropy is maximum when all values have the same probability mass. So, we can think that $H(X)$ is some sort of indicator of the expected information contained in a sample of random variable X . The more X is difficult to be randomly guessed, the more information a sample of it contains.

The previous formula may also be written as:

$$H(X) = E_{p(x)}[-\log_b p(x)]$$

(a.30)

But beware of the fact that in this latter formulation it is not explicit that we are dealing with a discrete distribution. In fact, when that formula is generalized to continuous distribution, $p(x)$ would be a probability density distribution, and $E_{p(x)}[-\log_b p(x)]$ would be equal to $-\int_{x \in \mathcal{X}} p(x) \log_b p(x)$, but this is not the same thing as the discrete case, since it can become negative !

For instance, in an example from [Rioul 2018], if we had a continuous random variable x , uniformly distributed in the interval (a, d) , we would have $p(x) = \frac{1}{(d-a)}$, and:

$$\begin{aligned}
 E_{p(x)}[-\log_b p(x)] &= \int_{x \in \mathcal{X}} -p(x) \log_b p(x) \\
 &= \int_{x \in \mathcal{X}} p(x) \log_b \frac{1}{p(x)} \\
 &= \int_{x \in \mathcal{X}} p(x) \log_b (d - a) \\
 &= \log_b (d - a)
 \end{aligned}
 \tag{a.31}$$

And in case $(d - a) < 1$ we would have a negative value. Also, taking the limits of a.31 as the interval $(d - a)$ tends to zero, the result is $= -\infty$. This not only is very different from the discrete variable case, but also loses its sense as a measure of information.

So, we cannot use Shannon's Information Entropy with continuous random variables and expect to have the same meaning.

Because of these differences, when X is a continuous random variable, the expression $E[-\log_b p(x)]$ is not properly called “*information entropy*” but rather “*differential entropy*” or “*continuous entropy*” and denoted with a lowercase $h(X)$:

with X continuous random variable

$$h(X) = E_{p(x)}[-\log_b p(x)] = - \int_{x \in \mathcal{X}} p(x) \log_b p(x)$$

(a.32)

Another example of the difference of behavior and meaning of $E_{p(x)}[-\log_b p(x)]$ between discrete and continuous random variables is in the case of change of variable. Consider having a discrete random variable $Y = T(X)$ which is the result of some invertible transformation T of the discrete random variable X . Let us refer to the probability mass of y as $q(y)$. The probability mass function $p(x)$ contains the same probability mass values of $q(y)$, the only difference is that $p(x)$ and $q(y)$ associate each probability mass value to different inputs. Having the same probability mass values implies that they have the same entropy, i.e. :

$$E_{q(y)}[-\log_b q(y)] = - \sum_{y \in \mathcal{Y}} q(y) \log_b q(y) = - \sum_{x \in \mathcal{X}} p(x) \log_b p(x) = E_{p(x)}[-\log_b p(x)]$$

(a.33)

Or, briefly, $H(Y) = H(X)$.

Things are different if X and Y are both continuous random variables, and $Y = T(X)$, with T some invertible transformation (whose inverted is denoted as $X = T^{-1}(Y)$) . In this case the invertible transformation must be a diffeomorphism . The derivative of $T(X)$ with respect to x is $T'(X) = \frac{dy}{dx}$. The density function of x is $p(x)$, and the density function of y is $q(y)$. We can denote the respective cumulative distribution functions as $P(x)$ and $Q(y)$.

With a change of variable from X to Y we must remember that (see [Casella and Berger 2002] theorem 2.1.3 and 2.1.5, or [Psu 2024]):

if $T(x)$ is an increasing function on \mathcal{X}

$$Q(y) = P(T^{-1}(y))$$

(a.34)

if $T(x)$ is a decreasing function on \mathcal{Y}

$$Q(y) = 1 - P(T^{-1}(y))$$

(a.35)

Equation a.34 and a.35 imply, because of the chain of rule of calculus:

$$q(y) = \frac{d}{dy} Q(y) = p(T^{-1}(y)) \left| \frac{d}{dy} T^{-1}(y) \right|$$

(a.36)

$$p(x) = \frac{d}{dx} P(x) = q(T(x)) \left| \frac{d}{dx} T(x) \right|$$

(a.37)

From a.36 substituting $x = T^{-1}(y)$ we can obtain:

$$q(y) = p(x) \left| \frac{dx}{dy} \right|$$

(a.38)

And analogously from a.37:

$$p(x) = q(y) \left| \frac{dy}{dx} \right|$$

(a.39)

Now to show the proof I will not use the compact integral notation as before, but instead I will write the full notation with differentials dx and dy .

$$E_{q(y)}[-\log_b q(y)] =$$

$$\int_{-\infty}^{\infty} q(y) \log_b \left(\frac{1}{q(y)} \right) dy =$$

$$\int_{-\infty}^{\infty} p(x) \left| \frac{dx}{dy} \right| \log_b \left(\frac{\left| \frac{dy}{dx} \right|}{p(x)} \right) dy$$

(a.40)

$$\text{if } T(x) \text{ is an increasing function on } \mathcal{X} \Rightarrow \left| \frac{dy}{dx} \right| = \frac{dy}{dx} \text{ and } \left| \frac{dx}{dy} \right| = \frac{dx}{dy}$$

$$= \int_{-\infty}^{\infty} p(x) \frac{dx}{dy} \log_b \left(\frac{\frac{dy}{dx}}{p(x)} \right) dy$$

$$\begin{aligned}
&= \int_{-\infty}^{\infty} p(x) \log_b \left(\frac{1}{p(x)} \right) dx + \int_{-\infty}^{\infty} p(x) \log_b \left(\frac{dy}{dx} \right) dx \\
&= \int_{-\infty}^{\infty} p(x) \log_b \left(\frac{1}{p(x)} \right) dx + \int_{-\infty}^{\infty} p(x) \log_b (T'(x)) dx \\
&= E_{p(x)} [-\log_b p(x)] + E_{p(x)} [\log_b T'(x)] \\
&\Rightarrow \\
&h(Y) = h(X) + E_{p(x)} [\log_b T'(x)]
\end{aligned}$$

(a.41)

if instead $T(x)$ is a decreasing function on $\mathcal{X} \Rightarrow \left| \frac{dy}{dx} \right| = -\frac{dy}{dx}$ and $\left| \frac{dx}{dy} \right| = -\frac{dx}{dy}$:

$$\begin{aligned}
&= - \int_{-\infty}^{\infty} p(x) \frac{dx}{dy} \log_b \left(\frac{-\frac{dy}{dx}}{p(x)} \right) dy \\
&= - \int_{-\infty}^{\infty} p(x) \log_b \left(\frac{1}{p(x)} \right) dx - \int_{-\infty}^{\infty} p(x) \log_b \left(-\frac{dy}{dx} \right) dx \\
&= - \int_{-\infty}^{\infty} p(x) \log_b \left(\frac{1}{p(x)} \right) dx - \int_{-\infty}^{\infty} p(x) \log_b (-T'(x)) dx \\
&= -E_{p(x)} [-\log_b p(x)] - E_{p(x)} [\log_b (-T'(x))] \\
&\Rightarrow
\end{aligned}$$

$$h(Y) = -h(X) - E_{p(x)} [\log_b (-T'(x))]$$

(a.42)

What we find with a.41 and a.42 shows that the meaning of information entropy for discrete variables $H(X)$ is changed when we are dealing with continuous information entropy $h(X)$.

It is easy to see that $h(Y) = h(X)$ only if $T(X)$ is an increasing function on \mathcal{X} and at the same time the derivative $T'(X)$ is one, that happens only when $T(X) = x + c$, with c constant. This is equivalent to say that differential entropy is translation invariant.

A.9 Differential Entropy of the Gaussian

An important property of the Normal probability distribution is that: for any fixed variance, among all possible continuous probability distributions, the Normal/Gaussian distribution has the greatest Differential Entropy.

Proof:

Let us denote the probability density of a gaussian with mean μ and variance σ^2 as $g(x) =$

$$\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}.$$

Its differential entropy $h(g)$ is:

$$\begin{aligned} & - \int_{-\infty}^{\infty} g(x) \log_b(g(x)) dx \\ &= - \int_{-\infty}^{\infty} g(x) \log_b \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \right) dx \\ &= - \int_{-\infty}^{\infty} g(x) \log_b \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) dx - \int_{-\infty}^{\infty} g(x) \left(-\frac{(x-\mu)^2}{2\sigma^2} \right) \log_b(e) dx \\ &= - \int_{-\infty}^{\infty} g(x) \log_b \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) dx - \log_b(e) \int_{-\infty}^{\infty} g(x) \left(-\frac{(x-\mu)^2}{2\sigma^2} \right) dx \\ &= - \int_{-\infty}^{\infty} g(x) \log_b \left((2\pi\sigma^2)^{-\frac{1}{2}} \right) dx - \log_b(e) \int_{-\infty}^{\infty} g(x) \left(-\frac{(x-\mu)^2}{2\sigma^2} \right) dx \\ &= \int_{-\infty}^{\infty} g(x) \frac{1}{2} \log_b(2\pi\sigma^2) dx - \log_b(e) \int_{-\infty}^{\infty} g(x) \left(-\frac{(x-\mu)^2}{2\sigma^2} \right) dx \end{aligned}$$

$$\begin{aligned}
&= \frac{1}{2} \log_b(2\pi\sigma^2) + \log_b(e) \frac{1}{2\sigma^2} \int_{-\infty}^{\infty} g(x)(x-\mu)^2 dx \\
&= \frac{1}{2} \log_b(2\pi\sigma^2) + \log_b(e) \frac{\sigma^2}{2\sigma^2} \\
&= \frac{1}{2} (\log_b(2\pi\sigma^2) + \log_b(e)) \\
&= \frac{1}{2} (\log_b(2\pi\sigma^2 e))
\end{aligned}
\tag{a.43}$$

Now, consider having another probability distribution whose pdf is denoted with $f(x)$, which has the same mean μ and variance σ^2 as the gaussian. Actually, since differential entropy is translation invariant (as seen before, in appendix A.7), if the gaussian had a different mean it could anyway be translated to have the same mean, and its differential entropy would be unchanged. For this reason, we are only putting one assumption: that $f(x)$ and $g(x)$ have the same variance σ^2 (we need to prove that “for any fixed variance, the Gaussian distribution has the greatest Differential Entropy”).

The Kullback-Leibler divergence of $f(x)$ from $g(x)$ is :

$$\begin{aligned}
D_{KL}(f||g) &= \int_{-\infty}^{\infty} f(x) \log_b \left(\frac{f(x)}{g(x)} \right) dx \\
&= -h(f) - \int_{-\infty}^{\infty} f(x) \log_b(g(x)) dx \\
&= -h(f) - \int_{-\infty}^{\infty} f(x) \log_b \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \right) dx
\end{aligned}
\tag{a.44}$$

The second term is $-\int_{-\infty}^{\infty} f(x) \log_b \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \right) dx$, can be simplified analogously to what we did with $h(g)$:

$$-\int_{-\infty}^{\infty} f(x) \log_b \left(\frac{1}{\sqrt{2\pi\sigma^2}} e^{-\frac{(x-\mu)^2}{2\sigma^2}} \right) dx =$$

$$= - \int_{-\infty}^{\infty} f(x) \log_b \left(\frac{1}{\sqrt{2\pi\sigma^2}} \right) dx - \int_{-\infty}^{\infty} f(x) \left(-\frac{(x-\mu)^2}{2\sigma^2} \right) \log_b(e) dx$$

$$= \frac{1}{2} \log_b(2\pi\sigma^2) + \log_b(e) \frac{1}{2\sigma^2} \int_{-\infty}^{\infty} f(x)(x-\mu)^2 dx$$

since $f(x)$ has same mean μ and variance σ^2 as the $g(x) \Rightarrow \int_{-\infty}^{\infty} f(x)(x-\mu)^2 dx = \sigma^2 \Rightarrow$

$$= \frac{1}{2} \log_b(2\pi\sigma^2) + \log_b(e) \frac{\sigma^2}{2\sigma^2}$$

$$= \frac{1}{2} (\log_b(2\pi\sigma^2 e))$$

$$= h(g)$$

(a.45)

Now plugging a.45 into a.44:

$$D_{KL}(f||g) = h(g) - h(f)$$

(a.46)

We know that Kullback-Leibler divergence is always non-negative. So:

$$h(g) - h(f) \geq 0$$

$$h(g) \geq h(f)$$

(a.47)

This proves that for a fixed variance σ^2 , the differential $h(g)$ entropy of a gaussian g is always greater than, or equal to, the differential entropy $h(f)$ of any other distribution f .

Let us also briefly enunciate the differential entropy when the gaussian is multivariate [Statproofbook 2021C], using the notation $|\mathbf{M}|$ for the determinant of a matrix \mathbf{M} , and the notation $tr(\mathbf{M})$ for the trace of a matrix \mathbf{M} . The variable x is now a vector of dimension n distributed normally with mean vector μ (analogously of dimension n) and covariance matrix Σ (of dimension $n \times n$).:

$$g: x \sim \mathcal{N}(\mu, \Sigma)$$

$$h(g) = \frac{n}{2} \log_b(2\pi e) + \frac{1}{2} \log_b(|\Sigma|)$$

(a.48)

A.10 Mutual Information

Consider two random variables X and Y , having respective marginal distributions $P(x)$ and $Q(y)$, and joint distribution $R(x, y)$.

Mutual Information is defined as:

$$I(X; Y) = D_{KL}(R(x, y) || P(x)Q(y))$$

(a.48)

That is the Kullback Leibler divergence from the joint distribution to the hypothetical joint distribution that would exist if the two variables were independent to each other (that is, the plain multiplication of the two marginal distributions).

It assesses the quantity of information about a variable that is obtainable observing the other. It is more versatile than Pearson's correlation coefficient because the latter is limited to real valued variables and linear correlations.

When the joint distribution coincides with the product of the two marginals, i.e. when the two variables are independent, the mutual information equals zero (it becomes a KL divergence from a distribution to itself).

It is easy to see that mutual information is symmetric with respect to the order of the variables.

Appendix B: Conjugate Gradient Algorithm

This Appendix provides a sketch of the Conjugate Gradient Algorithm [Hestenes and Stiefel 1952]. It will not be described in full detail, just what it is enough to understand its usage in Reinforcement Learning (as in Trust Region Policy Optimization). The reader is encouraged to find more on the topic in literature, starting for instance from the guides consulted to prepare this appendix like [Shewchuk 1994] and [Refsnæs 2009].

B.1 Aim

Conjugate Gradient is an algorithm that is typically used to solve linear systems such as:

$$\mathbf{A}x = b \tag{b.1}$$

Where \mathbf{A} is a matrix, b is a (column) vector, and x is the unknown (column) vector.

This problem may be solved also by other methods such as Gaussian Elimination, LU Decomposition, or Cholesky Decomposition, but when the matrix \mathbf{A} is big, those methods are slow and demand much memory. Conjugate Gradient allows to find an (approximate) solution in a faster way and with less memory usage, especially when the matrix \mathbf{A} is sparse. You do not even need to know the full matrix \mathbf{A} if somehow you manage to have a function that, given a vector v , returns the matrix-vector product $\mathbf{A} \cdot v$.

To use Conjugate Gradient, the matrix \mathbf{A} must be symmetrical and positive definite (the algorithm is able to converge also in case of symmetrical and positive semidefinite matrices [Hayami 2018]).

To have a quick refresh of algebra: a symmetric matrix \mathbf{A} of dimension $n \times n$ is positive definite if and only if:

$$v^T \cdot \mathbf{A} \cdot v > 0$$

for every column vector $v \neq 0$ in \mathbb{R}^n

(b.2)

An alternative definition is: a square matrix is positive definite if it is symmetric and all its eigenvalues are positive.

Property: a positive definite matrix is invertible, and its determinant is positive.

B.2 The Equivalent Problem

When the matrix \mathbf{A} is symmetrical and positive definite, finding the solution for $\mathbf{Ax} = \mathbf{b}$ is equivalent to minimize the function:

$$f(x) = \frac{1}{2}x^T \mathbf{A} x - x^T \mathbf{b}$$
(b.3)

The equivalence of the two problems is evident noting that one condition to find the minimum x^* of $f(x)$ is that it must be a critical point, which means that the gradient of $f(x)$ with respect to x must be the zero vector:

$$\nabla_x f(x^*) = \mathbf{A} x^* - \mathbf{b} = 0$$
(b.4)

That is equivalent to eq. b.1 .

To see that this critical point is also a minimum it is sufficient to note that the Hessian of $f(x)$ has all positive eigenvalues, hence it is a “concave up” function, that implies that a critical point must be the minimum. In fact the Hessian is:

$$\nabla_{x_i x_j} f(x^*) = \mathbf{A}$$
(b.5)

We know by assumption that \mathbf{A} is positive definite, that implies that all its eigenvalues are positive.

Therefore, we just proved that the vector x^* that minimizes eq. b.3 is also the solution to the linear system of eq. b.1 , under the condition b.2.

The Conjugate Gradient algorithm allows us to solve the minimization of function b.3 through a particular flavour of steepest descent algorithm, and that can be used whenever we want to solve a linear system in which the matrix is symmetrical positive definite.

B.3 The Steepest Descent

The Conjugate Gradient algorithm can be seen as a special kind of steepest descent. Let us see first thing a standard steepest descent algorithm and then we will see how the Conjugate Gradient improves on that.

In the steepest descent algorithm, we begin assigning a starting value for the vector x , and we iteratively add another vector to it until it reaches the solution x^* , or a close approximation of it. In formulas, starting with initial iteration $k = 0$ and an arbitrary vector $x_n \in \mathbb{R}^n$:

$$x_{k+1} = x_k + \alpha_k p_k \quad (\text{b.6})$$

where α_k is a scalar and $p_k \in \mathbb{R}^n$ is a direction vector. The initial vector x_0 may be set to the zero vector if we do not know any starting point who is allegedly closer to the solution than the zero vector.

The issue here is to find the right α_k and p_k such that x_{k+1} improves the objective compared to x_k . Since we are minimizing $f(x)$, that means:

$$f(x_{k+1}) < f(x_k) \quad (\text{b.7})$$

We do not want just that, we would also like the improvement to be the biggest that can be done, so to do as few iterations as possible.

If you already know methods like gradient descent (and you should, as a prerequisite to study Deep Reinforcement Learning), you are aware that the gradient of the function $f(x)$ is the direction of greatest increase of $f(x)$, hence the negative gradient is the direction of steepest descent. So we have found the direction, and we also already know that the gradient of $f(x)$ is $\mathbf{A}x - b$.

$$p_k = -(\mathbf{A}x_k - b) \quad (\text{b.8})$$

When the problem is solved we have $\mathbf{A}x = b$, when the problem is not yet solved $\mathbf{A}x$ is a vector different from b . The difference between b and this vector is called *residual*, and indicated with r_k :

$$r_k = b - \mathbf{A}x_k = p_k \quad (\text{b.9})$$

That implies:

$$x_{k+1} = x_k - \alpha_k \nabla_x f(x_k) = x_k - \alpha_k (\mathbf{A} x_k - b) = x_k + \alpha_k r_k \quad (\text{b.10})$$

Now we need to find the scalar α_k . To do that, we plug x_{k+1} of eq. b.6 into $f(x)$ of eq. b.3 and we minimize for α_k .

Doing so, we obtain:

$$\begin{aligned} f(x_k + \alpha_k p_k) &= \frac{1}{2} (x_k + \alpha_k p_k)^T \mathbf{A} (x_k + \alpha_k p_k) - (x_k + \alpha_k p_k)^T b \\ &= \frac{1}{2} p_k^T \mathbf{A} p_k \alpha_k^2 + \alpha_k p_k^T (\mathbf{A} x_k - b) + \dots \end{aligned}$$

where the three dots "..." represent other terms that do not contain α_k (b.11)

So, we obtained a quadratic at α_k , that is concave up, so we just need to set its first derivative to zero to minimize it:

$$p_k^T \mathbf{A} p_k \alpha_k + p_k^T (\mathbf{A} x_k - b) = 0$$

$$\alpha_k = -\frac{p_k^T (\mathbf{A} x_k - b)}{p_k^T \mathbf{A} p_k} = \frac{p_k^T r_k}{p_k^T \mathbf{A} p_k} \quad (\text{b.12})$$

This is a general formula. In steepest gradient we have also that $p_k = r_k$, so we have:

$$\alpha_k = \frac{r_k^T r_k}{r_k^T \mathbf{A} r_k} \quad (\text{b.13})$$

Doing this for enough iterations will converge to the solution.

One issue with this method is that it often produces a jagged path toward the solution (such as in fig. 6.2, but possibly even more jagged since in steepest descent two consecutive directions are orthogonal) that means that it may not be the fastest. The Conjugate Gradient algorithm, while following a similar iterative increment strategy, adopts some principled improvements that usually allows us to reach the solution quicker.

B.4 The Conjugate Gradient

Now we need to introduce the concept of “*conjugate vectors*”.

Given a positive definite $n \times n$ matrix \mathbf{A} , two vectors u and $v \in \mathbb{R}^n$ are said to be mutually \mathbf{A} – conjugate if and only if:

$$u^T \cdot \mathbf{A} \cdot v = 0 \quad (\text{b.15})$$

If u and v are \mathbf{A} – conjugate, then any two other vectors respectively parallel to u and v are also \mathbf{A} – conjugate: the direction matters, not the length. For this reason it is sometimes used the term “*conjugate directions*” rather than “*conjugate vectors*”. Also, sometimes it is used the term \mathbf{A} – orthogonal instead of \mathbf{A} – conjugate, that is motivated by the fact that u is orthogonal to the vector resulting from the multiplication $\mathbf{A} \cdot v$ (their dot product equals 0, by definition, and this implies orthogonality).

Property: any set of mutually \mathbf{A} – conjugate vectors is linearly independent.

So, if we have a set of n mutually \mathbf{A} – conjugate vectors $H = \{h_0, h_1, \dots, h_{n-1}\}$, that set is also a basis for \mathbb{R}^n . This allows us to express any vector as a linear combination of that basis, for instance we can express the difference between the solution x^* and the initial point x_0 as a linear combination:

$$x^* - x_0 = \gamma_0 h_0 + \gamma_1 h_1 + \dots + \gamma_{n-1} h_{n-1} \quad (\text{b.16})$$

With $\gamma_0 \dots \gamma_{n-1}$ scalars.

This may be rewritten as:

$$x^* = x_0 + \gamma_0 h_0 + \gamma_1 h_1 + \dots + \gamma_{n-1} h_{n-1} \quad (\text{b.17})$$

Written in this way it is noticeable that, starting from x_0 , the solution x^* can be reached in an iterative way summing a vector at each iteration, exactly like we were doing in steepest descent, but this time the vectors are forming a basis so we know that we need just n iterations (if we know the basis and the scalars). If we compare eq. b.17 with eq. b.6 we have γ_k in place of α_k and h_k in place of p_k . So, we can rewrite the formula replacing them (but keep note that they now will be computed in a different manner than in steepest descent):

$$x^* = x_0 + \alpha_0 p_0 + \alpha_1 p_1 + \dots + \alpha_{n-1} p_{n-1} \quad (\text{b.18})$$

At this point we need to find the basis p_k and the coefficients α_k . One way to find a basis could be to compute the eigenvectors, but that would be computationally expensive.

We introduced the concept of conjugate vectors because they have a useful property: each new \mathbf{A} – conjugate vector p_k can be calculated using the previous \mathbf{A} – conjugate vector p_{k-1} and the residual r_k , without needing any other previous \mathbf{A} – conjugate vector. The whole derivation of this property and its consequences would be very long so we will just write down the formulas here, and the reader may find proofs and theory in [Shewchuk 1994].

Remembering that $r_k = b - \mathbf{A} x_k$, we have:

$$p_k = r_k + \beta_k p_{k-1} \quad (\text{b.19})$$

Since r_k is the gradient of original function (b.3) to be minimized, and we are using conjugate vectors, the algorithm has been named “Conjugate Gradient”.

Because at the beginning of the computations we do not have any already existing vector to conjugate with, the formula b.15 does not apply to p_0 (clearly p_{-1} does not exist): to compute p_0 we just compute it as in steepest descent:

$$p_0 = r_0 \quad (\text{b.20})$$

Because p_k is now computed differently than in steepest descent, also α_k is different:

$$\alpha_k = \frac{r_k^T r_k}{p_k^T \mathbf{A} p_k} \quad (\text{b.21})$$

Now we need to compute β_k , using the condition of \mathbf{A} – conjugacy, and again the reader may find the detailed explanation on it in [Shewchuk 1994]:

$$\beta_k = \frac{r_k^T \mathbf{A} p_{k-1}}{p_{k-1}^T \mathbf{A} p_{k-1}} = \frac{r_k^T r_k}{r_{k-1}^T r_{k-1}} \quad (\text{b.22})$$

Moreover, please note that:

$$r_{k+1} = b - \mathbf{A} x_{k+1} = b - \mathbf{A} x_k - \alpha_k \mathbf{A} p_k = r_k - \alpha_k \mathbf{A} p_k \quad (\text{b.23})$$

Wrapping up, we may now write the algorithm:

Algorithm b.1 Conjugate Gradient

Require: $n \times n$ symmetric positive definite matrix \mathbf{A}

Require: b , a (column) vector of size n

Require: x_0 , a (column) vector of size n (optional, if absent assign the zero vector)

$$r_0 = b - \mathbf{A} x_0$$

$$p_0 = r_0$$

Do for $k = 0, 1, \dots$ until (or close to) convergence:

$$\alpha_k = \frac{r_k^T r_k}{p_k^T \mathbf{A} p_k}$$

$$x_{k+1} = x_k + \alpha_k p_k$$

$$r_{k+1} = r_k - \alpha_k \mathbf{A} p_k$$

$$\beta_k = \frac{r_{k+1}^T r_{k+1}}{r_k^T r_k}$$

$$p_{k+1} = r_{k+1} + \beta_k p_k$$

end of do

One notable aspect of the algorithm is that the matrix \mathbf{A} is used only to do one thing: multiply vectors of size n . So, hypothetically it is possible to use the Conjugate Gradient algorithm even if we do not know the full matrix \mathbf{A} , as long as we have the availability of a function that receives a vector v of size n as input and returns the multiplication $\mathbf{A} \cdot v$.

To speed up a little the algorithm it is easy to see in the inner loop that the matrix multiplication $\mathbf{A} p_k$ is computed twice, so it is possible to compute it once and store in a variable for reuse, and the same can be said of $r_k^T r_k$.

Appendix C: Importance Sampling

If we have a random variable X with probability density function $P(x)$, and we have a function of X named $g(x)$ (which is itself a random variable too, being a function of a random variable), we can compute the expectation of $g(x)$ through an integral:

$$E_{x \sim P(x)}[g(x)] = \int_{-\infty}^{\infty} g(x) P(x) dx \quad (\text{c.1})$$

If for some reason we cannot compute the integral or if it would be too expensive to do it, we can approximate it by sampling. We need the possibility to observe a process where X is distributed following $P(x)$ and then it is processed by $g(x)$, where for each sample x_i we can observe what is the resulting value of $g(x_i)$.

In that way we can compute the sample average of $g(x)$:

$$\bar{g}(x) = \frac{1}{N} \sum_{i=1}^N g(x_i)$$

with $x_i \sim P(x)$

(c.2)

Now, imagine that there is not any process to observe where X is distributed following $P(x)$. So, we cannot sample $g(x_i)$ and compute the average as in c.2. But suppose that we can instead observe a process where X is distributed following a known, different, probability density function $Q(x)$, and we can observe the corresponding samples x_i and the result of $g(x_i)$. Can we use that process to compute an approximation for the expected value $E_{x \sim P(x)}[g(x)]$? Sure we can! We just need to weigh each $g(x_i)$ by the ratio of $P(x_i)$ to $Q(x_i)$, to balance the fact that any sample may occur with different probability in $P(x_i)$ than in $Q(x_i)$. Intuitively, samples that would occur more frequently in $P(x_i)$ than in $Q(x_i)$ should be taken more in account than samples that would occur more rarely in $P(x_i)$ than in $Q(x_i)$. So, the ratio $P(x_i)/Q(x_i)$ is like representing the *importance* of the sample, hence the name “importance sampling”. The importance sampling estimate of the average of $g(x)$ is the following $\bar{g}(x)$:

$$\bar{g}(x) = \frac{1}{N} \sum_{i=1}^N \frac{P(x_i)}{Q(x_i)} g(x_i)$$

with $x_i \sim Q(x)$

(c.3)

A mathematical proof to justify that can be derived from c.1:

$$E_{x \sim P(x)}[g(x)] = \int_{-\infty}^{\infty} g(x) P(x) dx$$

since $\frac{Q(x)}{Q(x)}$ is equal to 1, we can insert it into the integral as a multiplicative factor:

$$= \int_{-\infty}^{\infty} g(x) P(x) \frac{Q(x)}{Q(x)} dx$$

$$= \int_{-\infty}^{\infty} \frac{P(x)}{Q(x)} g(x) Q(x) dx$$

$$= E_{x \sim Q(x)} \left[\frac{P(x)}{Q(x)} g(x) \right]$$

(c.4)

Now, if you sample from c.4 you obtain the equation c.3 .

There are some caveats with importance sampling: since the sampling distribution $Q(x)$ is different from the target distribution $P(x)$, importance sampling is theoretically guaranteed to give the right result only with an infinite number of samples, because it is unbiased but may have a big variance. In practice, with limited samples, if the two distributions differ much you could end up not having samples for values of x that have a great probability density function $P(x)$ but are sampled rarely because they have a low probability density function $Q(x)$. This will distort the result. The opposite is also bad: if an x with very low $Q(x)$ and very high $P(x)$ is sampled by chance, the ratio $P(x)/Q(x)$ could be a number very big and distort the computation of $\bar{g}(x)$.

References

- [Abbeel and Ng 2004] Pieter Abbeel, Andrew Y. Ng, "Apprenticeship Learning via Inverse Reinforcement Learning." In Proceedings of the 21st International Conference on Machine Learning (ICML), 2004
- [Achiam & OpenAi 2020A] Joshua Achiam & OpenAi
Vanilla policy gradient and Expected Grad-Log-Prob Lemma:
https://spinningup.openai.com/en/latest/spinningup/rl_intro3.html
- [Achiam & OpenAi 2020B] Joshua Achiam & OpenAi
Reward-to-go proof
https://spinningup.openai.com/en/latest/spinningup/extra_pg_proof1.html
- [Achiam & OpenAi 2020C] Joshua Achiam & OpenAi
Q-function in policy gradient proof:
https://spinningup.openai.com/en/latest/spinningup/extra_pg_proof2.html
- [Achiam & OpenAi 2020D] Joshua Achiam & OpenAi
TRPO: Matrix-vector product Hx
<https://spinningup.openai.com/en/latest/algorithms/trpo.html>
- [Achiam & OpenAi 2020E] Joshua Achiam & OpenAi
DDPG using gaussian noise.
<https://spinningup.openai.com/en/latest/algorithms/ddpg.html>
- [Amari 1998] Amari, S. "Natural gradient works efficiently in learning." *Neural Computation*, 10, 251. (1998).
- [Andrychowicz et al. 2020] Marcin Andrychowicz, Anton Raichuk, Piotr Stańczyk, Manu Orsini, Sertan Girgin, Raphael Marinier, Léonard Hussenot, Matthieu Geist, Olivier Pietquin, Marcin Michalski, Sylvain Gelly, Olivier Bachem, "What Matters In On-Policy Reinforcement Learning? A Large-Scale Empirical Study", <https://arxiv.org/abs/2006.05990>
- [Askell et al. 2021] A. Askell, Y. Bai, A. Chen, D. Drain, D. Ganguli, T. Henighan, A. Jones, N. Joseph, B. Mann, N. DasSarma, N. Elhage, Z. Hatfield-Dodds, D. Hernandez, J. Kernion, K. Ndousse, C. Olsson, D. Amodei, T. B. Brown, J. Clark, S. McCandlish, C. Olah, and J. Kaplan. "A general language assistant as a laboratory for alignment." CoRR, abs/2112.00861, 2021. URL <https://arxiv.org/abs/2112.00861>.
- [Bagnell and Schneider 2003] Bagnell, J., & Schneider, J. "Covariant policy search." In Proceedings of the international joint conference on artificial intelligence (pp. 1019–1024). (2003).
- [Bai et al. 2022a]
Y. Bai, A. Jones, K. Ndousse, A. Askell, A. Chen, N. DasSarma, D. Drain, S. Fort, D. Ganguli, T. Henighan, N. Joseph, S. Kadavath, J. Kernion, T. Conerly, S. El Showk, N. Elhage, Z. Hatfield-Dodds, D. Hernandez, T. Hume, S. Johnston, S. Kravec, L. Lovitt, N. Nanda, C. Olsson, D. Amodei, T. B. Brown, J. Clark, S. McCandlish, C. Olah, B. Mann, and J. Kaplan. "Training a helpful and harmless assistant with reinforcement learning from

human feedback." CoRR, abs/2204.05862, 2022. doi: 10.48550/arXiv.2204.05862. URL <https://doi.org/10.48550/arXiv.2204.05862>.

[Bai et al. 2022b] Y. Bai, S. Kadavath, S. Kundu, A. Askell, J. Kernion, A. Jones, A. Chen, A. Goldie, A. Mirhoseini, C. McKinnon, C. Chen, C. Olsson, C. Olah, D. Hernandez, D. Drain, D. Ganguli, D. Li, E. Tran-Johnson, E. Perez, J. Kerr, J. Mueller, J. Ladish, J. Landau, K. Ndousse, K. Lukosuite, L. Lovitt, M. Sellitto, N. Elhage, N. Schiefer, N. Mercado, No. DasSarma, R. Lasenby, R. Larson, S. Ringer, S. Johnston, S. Kravec, S. El Showk, S. Fort, T. Lanham, T. Telleen-Lawton, T. Conerly, T. Henighan, T. Hume, S. R. Bowman, Z. Hatfield-Dodds, B. Mann, D. Amodei, N. Joseph, S. McCandlish, T. Brown, J. Kaplan "Constitutional AI: Harmlessness from AI Feedback" <https://arxiv.org/abs/2212.08073>

[Baird 1993] Baird, L. C. "Advantage Updating." Wright Lab. 1993, Technical Report WL-TR-93-1146.

[Bellman 1952] Bellman, R. "On the Theory of Dynamic Programming". Proc Natl Acad Sci U S A. 38 (8): 716–9 , August 1952. doi:10.1073/pnas.38.8.716. PMC 1063639. PMID 16589166.

[Böhm et al. 2019] "Better rewards yield better summaries: Learning to summarise without references" Böhm, F., Gao, Y., Meyer, C. M., Shapira, O., Dagan, I., and Gurevych, I. arXiv preprint arXiv:1909.01214.

[Casella and Berger 2002] George Casella, Roger Berger "Statistical Inference" 2nd Edition, Duxbury 2002

[Christiano et al. 2017] "Deep reinforcement learning from human preferences." Christiano, P. F., Leike, J., Brown, T., Martic, M., Legg, S., and Amodei, D. In Advances in Neural Information Processing Systems, pages 4299–4307.

[DeepSeek-AI 2024] "DeepSeek-R1: Incentivizing Reasoning Capability in LLMs via Reinforcement Learning" https://github.com/deepseek-ai/DeepSeek-R1/blob/main/DeepSeek_R1.pdf

[Degris et al. 2012] Degris, T., White, M., and Sutton, R. S. "Linear off-policy actor-critic." In 29th International Conference on Machine Learning.

[Doersch 2016] Carl Doersch, "Tutorial on Variational Autoencoders" <https://arxiv.org/abs/1606.05908>

[Engstrom et al. 2020] Logan Engstrom, Andrew Ilyas, Shibani Santurkar, Dimitris Tsipras, Firdaus Janoos, Larry Rudolph, Aleksander Madry "Implementation Matters in Deep Policy Gradients: A Case Study on PPO and TRPO", <https://arxiv.org/abs/2005.12729>

[Fujimoto et al. 2018] Scott Fujimoto, Herke van Hoof, David Meger "Addressing Function Approximation Error in Actor-Critic Methods" <https://arxiv.org/pdf/1802.09477>

[Glaese et al. 2022] Amelia Glaese, Nat McAleese, Maja Trębacz, John Aslanides, Vlad Firoiu, Timo Ewalds, Maribeth Rauh, Laura Weidinger, Martin Chadwick, Phoebe Thacker,

Lucy Campbell-Gillingham, Jonathan Uesato, Po-Sen Huang, Ramona Comanescu, Fan Yang, Abigail See, Sumanth Dathathri, Rory Greig, Charlie Chen, Doug Fritz, Jaume Sanchez Elias, Richard Green, Soňa Mokrá, Nicholas Fernando, Boxi Wu, Rachel Foley, Susannah Young, Iason Gabriel, William Isaac, John Mellor, Demis Hassabis, Koray Kavukcuoglu, Lisa Anne Hendricks, Geoffrey Irving
"Improving alignment of dialogue agents via targeted human judgements"
<https://arxiv.org/abs/2209.14375>

[Glorot and Bengio 2010] Glorot, X., & Bengio, Y. "Understanding the difficulty of training deep feedforward neural networks. ", In Proceedings of the thirteenth international conference on artificial intelligence and statistics (pp. 249-256). JMLR Workshop and Conference Proceedings. (2010)

[Hafner and Riedmiller 2011] Roland Hafner, Martin Riedmiller
"Reinforcement learning in feedback control", Mach Learn (2011) 84:137–169, DOI 10.1007/s10994-011-5235-x
<https://link.springer.com/content/pdf/10.1007/s10994-011-5235-x.pdf>

[Haarnoja et al. 2018A] Tuomas Haarnoja, Aurick Zhou, Pieter Abbeel, Sergey Levine. "Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a Stochastic Actor".
<https://arxiv.org/abs/1801.01290>

[Haarnoja et al. 2018B] Tuomas Haarnoja, Aurick Zhou, Kristian Hartikainen, George Tucker, Sehoon Ha, Jie Tan, Vikash Kumar, Henry Zhu, Abhishek Gupta, Pieter Abbeel, Sergey Levine, "Soft Actor-Critic Algorithms and Applications"
<https://arxiv.org/abs/1812.05905>

[Hayami 2018] Ken Hayami, "Convergence of the Conjugate Gradient Method on Singular Systems", <https://arxiv.org/abs/1809.00793>

[Hestenes and Stiefel 1952] Hestenes, Magnus R.; Stiefel, Eduard. "Methods of Conjugate Gradients for Solving Linear Systems" (PDF). Journal of Research of the National Bureau of Standards. 49 (6): 409. (December 1952) doi:10.6028/jres.049.044.
<http://nvlpubs.nist.gov/nistpubs/jres/049/6/V49.N06.A08.pdf>

[Hu et al. 2024] Jian Hu, Xibin Wu, Weixun Wang, Xianyu, Dehao Zhang, Yu Cao, "OpenRLHF: An Easy-to-use, Scalable and High-performance RLHF Framework",
<https://arxiv.org/abs/2405.11143v1>

[Huang et al. 2022] Huang S., Dossa, R., Fernand J., Raffin A., Kanervisto A., Wang W., "The 37 Implementation Details of Proximal Policy Optimization"
<https://iclr-blog-track.github.io/2022/03/25/ppo-implementation-details/>

[Huang et al. 2023] S. C. Huang, T. Liu, L. von Werra, "The N Implementation Details of RLHF with PPO "
https://huggingface.co/blog/the_n_implementation_details_of_rlhf_with_ppo
Code: https://github.com/vwxyzjn/summarize_from_feedback_details

[Huang et al. 2024] S. Huang, M. Noukhovitch, A. Hosseini, K. Rasul, W. Wang, L. Tunstall "The N+ Implementation Details of RLHF with PPO: A Case Study on TL;DR Summarization", <https://arxiv.org/pdf/2403.17031>

[Ibarz et al. 2018] "Reward learning from human preferences and demonstrations in atari."

Ibarz, B., Leike, J., Pohlen, T., Irving, G., Legg, S., and Amodei, D. In Advances in neural information processing systems, pages 8011–8023.

[Ioffe and Szegedy 2015] Ioffe, Sergey; Szegedy, Christian. "Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift", ICML'15: Proceedings of the 32nd International Conference on International Conference on Machine Learning - Volume 37, July 2015 Pages 448–456

[Irpan 2018] Alex Irpan, "Deep Reinforcement Learning Doesn't Work Yet", (2018)
<https://www.alexirpan.com/2018/02/14/rl-hard.html>

[Jones 2021] Andy L. Jones, "Debugging RL, Without the Agonizing Pain"
<https://andyljones.com/posts/rl-debugging.html>

[Kakade 2002] Kakade Sham, "A natural policy gradient", Advances in Neural Information Processing Systems, pp. 1057–1063. MIT Press, 2002.

[Kakade and Langford 2002] Kakade, Sham and Langford, John. "Approximately optimal approximate reinforcement learning" In ICML, volume 2, pp. 267–274, 2002.

[Kingma and Welling 2014] Diederik P Kingma, Max Welling, "Auto-Encoding Variational Bayes"
<https://arxiv.org/pdf/1312.6114>

[Kristiadi 2018] Kristiadi Augustinus, "Fisher Information"
<https://agustinus.kristia.de/techblog/2018/03/11/fisher-information/>

[LeCun et al. 1998] LeCun, Y.; Bottou, L.; Bengio, Y.; Haffner, P. (1998). "Gradient-based learning applied to document recognition" (PDF). Proceedings of the IEEE. 86 (11): 2278–2324. doi:10.1109/5.726791. S2CID 14542261.

[Levine 2021a] Sergey Levine, CS 285 at UC Berkeley, Deep Reinforcement Learning, Lecture 9 (part 4)
<http://rail.eecs.berkeley.edu/deeprlcourse-fa21/static/slides/lec-9.pdf>
<https://www.youtube.com/watch?v=QWnpF0FaKL4>

[Levine 2021b] Sergey Levine, CS 285 at UC Berkeley, Deep Reinforcement Learning, Lecture 15 (part 2)
<https://rail.eecs.berkeley.edu/deeprlcourse/static/slides/lec-15.pdf>
<https://www.youtube.com/watch?v=9HrN6nHoxD8>

[Levine 2021c] Sergey Levine, CS 285 at UC Berkeley, Deep Reinforcement Learning, Lecture 5
<https://rail.eecs.berkeley.edu/deeprlcourse/static/slides/lec-5.pdf>
<https://www.youtube.com/watch?v=KZd508qGFt0>

[Levine 2021d] Sergey Levine, CS 285 at UC Berkeley, Deep Reinforcement Learning, Lecture 2
<https://rail.eecs.berkeley.edu/deeprlcourse/deeprlcourse/static/slides/lec-2.pdf>

[Levine and Koltun 2013] Sergey Levine and Vladlen Koltun, "Guided Policy Search", ICML 2013, PMLR 28(3):1-9, 2013.
<http://proceedings.mlr.press/v28/levine13.pdf>

[Levin et al. 2009] Levin, D. A., Peres, Y., and Wilmer, E. L. Markov chains and

mixing times. American Mathematical Society, 2009.

[Levine et al. 2020] Sergey Levine, Aviral Kumar, George Tucker, Justin Fu, "Offline Reinforcement Learning: Tutorial, Review, and Perspectives on Open Problems" <https://arxiv.org/abs/2005.01643>

[Lillicrap et al. 2016] Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver & Daan Wierstra "Continuous Control with Deep Reinforcement Learning" <https://arxiv.org/abs/1509.02971>

[Liu 2023] Jiacheng Liu, "The embarrassing redundancy of reward whitening and reward normalization in PPO" <https://liujch1998.github.io/2023/04/16/ppo-norm.html>

[Liu et al. 2019] Yao Liu, Adith Swaminathan, Alekh Agarwal, Emma Brunskill "Off-Policy Policy Gradient with State Distribution Correction" <https://arxiv.org/abs/1904.08473>

[Menick et al 2022] Jacob Menick, Maja Trebacz, Vladimir Mikulik, John Aslanides, Francis Song, Martin Chadwick, Mia Glaese, Susannah Young, Lucy Campbell-Gillingam, Geoffrey Irving and Nat McAleese, "Teaching language models to support answers with verified quotes", <https://arxiv.org/pdf/2203.11147>

[Mnih et al. 2013] Mnih, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., & Riedmiller, M. , "Playing Atari with Deep Reinforcement Learning." (2013), arXiv:1312.5602.

[Mnih et al. 2016] Volodymyr Mnih, Adrià Puigdomènech Badia, Mehdi Mirza, Alex Graves, Timothy P. Lillicrap, Tim Harley, David Silver, Koray Kavukcuoglu "Asynchronous Methods for Deep Reinforcement Learning" <https://arxiv.org/abs/1602.01783>

[Ng and Russell 2000] Andrew Y. Ng, Stuart Russell, "Algorithms for Inverse Reinforcement Learning". ICML '00: Proceedings of the Seventeenth International Conference on Machine Learning June 2000, Pages 663 - 670

[Ouyang et al. 2022] Long Ouyang, Jeff Wu, Xu Jiang, Diogo Almeida, Carroll L. Wainwright, Pamela Mishkin, Chong Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, John Schulman, Jacob Hilton, Fraser Kelton, Luke Miller, Maddie Simens, Amanda Askell, Peter Welinder, Paul Christiano, Jan Leike, Ryan Lowe, "Training language models to follow instructions with human feedback". <https://arxiv.org/abs/2203.02155>

[Peters 2007] Peters, J. "Machine learning of motor skills for robotics. "Ph.D. thesis University of Southern California, Los Angeles, CA, 90089, USA. (2007).

[Peters and Schaal 2008] Jan Peters, Stefan Schaal , "Reinforcement learning of motor skills with policy gradients." , Neural networks, 21(4), 682-697. (2008)

[Peters et al. 2003] Peters, J., Vijayakumar, S., & Schaal, S. "Reinforcement learning for humanoid robotics." In Proceedings of the IEEE-RAS international conference on humanoid robots (HUMANOIDS) (pp. 103–123).(2003)

[Pollard 2000] Pollard, David. "Asymptopia: an exposition of statistical asymptotic theory." 2000. URL <http://www.stat.yale.edu/~pollard/Books/Asymptopia>.

[Psu 2024] PennState University, "STAT 414: Introduction to Probability Theory"
<https://online.stat.psu.edu/stat414/lesson/22/22.2>

[Rahtz 2021] Matthew Rahtz - "Learning From Human Preferences, Implementation :
Adjusting softmax function"
<https://github.com/mrahtz/learning-from-human-preferences/issues/8>

[Ratliff 2013] Ratliff, Nathan. "Information Geometry and Natural Gradients" in Mathematics for Intelligent Systems.
<https://www.nathanratliff.com/pedagogy/mathematics-for-intelligent-systems>
<https://drive.google.com/file/d/1jDM9yZl1KrtH3JJznPE7S1QzNZZB2xs2/view>
https://web.archive.org/web/20200924055016/https://ipvs.informatik.uni-stuttgart.de/mlr/wp-content/uploads/2015/01/mathematics_for_intelligent_systems_lecture12_notes_l.pdf

[Ramamurthy et al. 2023] Rajkumar Ramamurthy, Prithviraj Ammanabrolu, Kianté Brantley, Jack Hessel, Rafet Sifa, Christian Bauckhage, Hannaneh Hajishirzi, Yejin Choi
"Is Reinforcement Learning (not) for Natural Language Processing: Benchmarks, Baselines and building Blocks for Natural Language Policy Optimization."
<https://arxiv.org/abs/2210.01241> , <https://rl4lms.apps.allenai.org/algorithms>

[Refsnæs 2009] Runar Heggelien Refsnæs, "A Brief Introduction to the Conjugate Gradient Method", 2009, <https://www.math.hkust.edu.hk/~mamu/courses/531/cg.pdf>

[Rioul 2018] Olivier Rioul, "This is IT: A Primer on Shannon's Entropy and Information", L'Information, Séminaire Poincaré XXIII (2018) 43 – 77
<https://seminaire-poincare.pages.math.cnrs.fr/rioul.pdf>

[Ross et al. 2010] Stephane Ross, Geoffrey J. Gordon, J. Andrew Bagnell, "A Reduction of Imitation Learning and Structured Prediction to No-Regret Online Learning"
<https://arxiv.org/abs/1011.0686>

[Saxe et al. 2013] Saxe, A. M., McClelland, J. L., & Ganguli, S. , "Exact solutions to the nonlinear dynamics of learning in deep linear neural networks." ICLR (2013), arXiv preprint arXiv:1312.6120.

[Schaal 1999] S. Schaal. "Is imitation learning the route to humanoid robots?" In Trends in Cognitive Sciences, 1999.

[Schulman et al. 2015] John Schulman, Sergey Levine, Philipp Moritz, Michael I. Jordan, Pieter Abbeel , "Trust Region Policy Optimization", In International conference on machine learning (pp. 1889-1897). PMLR, arXiv:1502.05477

[Schulman et al. 2016] John Schulman, Philipp Moritz, Sergey Levine, Michael I. Jordan and Pieter Abbeel, "High-Dimensional Continuous Control Using Generalized Advantage Estimation", ICLR 2016, arXiv:1506.02438

[Schulman et al. 2017] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017). , "Proximal policy optimization algorithms" , arXiv preprint arXiv:1707.06347.

[Schulman 2020] J. Schulman. "Approximating kl divergence" 2020.
<http://joschu.net/blog/kl-approx.html>

[Shannon 1948] C. E. Shannon, "A Mathematical Theory of Communication", the Bell System Technical Journal, Vol. 27, pp. 379–423, 623–656, July, October, 1948

<https://people.math.harvard.edu/~ctm/home/text/others/shannon/entropy/entropy.pdf>

[Shao et al. 2024] "Deepseekmath: Pushing the limits of mathematical reasoning in open language models"

Z. Shao, P. Wang, Q. Zhu, R. Xu, J. Song, M. Zhang, Y. Li, Y. Wu, and D. Guo. ,
arXiv:2402.03300, 2024.

[Shewchuk 1994] Jonathan Richard Shewchuk, "An Introduction to the Conjugate Gradient Method Without the Agonizing Pain" (1994)

<https://www.cs.cmu.edu/~quake-papers/painless-conjugate-gradient.pdf>

[Silver et al. 2014] Silver D., Lever G., Heess N., Degris T., Wierstra D., & Riedmiller M. "Deterministic policy gradient algorithms". In International conference on machine learning (pp. 387-395). PMLR. (2014, January)

<https://proceedings.mlr.press/v32/silver14.pdf>

<https://proceedings.mlr.press/v32/silver14-sup.pdf>

[Snell et al. 2023] Charlie Snell, Ilya Kostrikov, Yi Su, Mengjiao Yang, Sergey Levine, "Offline RL for Natural Language Generation with Implicit Language Q Learning",

<https://arxiv.org/abs/2206.11871> ,https://sea-snell.github.io/ILQL_site/

[Soemers 2019] Dennis Soemers, "Reward-to-go proof"

<https://ai.stackexchange.com/questions/9614/why-does-the-reward-to-go-trick-in-policy-gradient-methods-work/10369>

[Statproofbook 2021A] "Proof: Kullback-Leibler divergence for the normal distribution", The Book of Statistical Proofs <https://dx.doi.org/10.5281/zenodo.5820411>

<https://statproofbook.github.io/P/norm-kl.html>

[Statproofbook 2021B] "Proof: Kullback-Leibler divergence for the multivariate normal distribution", The Book of Statistical Proofs <https://dx.doi.org/10.5281/zenodo.5820411>

<https://statproofbook.github.io/P/mvn-kl.html>

[Statproofbook 2021C] "Proof: Differential entropy of the multivariate normal distribution"

<https://dx.doi.org/10.5281/zenodo.5820411>

<https://statproofbook.github.io/P/mvn-dent.html>

[Stiennon et al. 2020] "Learning to summarize from human feedback"

Stiennon, N., Ouyang, L., Wu, J., Ziegler, D. M., Lowe, R., Voss, C., Radford, A., Amodei, D., and Christiano, P. . arXiv preprint arXiv:2009.01325. Code:

<https://github.com/openai/summarize-from-feedback>

[Sutton & Barto 2018] Richard C. Sutton, Andrew C. Barto, "Reinforcement Learning - An introduction. 2nd Ed.", MIT Press - ISBN: 9780262039246

<http://incompleteideas.net/book/RLbook2020.pdf>

[Sutton et al. 1999] Sutton, R. S., McAllester, D. A., Singh, S. P., and Mansour, Y. (1999). "Policy gradient methods for reinforcement learning with function approximation." In Neural Information Processing Systems 12, pages 1057–1063.

[Szepesvári 2010] Csaba Szepesvári, "Algorithms for Reinforcement Learning", Morgan & Claypool 2010, ISBN: 9781608454921

[Thrun and Schwartz 1993] Thrun, S. and Schwartz, A. "Issues in using function approximation for reinforcement learning." In Proceedings of the

1993 Connectionist Models Summer School Hillsdale, NJ. Lawrence Erlbaum, 1993.
https://www.ri.cmu.edu/pub_files/pub1/thrun_sebastian_1993_1/thrun_sebastian_1993_1.pdf

[Touvron et al. 2023] H. Touvron, L. Martin, K. Stone, P. Albert, A. Almahairi, Y. Babaei, N. Bashlykov, S. Batra, P. Bhargava, S. Bhosale, et al.
“Llama 2: Open foundation and fine-tuned chat models.” arXiv preprint arXiv:2307.09288, 2023.

[van Hasselt 2010] Hado van Hasselt, “Double Q-learning.”
Advances in neural information processing systems, 23, 2613-2621. (2010)

[van Hasselt et al. 2016] Van Hasselt, Hado, Arthur Guez, and David Silver. "Deep reinforcement learning with double q-learning." Proceedings of the AAAI conference on artificial intelligence. Vol. 30. No. 1. 2016

[Wang et al. 2023] Peiyi Wang, Lei Li, Zhihong Shao, R.X. Xu, Damai Dai, Yifei Li, Deli Chen, Y.Wu, Zhifang Sui, "Math-Shepherd: Verify and Reinforce LLMs Step-by-step without Human Annotations", <https://arxiv.org/abs/2312.08935>

[Watanabe 2009] Watanabe Sumio, "Algebraic Geometry and Statistical Learning Theory", Cambridge University Press, 2009. ISBN-13 978-0-521-86467-1

[Watkins 1989] Christopher J. C. H. Watkins, “Learning from delayed rewards.”, PhD Thesis, King’s College (1989).

[Williams 1992] Williams, R. J. , “Simple statistical gradient-following algorithms for connectionist reinforcement learning.”, Machine learning, 8(3), 229-256. (1992)

[Wu et al. 2021] "Recursively summarizing books with human feedback."
Wu, J., Ouyang, L., Ziegler, D. M., Stiennon, N., Lowe, R., Leike, J., and Christiano, P. arXiv preprint arXiv:2109.10862.

[Yu et al. 2025] Qiying Yu, Zheng Zhang, Ruofei Zhu, Yufeng Yuan, Xiaochen Zuo, Yu Yue, Tiantian Fan, Gaohong Liu, Lingjun Liu, Xin Liu, Haibin Lin, Zhiqi Lin, Bole Ma, Guangming Sheng, Yuxuan Tong, Chi Zhang, Mofan Zhang, Wang Zhang, Hang Zhu, Jinhua Zhu, Jiaze Chen, Jiangjie Chen, Chengyi Wang, Hongli Yu, Weinan Dai, Yuxuan Song, Xiangpeng Wei, Hao Zhou, Jingjing Liu, Wei-Ying Ma, Ya-Qin Zhang, Lin Yan, Mu Qiao, Yonghui Wu, Mingxuan Wang, "DAPO: An Open-Source LLM Reinforcement Learning System at Scal" <https://arxiv.org/pdf/2503.14476>

[Ziegler et al. 2019] “Fine-Tuning Language Models from Human Preferences.”
Ziegler, D. M., Stiennon, N., Wu, J., Brown, T. B., Radford, A., Amodei, D., Christiano, P., and Irving, G. Fine-tuning language models from human preferences. arXiv preprint arXiv:1909.08593 , code: <https://github.com/openai/lm-human-preferences>